

Date:
March 30, 2014

**Delivered as part of Milestones
7.1, 7.2, 7.3**

User's Guide to FastForward Features in HDF5

FOR EXTREME-SCALE COMPUTING RESEARCH AND DEVELOPMENT (FAST FORWARD) STORAGE AND I/O

LLNS Subcontract No.	B599860
Subcontractor Name	Intel Federal LLC
Subcontractor Address	2200 Mission College Blvd. Santa Clara, CA 95052

NOTICE: THIS MANUSCRIPT HAS BEEN AUTHORED BY THE HDF GROUP UNDER THE INTEL SUBCONTRACT WITH LAWRENCE LIVERMORE NATIONAL SECURITY, LLC WHO IS THE OPERATOR AND MANAGER OF LAWRENCE LIVERMORE NATIONAL LABORATORY UNDER CONTRACT NO. DE-AC52-07NA27344 WITH THE U.S. DEPARTMENT OF ENERGY. THE UNITED STATES GOVERNMENT RETAINS AND THE PUBLISHER, BY ACCEPTING THE ARTICLE OF PUBLICATION, ACKNOWLEDGES THAT THE UNITED STATES GOVERNMENT RETAINS A NON-EXCLUSIVE, PAID-UP, IRREVOCABLE, WORLD-WIDE LICENSE TO PUBLISH OR REPRODUCE THE PUBLISHED FORM OF THIS MANUSCRIPT, OR ALLOW OTHERS TO DO SO, FOR UNITED STATES GOVERNMENT PURPOSES. THE VIEWS AND OPINIONS OF AUTHORS EXPRESSED HEREIN DO NOT NECESSARILY REFLECT THOSE OF THE UNITED STATES GOVERNMENT OR LAWRENCE LIVERMORE NATIONAL SECURITY, LLC.

Revision History.....	iv
1 Introduction	1
2 Definitions.....	1
3 EFF API reference manual pages	2
3.1 Startup and Shutdown APIs.....	3
3.1.1 <i>EFF_finalize</i>	4
3.1.2 <i>EFF_init</i>	5
3.1.3 <i>EFF_start_server</i>	6
4 HDF5 API reference manual pages	7
4.1 H5: General HDF5 APIs	8
4.1.1 <i>H5checksum</i>	9
4.2 H5A: Attribute APIs	11
4.2.1 <i>H5Aclose_ff</i>	13
4.2.2 <i>H5Acreate_ff</i>	14
4.2.3 <i>H5Acreate_by_name_ff</i>	16
4.2.4 <i>H5Adelete_ff</i>	18
4.2.5 <i>H5Adelete_by_name_ff</i>	19
4.2.6 <i>H5Aexists_ff</i>	21
4.2.7 <i>H5Aexists_by_name_ff</i>	23
4.2.8 <i>H5Aopen_ff</i>	25
4.2.9 <i>H5Aopen_by_name_ff</i>	26
4.2.10 <i>H5Aread_ff</i>	28
4.2.11 <i>H5Arename_ff</i>	29
4.2.12 <i>H5Arename_by_name_ff</i>	30
4.2.13 <i>H5Awrite_ff</i>	32
4.3 H5AS: Analysis Shipping APIs	34
4.3.1 <i>H5ASexecute</i>	35
4.4 H5D: Dataset APIs	37
4.4.1 <i>H5Dclose_ff</i>	38
4.4.2 <i>H5Dcreate_ff</i>	39
4.4.3 <i>H5Devict_ff</i>	41
4.4.4 <i>H5Dopen_ff</i>	43
4.4.5 <i>H5Dprefetch_ff</i>	45
4.4.6 <i>H5Dread_ff</i>	47
4.4.7 <i>H5Dset_extent_ff</i>	50
4.4.8 <i>H5Dwrite_ff</i>	51
4.5 H5ES: Event Stack APIs	54
4.5.1 <i>H5Ecancel</i>	55
4.5.2 <i>H5Ecancel_all</i>	56
4.5.3 <i>H5Eclear</i>	57
4.5.4 <i>H5Eclose</i>	58
4.5.5 <i>H5Ecreate</i>	59
4.5.6 <i>H5Eget_count</i>	60
4.5.7 <i>H5Eget_event_info</i>	61

Use or disclosure of data contained on this sheet is subject to the restriction on the title page of this document.

Copyright © The HDF Group, 2014. All rights reserved

4.5.8	<i>H5EStest</i>	63
4.5.9	<i>H5EStest_all</i>	64
4.5.10	<i>H5ESwait</i>	65
4.5.11	<i>H5ESwait_all</i>	66
4.6	H5F: File (Container) APIs.....	67
4.6.1	<i>H5Fclose_ff</i>	69
4.6.2	<i>H5Fcreate_ff</i>	70
4.6.3	<i>H5Fopen_ff</i>	72
4.7	H5G: Group APIs	74
4.7.1	<i>H5Gclose_ff</i>	75
4.7.2	<i>H5Gcreate_ff</i>	76
4.7.3	<i>H5Gevict_ff</i>	78
4.7.4	<i>H5Gopen_ff</i>	80
4.7.5	<i>H5Gprefetch_ff</i>	82
4.8	H5L: Link APIs.....	84
4.8.1	<i>H5Lcopy_ff</i>	85
4.8.2	<i>H5Lcreate_hard_ff</i>	87
4.8.3	<i>H5Lcreate_soft_ff</i>	89
4.8.4	<i>H5Ldelete_ff</i>	91
4.8.5	<i>H5Lexists_ff</i>	93
4.8.6	<i>H5Lmove_ff</i>	95
4.9	H5M: Map APIs	97
4.9.1	<i>H5Mclose_ff</i>	98
4.9.2	<i>H5Mcreate_ff</i>	99
4.9.3	<i>H5Mdelete_ff</i>	101
4.9.4	<i>H5Mevict_ff</i>	103
4.9.5	<i>H5Mexists_ff</i>	105
4.9.6	<i>H5Mget_ff</i>	107
4.9.7	<i>H5Mget_count_ff</i>	109
4.9.8	<i>H5Mget_types_ff</i>	110
4.9.9	<i>H5Miterate_ff</i>	111
4.9.10	<i>H5Mopen_ff</i>	113
4.9.11	<i>H5Mprefetch_ff</i>	115
4.9.12	<i>H5Mset_ff</i>	117
4.10	H5O: Object APIs.....	119
4.10.1	<i>H5Oclose_ff</i>	120
4.10.2	<i>H5Oexists_by_name_ff</i>	121
4.10.3	<i>H5Oget_comment_ff</i>	123
4.10.4	<i>H5Oget_comment_by_name_ff</i>	125
4.10.5	<i>H5Olink_ff</i>	127
4.10.6	<i>H5Oopen_ff</i>	129
4.10.7	<i>H5Oset_comment_ff</i>	130
4.10.8	<i>H5Oset_comment_by_name_ff</i>	132
4.11	H5P: Property APIs.....	134
4.11.1	<i>H5Pget_rcapl_version_request</i>	135
4.11.2	<i>H5Pget_trspl_num_peers</i>	136
4.11.3	<i>H5Pset_dxpl_checksum</i>	137

Use or disclosure of data contained on this sheet is subject to the restriction on the title page of this document.

Copyright © The HDF Group, 2014. All rights reserved

4.11.4	<i>H5Pset_dxpl_checksum_ptr</i>	138
4.11.5	<i>H5Pset_dxpl_inject_corruption</i>	139
4.11.6	<i>H5Pset_read_replica</i>	140
4.11.7	<i>H5Pset_evict_replica</i>	141
4.11.8	<i>H5Pset_fapl_iod</i>	142
4.11.9	<i>H5Pset_metadata_integrity_scope</i>	143
4.11.10	<i>H5Pset_rawdata_integrity_scope</i>	144
4.11.11	<i>H5Pset_rcapl_version_request</i>	145
4.11.12	<i>H5Pset_trspl_num_peers</i>	146
4.12	H5Q: Query APIs	147
4.12.1	<i>H5Qcreate</i>	148
4.12.2	<i>H5Qcombine</i>	150
4.12.3	<i>H5Qclose</i>	152
4.13	H5RC: Read Context APIs	153
4.13.1	<i>H5RCacquire</i>	154
4.13.2	<i>H5RCclose</i>	156
4.13.3	<i>H5RCcreate</i>	157
4.13.4	<i>H5RCget_version</i>	158
4.13.5	<i>H5RCpersist</i>	159
4.13.6	<i>H5RCrelease</i>	160
4.13.7	<i>H5RCsnapshot</i>	161
4.14	H5T: Datatype APIs	162
4.14.1	<i>H5Tclose_ff</i>	163
4.14.2	<i>H5Tcommit_ff</i>	164
4.14.3	<i>H5Tevict_ff</i>	166
4.14.4	<i>H5Topen_ff</i>	168
4.14.5	<i>H5Tprefetch_ff</i>	169
4.15	H5TR: Transaction Operation APIs	171
4.15.1	<i>H5TRabort</i>	172
4.15.2	<i>H5TRclose</i>	173
4.15.3	<i>H5TRcreate</i>	174
4.15.4	<i>H5TRfinish</i>	175
4.15.5	<i>H5TRget_trans_num</i>	177
4.15.6	<i>H5TRget_version</i>	178
4.15.7	<i>H5TRset_dependency</i>	179
4.15.8	<i>H5TRskip</i>	180
4.15.9	<i>H5TRstart</i>	181
5	Description of example programs	183
5.1	h5ff_server.c	183
5.2	h5ff_client_*c	184
6	Instructions for building and running demo code	185

Use or disclosure of data contained on this sheet is subject to the restriction on the title page of this document.

Copyright © The HDF Group, 2014. All rights reserved

Revision History

Date	Revision	Description	Authors
03/22/2013	1.0	Delivered to DOE stakeholders as part of Milestone 3.3 demonstration output.	Ruth Aydt, Mohamad Chaarawi, Quincey Koziol, Jerome Soumagne; The HDF Group
06/18/2013	2.0	Delivered to DOE stakeholders as part of Milestones 4.2, 4.3, 4.4 demonstration output.	Ruth Aydt, Mohamad Chaarawi, Quincey Koziol, Jerome Soumagne; The HDF Group
9/26/2013	3.0	<ul style="list-style-type: none"> • Updated Introduction section for Q5. • Changed placeholder ID to future ID. • Indicated that H5EQ and H5AO routines will be consolidated into H5ES; Added H5ES section and man pages. • Removed H5DO section. • Added H5RC and H5TR sections. • Added several new routines to H5P. • Updated RM section tables and man pages for Q5. See section table notes and History on man pages for details. • Updated instructions on how to build the EFF stack. • Updated test sections with the new testing framework. • Removed Appendix with Q4 example code. Did not include Q5 example code as it is much larger; see source distribution. • Delivered to DOE stakeholders as part of Milestones 5.6, 5.7 demonstration output. 	Ruth Aydt, Mohamad Chaarawi, Quincey Koziol; The HDF Group
12/20/2013	4.0.2	<ul style="list-style-type: none"> • Removed H5EQ and H5AO APIs – now consolidated into H5ES • Updated tables at beginning of each API section to reflect current status. • Updated list of example client programs and build instructions. • Fixed typos • Delivered to DOE stakeholders as part of Milestone 6.2, 6.3 demonstration output 	Ruth Aydt, Mohamad Chaarawi; The HDF Group
12/20/2013	4.0.3	<ul style="list-style-type: none"> • Added H5Qcreate/combine/close • Added H5ASexecute 	Quincey Koziol, Jerome Soumagne; The HDF Group
12/23/2013	4.0.4	<ul style="list-style-type: none"> • Final edit pass • Date on title page reconciled and highlight in section 1 removed without version change on 1/4/14. 	Quincey Koziol; The HDF Group
02/07/2014	5.1	<ul style="list-style-type: none"> • Update data markings and copyright year • Begin changes for Quarter 7 milestones 	Ruth Aydt; The HDF Group

Use or disclosure of data contained on this sheet is subject to the restriction on the title page of this document.

Copyright © The HDF Group, 2014. All rights reserved

		<ul style="list-style-type: none"> Added H5TRget_trans_num and H5TRget_version 	
02/26/2014	5.2	<ul style="list-style-type: none"> Added EFF API reference manual pages. (EFF_finalize, EFF_init, EFF_start_server) Add note to H5Dopen about using object id to access other CVs. 	Ruth Aydt, Mohamad Chaarawi; The HDF Group
03/10/2014	5.3	<ul style="list-style-type: none"> Added persist flag to H5Fclose_ff 	Mohamad Chaarawi, The HDF Group
03/30/2014	5.4	<ul style="list-style-type: none"> Added man pages for prefetch and evict routines and related properties. Noted that H5V and H5X reference manual pages will be added in Q8, with pointer to draft versions in the design doc. Updated list of client examples and build instructions. Updated status in Section 4.x tables Deliver to DOE as part of Q7 milestones. 	Ruth Aydt, Mohamad Chaarawi, The HDF Group

Use or disclosure of data contained on this sheet is subject to the restriction on the title page of this document.

Copyright © The HDF Group, 2014. All rights reserved

1 Introduction

This document is one of the deliverables for these Quarter 7 Milestones:

- Milestone 7.1: HDF5 Index Plugin API Demonstration
- Milestone 7.2: Burst Buffer and Data Integrity Demonstration
- Milestone 7.3: End-to-End Epoch Recovery Demonstration

The document reflects the status of the software stack at the time of the Quarter 7 milestone demonstrations, with the exception of the H5V (view) and H5X (index) APIs. The final reference manual pages for these API routines, which were first implemented in Quarter 7, will be included in the User Guide in Quarter 8. Please refer to Sections 5.2.11 and 5.2.12 in the *"Design and Implementation of FastForward Features in HDF5"* documents for draft reference manual pages for these APIs.

This document provides reference manual pages for API routines, high-level descriptions of example programs, and instructions for building and running both the HDF5 and example packages in the demonstration environment.

A separate document provides additional program notes for the Milestone 7.2 demo program (`h5ff_client_M7.2_demo.c`) and the VPIC I/O application used to demonstrate the corruption detection and end-to-end Epoch recovery for Milestones 7.2 and 7.3

Development will continue in Quarter 8, and both the document and the software will be updated as part of milestones for that quarter.

2 Definitions

ACG – Arbitrarily Connected Graphs

AXE – Asynchronous eXecution Engine

CN – Compute Node

EFF – Exascale Fast Forward

FS – Function Shipper, also known as Mercury

IOD – I/O Dispatcher

ION – I/O Node

Mercury – Name of Function Shipper

VOL – Virtual Object Layer

3 EFF API reference manual pages

This section contains the reference manual pages for new routines that are part of the Exascale Fast Forward I/O stack, but that are not specifically related to HDF5. These routines are used to start the Asynchronous eXecution Engine (AXE) and the function shipper (Mercury) server and clients – components that provide the foundation for the HDF5-IOD VOL plugin operations in the EFF stack.

3.1 Startup and Shutdown APIs

These routines start and stop components of the EFF stack such as the Asynchronous eXecution Engine (AXE) and the function shipper (Mercury).

Routine	Implemented	Notes
EFF_finalize	Quarter 4	
EFF_init	Quarter 4	
EFF_start_server	Quarter 4	

3.1.1 EFF_finalize

Name: EFF_finalize

Signature:

herr_t **EFF_finalize(void)**

Purpose:

Shut down the EFF stack.

Description:

EFF_finalize shuts down the EFF stack, stopping the Asynchronous eXecution Engine (AXE), finalizing IOD, and shutting down the function shipper (Mercury) that were started by a call to **EFF_init**.

The application must call **EFF_finalize** after all calls to the HDF5 routines had been completed. The normal sequence of calls at the end of an application running on the EFF stack is:

```
MPI_Barrier( MPI_COMM_WORLD );
EFF_finalize();
MPI_Finalize();
```

Parameters:

none

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

History:

Routine added in Quarter 4.

Man page added in Quarter 7.

Man Page Status:

No known issues.

3.1.2 EFF_init

Name: EFF_init

Signature:

herr_t **EFF_init(MPI_Comm comm, MPI_Info info)**

Purpose:

Initialize the EFF stack.

Description:

EFF_init initializes the EFF stack by setting up the connection to the Mercury server running the HDF5 VOL server, so that the HDF5 IOD VOL plugin at the client side can transparently forward requests to it. The call also triggers the initialization of the IOD library.

The application must call **EFF_init** after the function shipping (Mercury) server has been started and before the application makes any calls to the HDF5 routines.

comm is the MPI communicator for the compute processes.

info passes hints to the MPI library. **MPI_INFO_NULL** can be used if no hints need to be passed..

Parameters:

MPI_Comm **comm** IN: MPI communicator

MPI_Info **info** IN: MPI info.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

History:

Added in Quarter 4.

Man page added in Quarter 7.

Man Page Status:

No known issues.

3.1.3 EFF_start_server

Name: `EFF_start_server`

Signature:

`herr_t EFF_start_server(MPI_Comm comm, MPI_Info info)`

Purpose:

Start the function shipper server that is used with the EFF stack.

Description:

`EFF_start_server` starts the function shipper server that listens for incoming connections from clients. Once a connection has been established, the client will transparently forward calls to the server it has connected to.

`EFF_start_server` must be called in the server code before the user's application calls `EFF_init`.

`comm` is the MPI communicator for the compute processes.

`info` passes hints to the MPI library. `MPI_INFO_NULL` can be used if no hints need to be passed.

Parameters:

`MPI_Comm comm` IN: MPI communicator

`MPI_Info info` IN: MPI info.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

History:

Added in Quarter 4.

Man page added in Quarter 7.

Man Page Status:

No known issues.

4 HDF5 API reference manual pages

This section contains the reference manual pages for the modified and new HDF5 API routines that are part of the Fast Forward extensions to HDF5.

Tables at the beginning of each section also note routines that were implemented in the HDF5 IOD VOL client/server as part of the Fast Forward extensions to HDF5 for which the user interface is unchanged. Readers are directed to the standard HDF5 man pages, found at http://www.hdfgroup.org/HDF5/doc/RM/RM_H5Front.html, for those routines.

4.1 H5: General HDF5 APIs

These routines serve general-purpose needs of the HDF5 library and its users.

Routine	Implemented	Notes
H5checksum	Quarter 4	New for EFF.

4.1.1 H5checksum

Name: H5checksum

Signature:

```
uint32_t H5checksum( const void *buf, hsize_t length, H5_checksum_seed_t *cseed )
```

Purpose:

Generate a checksum.

Description:

H5checksum generates a checksum for the data in buf with size length bytes.

H5checksum will generate the same checksum for identical data regardless of whether the data is stored in a single contiguous buffer or in multiple non-contiguous buffers. For non-contiguous buffers, H5checksum must be called once for each buffer to “accumulate” the checksum for the complete data.

The checksum seed structure is defined as follows:

```
typedef struct H5_checksum_seed_t {
    uint32_t a;
    uint32_t b;
    uint32_t c;
    int32_t state;
    size_t total_length;
} H5_checksum_seed_t;
```

If the checksum is for a data in a contiguous buffer, call H5checksum with NULL for cseed parameter.

Otherwise, for checksumming a set of non-contiguous regions that will be passed in a single call to H5Dwrite_ff, cseed captures the internal state of the checksum generation algorithm, allowing a single checksum to be generated for the set of non-contiguous data that is identical to the checksum that would be generated if the same data was checksummed in one contiguous block.

When creating a checksum for a set of non-contiguous buffers, H5checksum should be called on each contiguous portion of the buffer with length set to that portion’s corresponding size in bytes. The a, b, c, and state fields in the checksum seed structure should be initialized to 0 for the first call to H5checksum, and total_length should be set to the total size in bytes of all the data to be checksummed (over all the non-contiguous sections of the data to checksum). Each call to H5checksum updates fields in cseed to capture the current internal state of the parameters used to compute the checksum and returns the updated checksum of the entire data that have been passed in so far to the routine with the same cseed parameter. The return values of the intermediate calls to H5checksum for a non-contiguous buffer can be discarded; only the last returned value is used as the checksum for the non-contiguous data. By using the same cseed structure in subsequent calls to H5checksum for each of the non-contiguous buffers, the overall checksum is computed step-by-step.

The generated checksum can be used with H5Pset_dxpl_checksum to attach a checksum to H5Dwrite_ff and H5Mput transfers and to verify the checksum returned by H5Dread_ff and H5Mget in the buffer specified by H5Pset_dxpl_checksum_ptr.

See *Design and Implementation of FastForward Features in HDF5* for information on the interaction between transformation operations, such as datatype conversion, and end-to-end integrity support with checksums.

Parameters:

const void *buf

IN: Buffer with data to be checksummed

Use or disclosure of data contained on this sheet is subject to the restriction on the title page of this document.
Copyright © The HDF Group, 2014. All rights reserved

<i>hsize_t</i> length	IN: Length, in bytes, of the data in buf
<i>H5_checksum_seed_t</i> * cseed	IN/OUT: Checksum state structure used to seed each invocation of H5checksum . Use NULL if data is in a contiguous buffer. If the data is in multiple non-contiguous buffers, use the same cseed structure for each call to H5checksum , which will be called once per buffer.

Returns:

Returns the checksum. In the case of checksumming data stored in non-contiguous buffers, the checksum returned for all calls but the last can be discarded

History:

Added in Quarter 4.

Quarter 5: Added note about interactions between checksums and data transformations.

Man Page Status:

No known issues.

4.2 H5A: Attribute APIs

These routines are used to operate on HDF5 Attribute Objects.

The routines ending in _ff have different signatures than the standard HDF5 library routines.

Man pages for routines whose user interface is unchanged from the standard HDF5 implementation can be found at: http://www.hdfgroup.org/HDF5/doc/RM/RM_H5A.html.

Routine	Implemented	Notes
H5Aclose_ff	Quarter 4	
H5Acreate_ff	Quarter 4	
H5Acreate_by_name_ff	Quarter 4	
H5Adelete_ff	Quarter 4	
H5Adelete_by_name_ff	Quarter 4	
H5Aexists_ff	Quarter 4	
H5Aexists_by_name_ff	Quarter 4	
H5Aopen_ff	Quarter 4	
H5Aopen_by_name_ff	Quarter 4	
H5Aread_ff	Quarter 4	
H5Arename_ff	Quarter 4	
H5Arename_by_name_ff	Quarter 4	
H5Awrite_ff	Quarter 4	
H5Aget_create_plist	Quarter 4	See standard HDF5 man page
H5Aget_name	Quarter 4	See standard HDF5 man page
H5Aget_space	Quarter 4	See standard HDF5 man page
H5Aget_type	Quarter 4	See standard HDF5 man page
H5Aevict_ff		Expect to implement in Q8
H5Aprefetch_ff		Probably will not implement in Prototype
H5Aget_info_ff		Probably will not implement in Prototype
H5Aget_info_by_name_ff		Probably will not implement in Prototype
H5Aiterate_ff		Will not implement
H5Aiterate_by_name_ff		Will not implement
H5Aget_storage_size_ff		Will not implement
H5Adelete_by_idx		Will not implement
H5Aget_info_by_idx		Will not implement

Use or disclosure of data contained on this sheet is subject to the restriction on the title page of this document.
Copyright © The HDF Group, 2014. All rights reserved

H5Aget_name_by_idx		Will not implement
H5Aopen_by_idx		Will not implement

Use or disclosure of data contained on this sheet is subject to the restriction on the title page of this document.
Copyright © The HDF Group, 2014. All rights reserved

4.2.1 H5Aclose_ff

Name: H5Aclose_ff

Signature:

herr_t H5Aclose_ff(hid_t attr_id, hid_t es_id)

Purpose:

Close the specified attribute, possibly asynchronously.

Description:

H5Aclose_ff terminates access to the attribute specified by attr_id by releasing the identifier.

Further use of a released attribute identifier is illegal; a function using such an identifier will fail.

The es_id parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in H5_EVENT_STACK_NULL for the es_id parameter.

Parameters:

<i>hid_t attr_id</i>	IN: Identifier of the attribute to release access to.
<i>hid_t es_id</i>	IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously. Use H5_EVENT_STACK_NULL for synchronous execution.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

When this routine is executed asynchronously, the return value indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

History:

Added in Quarter 4.

Quarter 5: Changed from event queue to event stack.

Man Page Status:

No known issues.

4.2.2 H5Acreate_ff

Name: H5Acreate_ff

Signature:

```
hid_t H5Acreate_ff(hid_t loc_id, const char *attr_name, hid_t type_id, hid_t space_id,
hid_t acpl_id, hid_t aapl_id, hid_t trans_id, hid_t es_id )
```

Purpose:

Create an attribute attached to the specified object, possibly asynchronously.

Description:

H5Acreate_ff creates an attribute, `attr_name`, which is attached to the object specified by the identifier `loc_id`. `loc_id` must be in scope for the transaction identified by `trans_id`.

The attribute name, `attr_name`, must be unique for the object.

The attribute is created with the specified datatype and dataspace, `type_id` and `space_id`, which are created with the H5T and H5S interfaces, respectively.

If `type_id` is either a fixed-length or variable-length string, it is important to set the string length when defining the datatype. String datatypes are derived from H5T_C_S1, which defaults to 1 character in size. See [H5Tset_size](#) and “[Creating variable-length string datatypes](#).”

The attribute creation and access property lists, `acpl_id` and `aapl_id`, are currently unused, but will be used in the future for optional attribute creation and access properties. These property lists should currently be `H5P_DEFAULT`.

`trans_id` indicates the transaction this operation is part of.

The `es_id` parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in `H5_EVENT_STACK_NULL` for the `es_id` parameter.

The attribute identifier returned by this function must be released with `H5Aclose_ff` or resource leaks will develop.

Parameters:

<code>hid_t loc_id</code>	IN: Location identifier May be any HDF5 object identifier (group, dataset, map, or named datatype) or file identifier that is in scope for the transaction. If <code>loc_id</code> is a file identifier, the attribute will be attached to that file's root group.
<code>const char *attr_name</code>	IN: Attribute name
<code>hid_t type_id</code>	IN: Attribute datatype identifier
<code>hid_t space_id</code>	IN: Attribute dataspace identifier
<code>hid_t acpl_id</code>	IN: Attribute creation property list. <i>Currently not used; specify H5P_DEFAULT.</i>
<code>hid_t aapl_id</code>	IN: Attribute access property list <i>Currently not used; specify H5P_DEFAULT.</i>

<i>hid_t trans_id</i>	IN: Transaction identifier specifying the transaction this operation is a part of.
<i>hid_t es_id</i>	IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously. Use H5_EVENT_STACK_NULL for synchronous execution.

Returns:

Returns an attribute identifier if successful; otherwise returns a negative value. When executed asynchronously, a future ID for the new attribute is returned initially. Upon completion of the asynchronous operation, the future ID will be transparently modified to be a “normal” attribute identifier.

When this routine is executed asynchronously, the return value indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

History:

Added in Quarter 4.

Quarter 5: Changed from transaction to transaction id and from event queue to event stack. Added scope requirement and map object.

Man Page Status:

No known issues.

4.2.3 H5Acreate_by_name_ff

Name: H5Acreate_by_name_ff

Signature:

```
hid_t H5Acreate_by_name_ff ( hid_t loc_id, const char *obj_name, const char *attr_name,
    hid_t type_id, hid_t space_id, hid_t acpl_id, hid_t aapl_id, hid_t lapl_id, hid_t trans_id,
    hid_t es_id )
```

Purpose:

Create an attribute attached to the specified object, possibly asynchronously.

Description:

H5Acreate_by_name_ff creates an attribute, `attr_name`, which is attached to the object specified by `loc_id` and `obj_name`.

`loc_id` is a location identifier; `obj_name` is the path to the object relative to `loc_id`. If `loc_id` fully specifies the object to which the attribute is to be attached, `obj_name` should be '.' (a dot). Both `loc_id` and `obj_name` must be in scope for the transaction identified by `trans_id`.

The attribute name, `attr_name`, must be unique for the object.

The attribute is created with the specified datatype and dataspace, `type_id` and `space_id`, which are created with the H5T and H5S interfaces, respectively.

The attribute creation and access property lists, `acpl_id` and `aapl_id`, are currently unused, but will be used in the future for optional attribute creation and access properties. These property lists should currently be `H5P_DEFAULT`.

The link access property list, `lapl_id`, may provide information regarding the properties of links required to access the object, `obj_name`. See “Link Access Properties” in the [H5P APIs](#). The link access property list currently has no effect in the EFF stack and should be set to `H5P_DEFAULT`.

`trans_id` indicates the transaction this operation is part of.

The `es_id` parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in `H5_EVENT_STACK_NULL` for the `es_id` parameter.

The attribute identifier returned by this function must be released with `H5Aclose_ff` or resource leaks will develop.

Parameters:

<code>hid_t loc_id</code>	IN: Location identifier May be any HDF5 object identifier (group, dataset, map, or named datatype) or file identifier that is in scope for the transaction.
<code>const char *obj_name</code>	IN: Object name The object name (path to the object) can be specified relative to <code>loc_id</code> , absolute from the file’s root group, or '.' (a dot), and must be in scope for the transaction.
<code>const char *attr_name</code>	IN: Attribute name
<code>hid_t type_id</code>	IN: Attribute datatype identifier

<i>hid_t space_id</i>	IN: Attribute dataspace identifier
<i>hid_t acpl_id</i>	IN: Attribute creation property list <i>Currently not used; specify H5P_DEFAULT.</i>
<i>hid_t aapl_id</i>	IN: Attribute access property list <i>Currently not used; specify H5P_DEFAULT.</i>
<i>hid_t lapl_id</i>	IN: Link access property list <i>Currently not used in EFF; specify H5P_DEFAULT.</i>
<i>hid_t trans_id</i>	IN: Transaction identifier specifying the transaction this operation is a part of.
<i>hid_t es_id</i>	IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously. Use H5_EVENT_STACK_NULL for synchronous execution

Returns:

Returns an attribute identifier if successful; otherwise returns a negative value. When executed asynchronously, a future ID for the new attribute is returned initially. Upon completion of the asynchronous operation, the future ID will be transparently modified to be a “normal” attribute identifier.

When this routine is executed asynchronously, the return value indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

History:

Added in Quarter 4.

Quarter 5: Changed from transaction to transaction id and from event queue to event stack. Added scope requirement and map object. Noted lapl not used in EFF.

Man Page Status:

No known issues.

4.2.4 H5Adelete_ff

Name: H5Adelete_ff

Signature:

```
herr_t H5Adelete_ff( hid_t loc_id, const char *attr_name, hid_t trans_id, hid_t es_id )
```

Purpose:

Delete an attribute from a specified location, possibly asynchronously.

Description:

H5Adelete_ff removes the attribute specified by its name, `attr_name`, from the object specified by `loc_id`. This function should not be used when attribute identifiers are open on `loc_id` as it may cause the internal indexes of the attributes to change and future writes to the open attributes to produce incorrect results. `loc_id` must be in scope for the transaction identified by `trans_id`.

`trans_id` indicates the transaction this operation is part of.

The `es_id` parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in `H5_EVENT_STACK_NULL` for the `es_id` parameter.

Parameters:

<code>hid_t loc_id</code>	May be any HDF5 object identifier (group, dataset, map, or named datatype) or file identifier that is in scope for the transaction. If <code>loc_id</code> is a file identifier, the attribute will be deleted from that file's root group.
<code>const char *attr_name</code>	IN: Name of the attribute to delete
<code>hid_t trans_id</code>	IN: Transaction identifier specifying the transaction this operation is a part of.
<code>hid_t es_id</code>	IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously. Use <code>H5_EVENT_STACK_NULL</code> for synchronous execution.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

When this routine is executed asynchronously, the return value indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

History:

Added in Quarter 4.

Quarter 5: Changed from transaction to transaction id and from event queue to event stack. Added scope requirement and map object.

Man Page Status:

No known issues.

4.2.5 H5Adelete_by_name_ff

Name: H5Adelete_by_name_ff

Signature:

```
herr_t H5Adelete_by_name_ff( hid_t loc_id, const char *obj_name, const char *attr_name,
hid_t lapl_id, hid_t trans_id, hid_t es_id )
```

Purpose:

Delete an attribute from a specified location, possibly asynchronously.

Description:

H5Adelete_by_name_ff removes the attribute attr_name from an object specified by loc_id and obj_name.

loc_id is a location identifier; obj_name is the path to the object relative to loc_id. If loc_id fully specifies the object from which the attribute is to be deleted, obj_name should be '.' (a dot). Both loc_id and obj_name must be in scope for the transaction identified by trans_id.

The link access property list, lapl_id, may provide information regarding the properties of links required to access the object, obj_name. See “Link Access Properties” in the [H5P APIs](#). The link access property list currently has no effect in the EFF stack and should be set to H5P_DEFAULT.

trans_id indicates the transaction this operation is part of.

The es_id parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in H5_EVENT_STACK_NULL for the es_id parameter.

Parameters:

<i>hid_t loc_id</i>	IN: May be any HDF5 object identifier (group, dataset, map, or named datatype) or file identifier that is in scope for the transaction.
<i>const char *obj_name</i>	IN: Name of object from which attribute is to be removed. The object name (path to the object) can be specified relative to loc_id, absolute from the file's root group, or '.' (a dot), and must be in scope for the transaction.
<i>const char *attr_name</i>	In: Name of attribute to delete
<i>hid_t lapl_id</i>	IN: Link access property list <i>Currently not used in EFF; specify H5P_DEFAULT.</i>
<i>hid_t trans_id</i>	IN: Value used to indicate transaction this operation is a part of.
<i>hid_t es_id</i>	IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously. Use H5_EVENT_STACK_NULL for synchronous execution.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

When this routine is executed asynchronously, the return value indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

Use or disclosure of data contained on this sheet is subject to the restriction on the title page of this document.
Copyright © The HDF Group, 2014. All rights reserved

History:

Added in Quarter 4.

Quarter 5: Changed from transaction to transaction id and from event queue to event stack. Added scope requirement and map object. Noted lapl not used in EFF.

Man Page Status:

No known issues.

4.2.6 H5Aexists_ff

Name: H5Aexists_ff

Signature:

```
herr_t H5Aexists_ff(hid_t loc_id, const char *attr_name, hbool_t *exists,  
hid_t rcntxt_id, hid_t es_id)
```

Purpose:

Determine whether an attribute with a given name exists for an object, possibly asynchronously.

Description:

H5Aexists_ff determines whether the attribute, `attr_name`, exists for the object specified by `loc_id`. `loc_id` must be in scope for the read context identified by `rcntxt_id`.

The results of the existence test are put in `exists`. Note that this value will not be set until after the function completes, which may be later than when the call returns for asynchronous execution. Traditionally the existence results were indicated by the function call's return value, but with the support for asynchronous execution that is no longer possible and the FastForward version included the new `exists` parameter.

`rcntxt_id` indicates the read context for this operation.

The `es_id` parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in `H5_EVENT_STACK_NULL` for the `es_id` parameter.

Parameters:

<code>hid_t loc_id</code>	IN: May be any HDF5 object identifier (group, dataset, map, or named datatype) or file identifier that is in scope for the read context.
<code>const char *attr_name</code>	IN: Attribute name
<code>hbool_t *exists</code>	OUT: Pointer to returned results indicating existence of attribute. When successful, will be a positive value to indicate the attribute exists and 0 (zero) if the attribute does not exist. The value pointed to will not be modified if the existence test failed for some reason.
<code>hid_t rcntxt_id</code>	IN: Read context identifier indicating the read context for this operation.
<code>hid_t es_id</code>	IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously. Use <code>H5_EVENT_STACK_NULL</code> for synchronous execution.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

When this routine is executed asynchronously, the return value indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack. When the asynchronous execution completes successfully, `exists` will contain the existence test results.

History:

Added in Quarter 4.

Quarter 5: Changed from obj_id to loc_id, from transaction to read context id, and from event queue to event stack. Added scope requirement and map object. Changed from htri_t to hbool_t for exists parameter type and updated description of value returned on test failure.

Man Page Status:

No known issues.

4.2.7 H5Aexists_by_name_ff

Name: H5Aexists_by_name_ff

Signature:

```
herr_t H5Aexists_by_name_ff(hid_t loc_id, const char *obj_name, const char *attr_name,  
hid_t lapl_id, hbool_t *exists, hid_t rcntxt_id, hid_t es_id)
```

Purpose:

Determine whether an attribute with a given name exists for an object, possibly asynchronously.

Description:

H5Aexists_ff determines whether the attribute, attr_name, exists for the object specified by loc_id and obj_name.

loc_id is a location identifier; obj_name is the path to the object relative to loc_id. If loc_id fully specifies the object the attribute may be attached to, obj_name should be '.' (a dot). Both loc_id and obj_name must be in scope for the read context identified by rcntxt_id.

The link access property list, lapl_id, may provide information regarding the properties of links required to access the object, obj_name. See “Link Access Properties” in the [H5P APIs](#). The link access property list currently has no effect in the EFF stack and should be set to H5P_DEFAULT.

The results of the existence test are put in exists. Note that this value will not be set until after the function completes, which may be later than when the call returns for asynchronous execution. Traditionally the existence results were indicated by the function call’s return value, but with the support for asynchronous execution that is no longer possible and the FastForward version included the new exists parameter.

rcntxt_id indicates the read context for this operation.

The es_id parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in H5_EVENT_STACK_NULL for the es_id parameter.

Parameters:

hid_t loc_id	IN: Location identifier May be any HDF5 object identifier (group, dataset, map, or named datatype) or file identifier that is in scope for the read context.
const char *obj_name	IN: Object name The object name (path to the object) can be specified relative to loc_id, absolute from the file’s root group, or '.' (a dot), and must be in scope for the read context.
const char *attr_name	IN: Attribute name
hid_t lapl_id	IN: Link property access list <i>Currently not used in EFF; specify H5P_DEFAULT.</i>
hbool_t *exists	OUT: Pointer to returned results indicating existence of attribute. When successful, will be a positive value to indicate the attribute exists and 0 (zero) if the attribute does not exist. The value pointed to will not be modified if the existence test failed for some reason.

<i>hid_t rcntxt_id</i>	IN: Read context identifier indicating the read context for this operation.
<i>hid_t es_id</i>	IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously. Use H5_EVENT_STACK_NULL for synchronous execution.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

When this routine is executed asynchronously, the return value indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack. When the asynchronous execution completes successfully, `exists` will contain the existence test results.

History:

Added in Quarter 4.

Quarter 5: Changed from `obj_id` to `loc_id`, from transaction to read context id, and from event queue to event stack. Added scope requirement and map object. Noted `lapl` currently not used in EFF. Changed from `htri_t` to `hbool_t` for `exists` parameter type and updated description of value returned on test failure.

Man Page Status:

No known issues.

4.2.8 H5Aopen_ff

Name: H5Aopen_ff

Signature:

```
hid_t H5Aopen_ff(hid_t loc_id, const char *attr_name, hid_t aapl_id, hid_t rcntxt_id,  
hid_t es_id)
```

Purpose:

Open an attribute for an object, possibly asynchronously.

Description:

H5Aopen_ff opens an attribute, `attr_name`, which is attached to the object specified by `loc_id`.

`loc_id` is a location identifier and must be in scope for the read context specified by `rcntxt_id`.

The attribute access property list, `aapl_id`, is currently unused and should be H5P_DEFAULT.

`rcntxt_id` indicates the read context for this operation.

The `es_id` parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in H5_EVENT_STACK_NULL for the `es_id` parameter.

The attribute identifier returned by this function must be released with H5Aclose_ff or resource leaks will develop.

Parameters:

<code>hid_t loc_id</code>	IN: May be any HDF5 object identifier (group, dataset, map, or named datatype) or file identifier that is in scope for the read context.
<code>const char *attr_name</code>	IN: Attribute name
<code>hid_t aapl_id</code>	IN: Attribute access property list. Currently not used.
<code>hid_t rcntxt_id</code>	IN: Read context identifier indicating the read context for this operation.
<code>hid_t es_id</code>	IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously. Use H5_EVENT_STACK_NULL for synchronous execution.

Returns:

Returns an attribute identifier if successful; otherwise returns a negative value. When executed asynchronously, a future ID for the new attribute is returned initially. Upon completion of the asynchronous operation, the future ID will be transparently modified to be a “normal” attribute identifier.

When this routine is executed asynchronously, the return value indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

History:

Added in Quarter 4.

Quarter 5: Changed from transaction to read context id and from event queue to event stack. Added scope requirement and map object.

Man Page Status:

No known issues.

Use or disclosure of data contained on this sheet is subject to the restriction on the title page of this document.
Copyright © The HDF Group, 2014. All rights reserved

4.2.9 H5Aopen_by_name_ff

Name: H5Aopen_by_name_ff

Signature:

```
hid_t H5Aopen_by_name_ff(hid_t loc_id, const char *obj_name, const char *attr_name,  
hid_t aapl_id, hid_t lapl_id, hid_t rcntxt_id, hid_t es_id)
```

Purpose:

Open an attribute for an object, possibly asynchronously.

Description:

H5Aopen_by_name_ff opens an attribute, `attr_name`, that is attached to an object specified by `loc_id` and `obj_name`.

`loc_id` is a location identifier; `obj_name` is the path to the object relative to `loc_id`. If `loc_id` fully specifies the object to which the attribute is attached, `obj_name` should be '.' (a dot). Both `loc_id` and `obj_name` must be in scope for the read context identified by `rcntxt_id`.

The attribute access property list, `aapl_id`, is currently unused and should be `H5P_DEFAULT`.

The link access property list, `lapl_id`, may provide information regarding the properties of links required to access the object, `obj_name`. See “Link Access Properties” in the [H5P APIs](#). The link access property list currently has no effect in the EFF stack and should be set to `H5P_DEFAULT`.

`rcntxt_id` indicates the read context for this operation.

The `es_id` parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in `H5_EVENT_STACK_NULL` for the `es_id` parameter.

The attribute identifier returned by this function must be released with `H5Aclose_ff` or resource leaks will develop.

Parameters:

<code>hid_t loc_id</code>	IN: Location identifier May be any HDF5 object identifier (group, dataset, map, or named datatype) or file identifier that is in scope for the read context.
<code>const char *obj_name</code>	IN: Object name Either relative to <code>loc_id</code> , absolute from the file’s root group, or '.' (a dot)
<code>hid_t aapl_id</code>	IN: Attribute access property list. <i>Currently not used; specify H5P_DEFAULT.</i>
<code>hid_t lapl_id</code>	IN: Link access property list. <i>Currently not used in EFF; specify H5P_DEFAULT.</i>
<code>hid_t rcntxt_id</code>	IN: Read context identifier indicating the read context for this operation.
<code>hid_t es_id</code>	IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously. Use <code>H5_EVENT_STACK_NULL</code> for synchronous execution.

Returns:

Use or disclosure of data contained on this sheet is subject to the restriction on the title page of this document.
Copyright © The HDF Group, 2014. All rights reserved

Returns an attribute identifier if successful; otherwise returns a negative value. When executed asynchronously, a future ID for the new attribute is returned initially. Upon completion of the asynchronous operation, the future ID will be transparently modified to be a “normal” attribute identifier.

When this routine is executed asynchronously, the return value indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

History:

Added in Quarter 4.

Quarter 5: Changed from transaction to read context id and from event queue to event stack. Added scope requirement and map object. Noted lapl not used in EFF.

Man Page Status:

No known issues.

4.2.10 H5Aread_ff

Name: H5Aread_ff

Signature:

```
herr_t H5Aread_ff(hid_t attr_id, hid_t mem_type_id, void *buf, hid_t rcntxt_id,  
hid_t es_id )
```

Purpose:

Read an attribute, possibly asynchronously.

Description:

H5Aread_ff reads an attribute, specified by `attr_id`. `attr_id` must be in scope for the read context identified by `rcntxt_id`.

The attribute's memory datatype is specified with `mem_type_id`. The entire attribute is read into `buf` from the file.

Datatype conversion takes place at the time of a read and is automatic. See the [Data Conversion](#) section of *The Data Type Interface (H5T)* in the *HDF5 User's Guide* for a discussion of data conversion, including the range of conversions currently supported by the HDF5 libraries.

`rcntxt_id` indicates the read context for this operation.

The `es_id` parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in `H5_EVENT_STACK_NULL` for the `es_id` parameter.

Parameters:

`hid_t attr_id` IN: Identifier of the attribute to read. Must be in scope for the read context.

`hid_t mem_type_id` IN: identifier of the attribute datatype (in memory)

`void *buf` OUT: Buffer to receive data read from file

`hid_t rcntxt_id` IN: Read context identifier indicating the read context for this operation.

`hid_t eq_id` IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously. Use `H5_EVENT_STACK_NULL` for synchronous execution.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

When this routine is executed asynchronously, the return value indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

History:

Added in Quarter 4.

Quarter 5: Changed from transaction to read context id and from event queue to event stack. Added scope requirement.

Man Page Status:

No known issues.

Use or disclosure of data contained on this sheet is subject to the restriction on the title page of this document.
Copyright © The HDF Group, 2014. All rights reserved

4.2.11 H5Arename_ff

Name: H5Arename_ff

Signature:

```
herr_t H5Arename_ff(hid_t loc_id, const char *old_attr_name, const char *new_attr_name,  
hid_t trans_id, hid_t es_id )
```

Purpose:

Rename an attribute, possibly asynchronously.

Description:

H5Arename_ff changes the name of the attribute attached to the object specified by `loc_id`.

`loc_id` is a location identifier and must be in scope for the transaction identified by `trans_id`.

The old name, `old_attr_name`, is changed to the new name, `new_attr_name`.

`trans_id` indicates the transaction this operation is part of.

The `es_id` parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in `H5_EVENT_STACK_NULL` for the `es_id` parameter.

Parameters:

<code>hid_t loc_id</code>	IN: Location identifier for the object whose attribute is to be renamed. May be any HDF5 object identifier (group, dataset, map, or named datatype) or file identifier that is in scope for the transaction.
<code>const char *old_attr_name</code>	IN: Name of the attribute to be changed.
<code>const char *new_attr_name</code>	IN: New name for the attribute.
<code>hid_t trans_id</code>	IN: Transaction identifier specifying the transaction this operation is a part of.
<code>hid_t es_id</code>	IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously. Use <code>H5_EVENT_STACK_NULL</code> for synchronous execution.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

When this routine is executed asynchronously, the return value indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

History:

Added in Quarter 4.

Quarter 5: Changed from transaction to transaction id and from event queue to event stack. Added scope requirement and map object.

Man Page Status:

No known issues.

4.2.12 H5Arename_by_name_ff

Name: H5Arename_by_name_ff

Signature:

```
herr_t H5Arename_by_name_ff(hid_t loc_id, const char* obj_name,  
const char *old_attr_name, const char *new_attr_name, hid_t lapl_id, hid_t trans_id,  
hid_t es_id )
```

Purpose:

Rename an attribute for an object, possibly asynchronously.

Description:

H5Arename_by_name_ff changes the name of the attribute that is attached to the object specified by `loc_id` and `obj_name`.

`loc_id` is a location identifier; `obj_name` is the path to the object relative to `loc_id`. If `loc_id` fully specifies the object to which the attribute is attached, `obj_name` should be '.' (a dot). Both `loc_id` and `obj_name` must be in scope for the transaction identified by `trans_id`.

The attribute named `old_attr_name` is renamed `new_attr_name`.

The link access property list, `lapl_id`, may provide information regarding the properties of links required to access `obj_name`. See “Link Access Properties” in the [H5P APIs](#). The link access property list currently has no effect in the EFF stack and should be set to `H5P_DEFAULT`.

`trans_id` indicates the transaction this operation is part of.

The `es_id` parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in `H5_EVENT_STACK_NULL` for the `es_id` parameter.

Parameters:

<code>hid_t loc_id</code>	IN: Location identifier for the object whose attribute is to be renamed. May be any HDF5 object identifier (group, dataset, map, or named datatype) or file identifier that is in scope for the transaction.
<code>const char *obj_name</code>	IN: Name of the object whose attribute is to be renamed. The object name (path to the object) can be specified relative to <code>loc_id</code> , absolute from the file’s root group, or '.' (a dot), and must be in scope for the transaction.
<code>const char *old_attr_name</code>	IN: Name of the attribute to be changed.
<code>const char *new_attr_name</code>	IN: New name for the attribute.
<code>hid_t lapl_id</code>	IN: Link access property list identifier. <i>Currently not used in EFF; specify H5P_DEFAULT.</i>
<code>hid_t trans_id</code>	IN: Transaction identifier specifying the transaction this operation is a part of.
<code>hid_t es_id</code>	IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously.

Use H5_EVENT_STACK_NULL for synchronous execution

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

When this routine is executed asynchronously, the return value indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

History:

Added in Quarter 4.

Quarter 5: Changed from transaction to transaction id and from event queue to event stack. Added scope requirement and map object. Noted LAPL not used in EFF.

Man Page Status:

No known issues.

4.2.13 H5Awrite_ff

Name: H5Awrite_ff

Signature:

```
herr_t H5Awrite_ff( hid_t attr_id, hid_t mem_type_id, const void *buf, hid_t trans_id,  
hid_t es_id )
```

Purpose:

Write data to an attribute, possibly asynchronously.

Description:

H5Awrite_ff writes an attribute specified with `attr_id`. `attr_id` must be in scope for the transaction identified by `trans_id`.

The attribute's memory datatype is specified with `mem_type_id`. The entire attribute is written from `buf` into the file.

If `mem_type_id` is either a fixed-length or variable-length string, it is important to set the string length when defining the datatype. String datatypes are derived from `H5T_C_S1`, which defaults to 1 character in size. See [H5Tset_size](#) and “[Creating variable-length string datatypes](#).”

Datatype conversion takes place at the time of a write and is automatic. See the [Data Conversion](#) section of *The Data Type Interface (H5T)* in the *HDF5 User's Guide* for a discussion of data conversion, including the range of conversions currently supported by the HDF5 libraries.

`trans_id` indicates the transaction this operation is part of.

The `es_id` parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in `H5_EVENT_STACK_NULL` for the `es_id` parameter.

Parameters:

<code>hid_t attr_id</code>	IN: Identifier of the attribute to write. Must be in scope for the transaction.
<code>hid_t mem_type_id</code>	IN: identifier of the attribute datatype (in memory)
<code>void *buf</code>	IN: Buffer with data to be written
<code>hid_t trans</code>	IN: Transaction identifier specifying the transaction this operation is a part of.
<code>hid_t es_id</code>	IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously. Use <code>H5_EVENT_STACK_NULL</code> for synchronous execution.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

When this routine is executed asynchronously, the return value indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

History:

Added in Quarter 4.

Quarter 5: Changed from transaction to transaction id and from event queue to event stack. Added scope requirement.

Use or disclosure of data contained on this sheet is subject to the restriction on the title page of this document.
Copyright © The HDF Group, 2014. All rights reserved

Man Page Status:

No known issues.

Use or disclosure of data contained on this sheet is subject to the restriction on the title page of this document.
Copyright © The HDF Group, 2014. All rights reserved

4.3 H5AS: Analysis Shipping APIs

These routines are used to perform an analysis shipping operation on data contained in HDF5 Files (containers) in the Exascale Fast Forward stack.

Analysis shipping operations apply Python operations on data *in situ*, restricted to the results matching an HDF5 query operation. As of the deliverables for Q6, analysis shipping operations can only be performed on HDF5 dataset objects, but future deliverables in Q8 will expand analysis shipping capabilities to allow analysis operations on attribute and group objects as well.

These routines did not exist prior to the Exascale FastForward project.

Routine	Implemented	Notes
H5ASexecute	Quarter 6	Restricted to operations on HDF5 dataset objects currently.

4.3.1 H5ASexecute

Name: H5ASexecute

Signature:

```
herr_t H5ASexecute ( const char *file_name, const char *dataset_name, hid_t query_id, const
                      char *split_script, const char *combine_script, hid_t estack_id)
```

Purpose:

Perform an analysis shipping operation on a HDF5 dataset.

Description:

H5ASexecute ships to a remote server node an analysis operation that executes on data locally stored and selected by the user-defined query query_id. The analysis operation is received on the master storage server node, which first queries the layout of the dataset dataset_name in the container named file_name. The query identified by query_id and the split_script are then farmed to the storage server nodes that locally own a piece of the dataset. The storage server nodes select data that matches the query, run the split script and send the resulting data back to the master server node. The master server node executes combine_script on the data it has gathered from the storage server nodes and sends the result back to the application that initiated the analysis operation.

For example, to compute the average of the data elements that are greater than 53.4 in the dataset “D1”, which is located in the file “test.h5”, one would first create a query:

```
float f = 53.4;
hid_t query_id = H5Qcreate(H5Q_TYPE_DATA_ELEM, H5Q_MATCH_GREATER_THAN,
                           H5T_NATIVE_FLOAT, &f);
```

Then define two python functions split and combine (in this case we assume that the user has passed a stringified version of these functions using the split_script and combine_script parameters):

```
import numpy as np

def split(array):
    return np.array([np.average(array)])


import numpy as np

def combine(arrays):
    global_average = 0
    for a in arrays:
        global_average += a[0]
    global_average /= len(arrays)
    return np.array([global_average])
```

Use or disclosure of data contained on this sheet is subject to the restriction on the title page of this document.
Copyright © The HDF Group, 2014. All rights reserved

And finally call `H5ASexecute` to ship the analysis to the data:

```
H5ASexecute("test.h5", "D1", query_id, split_script, combine_script,  
H5_EVENT_STACK_NULL);
```

Note that the final data output result is not directly returned by this call, one would need to either explicitly request it (assuming that it can be easily transferred) or write it to a separate container on the server. Additional routines to handle returning out from the combine script will be defined in the next milestone.

Parameters:

<code>const char *file_name</code>	IN: Name of the file.
<code>const char *dataset_name</code>	IN: Name of the object/dataset.
<code>hid_t query_id</code>	IN: Identifier of the query that will be applied to create a selection on the dataset elements of <code>dataset_name</code> .
<code>const char *split_script</code>	IN: String defining a python script. The script must define a python function named <code>split</code> , which will be applied locally on the server nodes that contain data matching the query.
<code>const char *combine_script</code>	IN: String defining a <code>combine</code> python script. The script must define a python function named <code>combine</code> , which will be applied on the master server node once the data from the farmed split operations has been gathered.
<code>hid_t estack_id</code>	IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously. Use <code>H5_EVENT_STACK_NULL</code> for synchronous execution.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

When this routine is executed asynchronously, the return value indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

History:

Added in Quarter 6.

Man Page Status:

No known issues.

4.4 H5D: Dataset APIs

These routines are used to operate on HDF5 Datasets Objects.

The routines ending in _ff have different signatures than the standard HDF5 library routines.

Man pages for routines whose user interface is unchanged from the standard HDF5 implementation can be found at: http://www.hdfgroup.org/HDF5/doc/RM/RM_H5D.html.

Routine	Implemented	Notes
H5Dclose_ff	Quarter 4	
H5Dcreate_ff	Quarter 3	
H5Devict_ff	Quarter 7	
H5Dopen_ff	Quarter 3	
H5Dprefetch_ff	Quarter 7	
H5Dread_ff	Quarter 3	
H5Dset_extent_ff	Quarter 4	
H5Dwrite_ff	Quarter 3	
H5Dget_access_plist	Quarter 3	See standard HDF5 man page
H5Dget_create_plist	Quarter 3	See standard HDF5 man page
H5Dget_space	Quarter 3	See standard HDF5 man page
H5Dget_type	Quarter 3	See standard HDF5 man page
H5Dget_space_status		Will not implement
H5Dget_storage_size_ff		Will not implement
H5Dcreate_anon		Will not implement
H5Dget_offset		Doesn't make sense for FF storage; will not implement

4.4.1 H5Dclose_ff

Name: H5Dclose_ff

Signature:

hid_t H5Dclose_ff(hid_t dset_id, hid_t es_id)

Purpose:

Close the specified dataset, possibly asynchronously.

Description:

H5Dclose_ff ends access to a dataset specified by *dset_id* and releases resources used by it. Further use of the dataset identifier is illegal in calls to the dataset API.

The *es_id* parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in H5_EVENT_STACK_NULL for the *es_id* parameter.

Parameters:

hid_t dset_id IN: Identifier of the dataset to close access to.

hid_t es_id IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously. Use H5_EVENT_STACK_NULL for synchronous execution.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

When this routine is executed asynchronously, the return value indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

History:

Added in Quarter 4.

Quarter 5: Changed from event queue to event stack.

Man Page Status:

No known issues.

4.4.2 H5Dcreate_ff

Name: H5Dcreate_ff

Signature:

```
hid_t H5Dcreate_ff ( hid_t loc_id, const char *name, hid_t dtype_id, hid_t space_id,
                      hid_t lcpl_id, hid_t dcpl_id, hid_t dapl_id, hid_t trans_id, hid_t es_id )
```

Purpose:

Create a new dataset and link it into the file, possibly asynchronously.

Description:

H5Dcreate_ff creates a new dataset named name at the location specified by loc_id, and associates constant and initial persistent properties with that dataset, including dtype_id, the datatype of each data element as stored in the file; space_id, the dataspace of the dataset; and other initial properties as defined in the dataset creation property and access property lists, dcpl_id and dapl_id, respectively. Once created, the dataset is opened for access.

loc_id may be a file identifier or a group identifier. name may be either an absolute path in the file or a relative path from loc_id naming the dataset. Both loc_id and name must be in scope for the transaction identified by trans_id.

If dtype_id is either a fixed-length or variable-length string, it is important to set the string length when defining the datatype. String datatypes are derived from H5T_C_S1, which defaults to 1 character in size. See [H5Tset_size](#) and “[Creating variable-length string datatypes](#).”

Currently in the Exascale Fast Forward stack you cannot have nested variable-length datatypes, nor a complex or array datatype with variable-length elements.

In the Exascale Fast Forward stack, only the first dimension in the dataset can be unlimited. The maximum dimensions for the dataspace specified by space_id control this constraint.

The link creation property list, lcpl_id, governs creation of the link(s) by which the new dataset is accessed and the creation of any intermediate groups that may be missing. In the EFF stack, automatic creation of missing intermediate groups (controlled by H5Pset_create_intermediate_group) is not supported. The link access property list currently has no effect in the EFF stack and should be set to H5P_DEFAULT.

trans_id indicates the transaction this operation is part of.

The es_id parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in H5_EVENT_STACK_NULL for the es_id parameter.

The datatype and dataspace properties and the dataset creation and access property lists are attached to the dataset, so the caller may derive new datatypes, dataspaces, and creation and access properties from the old ones, and reuse them in calls to create additional datasets.

Once created, the dataset is ready to receive raw data. Immediately attempting to read raw data from the dataset will return zeros. Note that this behavior differs from traditional HDF5 files, where the fill value will be returned.

To conserve and release resources, the dataset should be closed when access is no longer required.

Parameters:

hid_t loc_id	IN: Location identifier May be any HDF5 group identifier or file identifier that is in scope for the transaction
--------------	---

Use or disclosure of data contained on this sheet is subject to the restriction on the title page of this document.
Copyright © The HDF Group, 2014. All rights reserved

<i>const char *name</i>	IN: Dataset name The dataset name (path to the dataset) can be specified relative to <i>loc_id</i> or absolute from the file's root group, and must be in scope for the transaction.
<i>hid_t dtype_id</i>	IN: Datatype identifier
<i>hid_t space_id</i>	IN: Dataspace identifier
<i>hid_t lcpl_id</i>	IN: Link creation property list <i>Currently not used in EFF; specify H5P_DEFAULT.</i>
<i>hid_t dcpl_id</i>	IN: Dataset creation property list
<i>hid_t dapl_id</i>	IN: Dataset access property list
<i>hid_t trans_id</i>	IN: Transaction identifier specifying the transaction this operation is a part of.
<i>hid_t es_id</i>	IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously. Use <code>H5_EVENT_STACK_NULL</code> for synchronous execution.

Returns:

Returns a dataset identifier if successful; otherwise returns a negative value. When executed asynchronously, a future ID for the new dataset is returned initially. Upon completion of the asynchronous operation, the future ID will be transparently modified to be a “normal” dataset identifier.

When this routine is executed asynchronously, the return value indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

History:

Added in Quarter 3.
 Quarter 4: Updated to reflect switch from parameter *H5_request_t request_ptr* to *hid_t eq_id*.
 Quarter 5: Changed from transaction to transaction id and from event queue to event stack. Added scope requirement. Noted that lapl is currently not used in EFF. Noted current restrictions on support for variable-length types in EFF stack. Noted that only first dimension can be extensible in EFF.

Man Page Status:

No known issues.

4.4.3 H5Devict_ff

Name: H5Devict_ff

Signature:

```
herr_t H5Devict_ff( hid_t dset_id, uint64_t *container_version, hid_t dapl_id,
hid_t es_id )
```

Purpose:

Evict dataset from the burst buffer, possibly asynchronously.

Description:

H5Devict_ff evicts data associated with a dataset from the burst buffer.

The data to be evicted may be resident in the burst buffer under two different scenarios

In the first scenario, the data is resident in the burst buffer as the result of updates to the dataset made via the EFF transaction model. For example, through calls to H5Dcreate_ff or H5Dwrite_ff. These calls add updates to a transaction that are atomically applied to the dataset when the transaction is committed, and we refer to this data as *transaction update data*.

When evicting transaction update data, the container version being evicted should first be persisted to permanent storage (DAOS), with the H5RCpersist command. The dataset's transaction update data for the specified container version, as well as the dataset's transaction update data for all lower-numbered container versions that has not yet been evicted from the burst buffer, will be evicted as the result of this call. If evicting the data would result in container versions with open read contexts becoming inaccessible, the evict will fail.

In the second scenario, the data is resident in the burst buffer as the result of a call to H5Dprefetch_ff. This call replicates data from persistent storage (DAOS) to the burst buffer, and we refer to this data as *replica data*. When replica data is evicted, only data in the burst buffer as a result of the exact replica specified is evicted – transaction update data and other replicas for the dataset remain in the burst buffer.

The `dset_id` parameter specifies the dataset whose data is to be evicted.

The version of the dataset to be evicted is specified by the `container_version` property.

The evict replica property in the access property list, `dapl_id`, is used to specify that a dataset replica is to be evicted. `H5Pset_evict_replica()` sets the evict replica property. If this is set, then replica data will be evicted, otherwise transaction update data will be evicted.

The `es_id` parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in `H5_EVENT_STACK_NULL` for the `es_id` parameter.

Limitations / Future Considerations:

In Quarter 7, only the primary IOD Array object associated with the dataset is evicted. Auxiliary IOD objects will also be evicted in Quarter 8.

When evicting a replica, the `container_version` argument is redundant, as the replica identifier fully specifies the data to be evicted. Consider revisiting and possibly revising the API prior to production release. Possibly have separate evict commands for eviction of transaction update data and of replicas.

For other potential extensions that are beyond the scope of the EFF prototype project, refer to the document *Burst Buffer Space Management – Prototype to Production*.

Parameters:

Use or disclosure of data contained on this sheet is subject to the restriction on the title page of this document.
Copyright © The HDF Group, 2014. All rights reserved

<i>hid_t dset_id</i>	IN: Identifier of the dataset being evicted.
<i>uint64_t container_version</i>	IN: Container version specifying what version of the dataset to evict.
<i>hid_t dapl_id</i>	IN: Identifier of an access property list. If the access property list contains an evict replica property (set via <code>H5Pset_evict_replica()</code>), then the replica_id specified by that property will be evicted.
<i>hid_t es_id</i>	IN: The <code>es_id</code> parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in <code>H5_EVENT_STACK_NULL</code> for the <code>es_id</code> parameter.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

Note that when this routine is executed asynchronously, the return value from the routine only indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

History:

Added in Quarter 7.

Man Page Status:

Will need updates in Quarter 8 when additional features are implemented.

4.4.4 H5Dopen_ff

Name: H5Dopen_ff

Signature:

```
hid_t H5Dopen_ff( hid_t loc_id, const char *name, hid_t dapl_id, hid_t rcntxt_id,  
hid_t es_id)
```

Purpose:

Open an existing dataset, possibly asynchronously.

Description:

H5Dopen_ff opens the existing dataset specified by `loc_id` and `name`.

`loc_id` may be a file identifier or a group identifier. `name` may be either an absolute path in the file or a relative path from `loc_id` naming the dataset. Both `loc_id` and `name` must be in scope for the read context identified by `rcntxt_id`.

The dataset access property list, `dapl_id`, provides information regarding access to the dataset.

`rcntxt_id` indicates the read context for this operation. Note that the returned dataset identifier can be used in other calls that operate on different read contexts, as specified by the read context id passed to those calls.

The `es_id` parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in `H5_EVENT_STACK_NULL` for the `es_id` parameter.

To conserve and release resources, the dataset should be closed with `H5Dclose_ff` when access is no longer required.

Parameters:

`hid_t loc_id` IN: Location identifier
May be any HDF5 group identifier or file identifier that is in scope for the read context.

`const char *name` IN: Dataset name
The dataset name (path to the dataset) can be specified relative to `loc_id` or absolute from the file's root group, and must be in scope for the read context.

`hid_t dapl_id` IN: Dataset access property list

`hid_t rcntxt_id` IN: Read context identifier indicating the read context for this operation.

`hid_t eq_id` IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously. Use `H5_EVENT_STACK_NULL` for synchronous execution.

Returns:

Returns a dataset identifier if successful; otherwise returns a negative value. When executed asynchronously, a future ID for the dataset is returned initially. Upon completion of the asynchronous operation, the future ID will be transparently modified to be a "normal" dataset identifier.

The dataset identifier returned by this call can be used to access the dataset at different container versions, even though a specific read context (and associated container version) is passed in this routine.

When this routine is executed asynchronously, the return value indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

History:

Added in Quarter 3.

Quarter 4: Updated to reflect switch from parameter *H5_request_t request_ptr* to *hid_t eq_id*.

Quarter 5: Changed from transaction to read context id and from event queue to event stack. Added scope requirement.

Quarter 7: Added note that the returned dataset ID can be used to access the dataset at other container versions – not just the one associated with the read context specified in this open call.

Man Page Status:

No known issues.

4.4.5 H5Dprefetch_ff

Name: H5Dprefetch_FF

Signature:

```
herr_t H5Dprefetch_ff( hid_t dset_id, hid_t rcntxt_id, hid_t *replica_id, hid_t dapl_id,  
hid_t es_id )
```

Purpose:

Prefetch all or part of a dataset from persistent storage to burst buffer storage, possibly asynchronously.

Description:

H5Dprefetch_ff prefetches a (partial) dataset, specified by its identifier `dset_id`, from persistent storage (DAOS) into burst buffer storage.

`rcntxt_id` indicates the read context for this operation.

`replica_id`, the replica identifier, is set to indicate where the pre-fetched data can be found in the burst buffer, and is passed to subsequent H5Dread_ff and H5Devict_ff calls.

`dapl_id`, a data access property list identifier, is used to specify partial datasets (hyperslab selection) and control layout of the fetched data on the burst buffers.

The `es_id` parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in `H5_EVENT_STACK_NULL` for the `es_id` parameter.

The `replica_id` is not valid until the operation has completed, if it is executing asynchronously.

Limitations / Future Considerations:

For the EFF prototype project, only the primary IOD Array object associated with the HDF5 dataset will be prefetched; auxiliary IOD objects (including the Blob objects that hold variable-length data) remain on persistent storage (DAOS). For more information, and other potential extensions, refer to the document *Burst Buffer Space Management – Prototype to Production*.

The current API takes a read context identifier rather than simply specifying the container version to be prefetched. Although this seems reasonable in the context of an application that will presumably do a read (on a read context) after a prefetch, in the case of a scheduler that might prefetch data on behalf of the application, a container version would be more appropriate. Revisit and possibly revise the API prior to production release.

In Quarter 7, only full datasets and the default layout are supported.

Parameters:

<code>hid_t dset_id</code>	IN: Identifier of the dataset being prefetched. <code>dset_id</code> must be in scope for the read context.
<code>hid_t rcntxt_id</code>	IN: Read context identifier indicating the read context for this operation.
<code>hid_t * replica_id</code>	IN: Identifier of the replicated data in the burst buffer.
<code>hid_t dapl_id</code>	IN: Identifier of an access property list for this I/O operation.
<code>hid_t es_id</code>	IN: The <code>es_id</code> parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in <code>H5_EVENT_STACK_NULL</code> for the <code>es_id</code> parameter.

Use or disclosure of data contained on this sheet is subject to the restriction on the title page of this document.
Copyright © The HDF Group, 2014. All rights reserved

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

Note that when this routine is executed asynchronously, the return value from the routine only indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

History:

Added in Quarter 7.

Man Page Status:

Will need updates in Quarter 8 when additional features are implemented.

4.4.6 H5Dread_ff

Name: H5Dread_FF

Signature:

```
herr_t H5Dread_ff( hid_t dset_id, hid_t mem_type_id, hid_t mem_space_id, hid_t  
file_space_id, hid_t dxpl_id, void * buf, hid_t rcntxt_id,  
hid_t es_id )
```

Purpose:

Read raw data from a dataset into a buffer, possibly asynchronously.

Description:

H5Dread_ff reads a (partial) dataset, specified by its identifier `dset_id`, from the file into an application memory buffer `buf`. Data transfer properties are defined by the argument `dxpl_id`. The memory datatype of the (partial) dataset is identified by the identifier `mem_type_id`. The part of the dataset to read is defined by `mem_space_id` and `file_space_id`.

`file_space_id` is used to specify only the selection within the file dataset's dataspace. Any dataspace specified in `file_space_id` is ignored by the library and the dataset's dataspace is always used.

`file_space_id` can be the constant `H5S_ALL`, which indicates that the entire file dataspace, as defined by the current dimensions of the dataset, is to be selected.

`mem_space_id` is used to specify both the memory dataspace and the selection within that dataspace.

`mem_space_id` can be the constant `H5S_ALL`, in which case the file dataspace is used for the memory dataspace and the selection defined with `file_space_id` is used for the selection within that dataspace.

If raw data storage space has not been allocated for the dataset, the returned buffer `buf` is filled with zeros. Note that fill values are not supported with the Exascale FastForward storage stack.

The behavior of the library for the various combinations of valid dataspace identifiers and `H5S_ALL` for the `mem_space_id` and the `file_space_id` parameters is described below:

<code>mem_space_id</code>	<code>file_space_id</code>	Behavior
valid dataspace identifier	valid dataspace identifier	<code>mem_space_id</code> specifies the memory dataspace and the selection within it. <code>file_space_id</code> specifies the selection within the file dataset's dataspace.
<code>H5S_ALL</code>	valid dataspace identifier	The file dataset's dataspace is used for the memory dataspace and the selection specified with <code>file_space_id</code> specifies the selection within it. The combination of the file dataset's dataspace and the selection from <code>file_space_id</code> is used for memory also.
valid dataspace identifier	<code>H5S_ALL</code>	<code>mem_space_id</code> specifies the memory dataspace and the selection within it. The selection within the file dataset's dataspace is set to the "all" selection.
<code>H5S_ALL</code>	<code>H5S_ALL</code>	The file dataset's dataspace is used for the memory dataspace and the selection within the memory dataspace is set to the "all" selection. The selection within the file dataset's dataspace is set to the "all" selection.

Use or disclosure of data contained on this sheet is subject to the restriction on the title page of this document.
Copyright © The HDF Group, 2014. All rights reserved

Using `H5S_ALL` for the file dataspace selection indicates that the entire dataspace, as defined by the dimensions of the dataset's dataspace for the specified read context, will be selected. The number of elements selected in the memory dataspace must match the number of elements selected in the file dataspace.

`dpxl_id` can be the constant `H5P_DEFAULT`, in which case the default data transfer properties are used. To read from a replica that was previously brought into the burst buffer by a call to `H5Dprefetch_ff`, use the read replica property in the data transfer property list. The call `H5Pset_read_replica()` sets the read replica property.

`rcntxt_id` indicates the read context for this operation.

The `es_id` parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in `H5_EVENT_STACK_NULL` for the `es_id` parameter.

The selected elements in the buffer are not valid until the operation has completed, if it is executing asynchronously.

Data is automatically converted from the file datatype and dataspace to the memory datatype and dataspace at the time of the read. See the [Data Conversion](#) section of *The Data Type Interface (H5T)* in the *HDF5 User's Guide* for a discussion of data conversion, including the range of conversions currently supported by the HDF5 libraries.

See *Design and Implementation of FastForward Features in HDF5* for information on the interaction between transformation operations, such as datatype conversion, and end-to-end integrity support with checksums.

Parameters:

<code>hid_t dset_id</code>	IN: Identifier of the dataset read from. <code>dset_id</code> must be in scope for the read context.
<code>hid_t mem_type_id</code>	IN: Identifier of the memory datatype.
<code>hid_t mem_space_id</code>	IN: Identifier of the memory dataspace.
<code>hid_t file_space_id</code>	IN: Identifier of the dataset's dataspace in the file.
<code>hid_t dpxl_id</code>	IN: Identifier of a transfer property list for this I/O operation. If specified, the read replica property directs the read to access pre-fetched data.
<code>void * buf</code>	OUT: Buffer to receive data read from file.
<code>hid_t rcntxt_id</code>	IN: Read context identifier indicating the read context for this operation.
<code>hid_t es_id</code>	IN: The <code>es_id</code> parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in <code>H5_EVENT_STACK_NULL</code> for the <code>es_id</code> parameter.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

Note that when this routine is executed asynchronously, the return value from the routine only indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

History:

Use or disclosure of data contained on this sheet is subject to the restriction on the title page of this document.
Copyright © The HDF Group, 2014. All rights reserved

Added in Quarter 3.

Quarter 4: Updated to reflect switch from parameter *H5_request_t request_ptr* to *hid_t eq_id*.

Quarter 5: Changed from transaction to read context id and from event queue to event stack. Added scope requirement.

Quarter 7: Added information about reads from prefetched replicas.

Man Page Status:

No known issues.

4.4.7 H5Dset_extent_ff

Name: H5Dset_extent_ff

Signature:

`hid_t H5Dset_extent_ff(hid_t dset_id, const hsize_t size[], hid_t trans_id, hid_t es_id)`

Purpose:

Change the sizes of a dataset's dimensions.

Description:

H5Dset_extent_ff sets the current dimensions of the dataset `dset_id` to the sizes specified in `size`.

`size` is a 1-dimensional array with n elements, where n is the rank of the dataset's dataspace.

`trans_id` indicates the transaction this operation is part of.

The `es_id` parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in `H5_EVENT_STACK_NULL` for the `es_id` parameter.

Notes:

If the sizes specified in `size` are smaller than the dataset's current dimension sizes, H5Dset_extent_ff will reduce the dataset's dimension sizes to the specified values. In the EFF stack, with support for transactions and asynchronous operations, it is important to remember that the new sizes will not become part of the transaction until the asynchronous operation completes successfully, and will not be registered in the file until the transaction is committed.

It is the user application's responsibility to ensure that valuable data is not lost, as H5Dset_extent_ff does not check that the dataset size is growing.

Parameters:

`hid_t dset_id` IN: Dataset identifier.
`dset_id` must be in scope for the transaction.

`const hsize_t size[]` IN: Array containing the new magnitude of each dimension of the dataset.

`hid_t trans_id` IN: Transaction identifier specifying the transaction this operation is a part of.

`hid_t es_id` IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously.
Use `H5_EVENT_STACK_NULL` for synchronous execution.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

When this routine is executed asynchronously, the return value indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

History:

Added in Quarter 4.

Quarter 5: Changed from transaction to transaction id and from event queue to event stack. Added scope requirement. Removed discussion of space allocation time and fill values as not relevant for EFF.

Man Page Status:

No known issues.

Use or disclosure of data contained on this sheet is subject to the restriction on the title page of this document.
Copyright © The HDF Group, 2014. All rights reserved

4.4.8 H5Dwrite_ff

Name: H5Dwrite_ff

Signature:

```
herr_t H5Dwrite_ff(hid_t dset_id, hid_t mem_type_id, hid_t mem_space_id,  
hid_t file_space_id, hid_t dxpl_id, const void *buf, hid_t trans_id,  
hid_t es_id)
```

Purpose:

Write raw data from a buffer to a dataset, possibly asynchronously.

Description:

H5Dwrite_ff writes a (partial) dataset, specified by its identifier `dset_id`, from the application memory buffer `buf` into the file. Data transfer properties are defined by the argument `dxpl_id`. The memory datatype of the (partial) dataset is identified by the identifier `mem_type_id`. The part of the dataset to write is defined by `mem_space_id` and `file_space_id`.

If `mem_type_id` is either a fixed-length or variable-length string, it is important to set the string length when defining the datatype. String datatypes are derived from `H5T_C_S1`, which defaults to 1 character in size. See [H5Tset_size](#) and “[Creating variable-length string datatypes](#).”

`file_space_id` is used to specify only the selection within the file dataset's dataspace. Any dataspace specified in `file_space_id` is ignored by the library and the dataset's dataspace is always used.

`file_space_id` can be the constant `H5S_ALL`, which indicates that the entire file dataspace, as defined by the current dimensions of the dataset, is to be selected.

`mem_space_id` is used to specify both the memory dataspace and the selection within that dataspace.

`mem_space_id` can be the constant `H5S_ALL`, in which case the file dataspace is used for the memory dataspace and the selection defined with `file_space_id` is used for the selection within that dataspace.

The behavior of the library for the various combinations of valid dataspace IDs and `H5S_ALL` for the `mem_space_id` and the `file_space_id` parameters is described below:

mem_space_id	file_space_id	Behavior
valid dataspace identifier	valid dataspace identifier	<code>mem_space_id</code> specifies the memory dataspace and the selection within it. <code>file_space_id</code> specifies the selection within the file dataset's dataspace.
<code>H5S_ALL</code>	valid dataspace identifier	The file dataset's dataspace is used for the memory dataspace and the selection specified with <code>file_space_id</code> specifies the selection within it. The combination of the file dataset's dataspace and the selection from <code>file_space_id</code> is used for memory also.
valid dataspace identifier	<code>H5S_ALL</code>	<code>mem_space_id</code> specifies the memory dataspace and the selection within it. The selection within the file dataset's dataspace is set to the "all" selection.
<code>H5S_ALL</code>	<code>H5S_ALL</code>	The file dataset's dataspace is used for the memory dataspace and the selection within the memory dataspace is set to the "all" selection. The selection within the file dataset's dataspace is set to

the "all" selection.

Using H5S_ALL for the file dataspace selection indicates that the entire dataspace, as defined by the dimensions of the dataset's dataspace for the read context associated with the transaction, will be selected. The number of elements selected in the memory dataspace must match the number of elements selected in the file dataspace.

`dxpl_id` can be the constant H5P_DEFAULT, in which case the default data transfer properties are used.

`trans_id` indicates the transaction this operation is part of.

The `es_id` parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in H5_EVENT_STACK_NULL for the `es_id` parameter.

Writing to a dataset will fail if the HDF5 file was not opened with write access permissions.

The selected elements in the buffer must not be modified until the operation has completed, if it is executing asynchronously.

Data is automatically converted from the memory datatype and dataspace to the file datatype and dataspace at the time of the write. See the [Data Conversion](#) section of *The Data Type Interface (H5T)* in the *HDF5 User's Guide* for a discussion of data conversion, including the range of conversions currently supported by the HDF5 libraries.

See *Design and Implementation of FastForward Features in HDF5* for information on the interaction between transformation operations, such as datatype conversion, and end-to-end integrity support with checksums.

Parameters:

<code>hid_t dset_id</code>	IN: Identifier of the dataset to write to. <code>dset_id</code> must be in scope for the transaction.
<code>hid_t mem_type_id</code>	IN: Identifier of the memory datatype.
<code>hid_t mem_space_id</code>	IN: Identifier of the memory dataspace.
<code>hid_t file_space_id</code>	IN: Identifier of the dataset's dataspace in the file.
<code>hid_t dxpl_id</code>	IN: Identifier of a transfer property list for this I/O operation.
<code>const void * buf</code>	IN: Buffer with data to be written to the file.
<code>hid_t trans_id</code>	IN: Transaction identifier specifying the transaction this operation is a part of
<code>hid_t es_id</code>	IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously. Use H5_EVENT_STACK_NULL for synchronous execution.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

When this routine is executed asynchronously, the return value indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

History:

Added in Quarter 3.

Quarter 4: Updated to reflect switch from parameter `H5_request_t request_ptr` to `hid_t eq_id`.

Quarter 5: Changed from transaction to transaction id and from event queue to event stack. Added scope requirement. Noted the interaction between transformations and checksums. Removed space allocation, fill values, and compact representation text. Noted that “all” is relative to the read context associated with the transaction.

Man Page Status:

No known issues.

4.5 H5ES: Event Stack APIs

These routines allow the application to create and work with event stacks. An event stack provides an organizing structure for managing and monitoring functions that have been called asynchronously.

Once an event stack is created, its identifier can be passed to other HDF5 APIs that will be run asynchronously. The event associated with an asynchronous operation will be pushed onto the event stack whose identifier was passed as a parameter to the function. The application can use the H5ES APIs to monitor the completion status of an individual event or all events in an event stack. One or more of the events in an event stack can be cancelled. Monitoring or cancelling an event is equivalent to monitoring or cancelling the asynchronous operation the event is associated with.

Event stacks are per-process and there may be multiple event stacks for any given process. The most recent event added to an event stack will always be at index 0 (useful for calls to H5EScancel, H5ESTest, H5ESwait).

The order in which asynchronous functions execute and complete is independent of their associated event objects' location in any event stack.

These routines did not exist prior to the Exascale FastForward project.

Routine	Implemented	Notes
H5EScancel		Deferred indefinitely
H5EScancel_all		Deferred indefinitely
H5ESclear	Quarter 6	
H5ESclose	Quarter 6	
H5EScreate	Quarter 6	
H5ESget_count	Quarter 6	
H5ESget_event_info		Deferred indefinitely
H5ESTest	Quarter 6	
H5ESTest_all	Quarter 6	
H5ESwait	Quarter 6	
H5ESwait_all	Quarter 6	

4.5.1 H5EScancel

Name: H5EScancel

Signature:

```
herr_t H5EScancel(hid_t es_id, size_t event_idx, H5ES_status_t *status )
```

Purpose:

Cancel a particular event in an event stack.

Description:

`H5EScancel` cancels the in-progress event specified by `event_idx` in the event stack specified by `es_id` and indicates the event's status.

`event_idx` is the index of the event to be cancelled. The most recent event pushed on to the event stack is at index 0.

`status` is set to indicate whether the event was cancelled or had already completed. Possible values for `status` and their meanings are:

`H5ES_STATUS_CANCEL` Event was cancelled, either by this call or previously.

`H5ES_STATUS_SUCCEED` Event completed successfully.

`H5ES_STATUS_FAIL` Event completed unsuccessfully.

Cancelled (and completed) events remain on the event stack until the stack is cleared.

Parameters:

`hid_t es_id` IN: Identifier of the event stack.

`size_t event_idx` IN: Index of the event to be cancelled.

`H5ES_status_t *status` OUT: Status of the event upon completion of the call.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

History:

Quarter 5: Designed and documented but not yet implemented.

Man Page Status:

No known issues.

4.5.2 H5EScancel_all

Name: H5EScancel_all

Signature:

*hid_t H5EScancel_all(hid_t es_id, H5ES_status_t *status)*

Purpose:

Cancel all events in an event stack that have not yet completed.

Description:

H5EScancel_all cancels the in-progress events in the event stack specified by `es_id` and indicates a status.

`status` is set to indicate the overall status of the events in the event stack. Possible values for `status` and their meanings are:

`H5ES_STATUS_CANCEL` All events were cancelled, either by this call or previously.

`H5ES_STATUS_SUCCEED` Some or all events completed, and all completed events succeeded.

`H5ES_STATUS_FAIL` Some or all events completed, and one or more completed event failed.

Cancelled (and completed) events remain on the event stack until the stack is cleared.

Parameters:

`hid_t es_id` IN: Identifier of the event stack.

`H5ES_status_t *status` OUT: Overall status of the events in the event stack upon completion of the call.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

History:

Quarter 5: Designed and documented but not yet implemented.

Man Page Status:

No known issues.

4.5.3 H5ESclear

Name: H5ESclear

Signature:

herr_t H5ESclear(hid_t es_id)

Purpose:

Clear all events from an event stack.

Description:

H5ESclear clears all event objects from the event stack specified by `es_id`, after first confirming that no in-progress events remain in the stack. If the stack does contain one or more in-progress events, no events are cleared and the call fails (a negative value is returned).

If there are in-progress events on the stack, `H5ESwait_all` can be used to wait for their completion before calling `H5ESclear`. Alternatively, `H5EScancel_all` can be used to cancel the in-progress events before calling `H5ESclear`.

Parameters:

`hid_t es_id` IN: Identifier of the event stack to be cleared.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value. If the event contains one or more in-progress events a negative value is returned.

History:

Quarter 5: Designed and documented but not yet implemented.

Quarter 6: Implemented.

Man Page Status:

No known issues.

4.5.4 H5ESclose

Name: H5ESclose

Signature:

herr_t H5ESclose(hid_t es_id)

Purpose:

Close an event stack.

Description:

`H5ESclose` closes the event stack specified by `es_id`, after first confirming that no in-progress events remain in the stack. If the stack does contain one or more in-progress events, the event stack is not closed and the call fails (a negative value is returned).

If there are in-progress events on the stack, `H5ESwait_all` can be used to wait for their completion before calling `H5ESclose`. Alternatively, `H5EScancel_all` can be used to cancel the in-progress events before calling `H5ESclose`.

Upon successful completion, the event stack identifier, `es_id`, is no longer valid.

Parameters:

hid_t es_id IN: Identifier of the event stack to be closed.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

History:

Quarter 5: Designed and documented but not yet implemented.
Quarter 6: Implemented.

Man Page Status:

No known issues.

4.5.5 H5EScreate

Name: H5EScreate

Signature:

hid_t H5EScreate(void)

Purpose:

Create an event stack.

Description:

H5EScreate creates an event stack to monitor and manage the events associated with asynchronous function execution. Multiple event stacks can be created and used concurrently in a program. Events added to the event stack are bundled and tracked together, although they may also be monitored individually.

The identifier for an event stack is used in all asynchronous function calls, which push the event associated with the function's execution onto the event stack. The order in which asynchronous functions execute and complete is independent of their associated event's location in any event stack.

Parameters:

none

Returns:

Returns an event stack identifier if successful; otherwise returns a negative value.

History:

Quarter 5: Designed and documented but not yet implemented.

Quarter 6: Implemented.

Man Page Status:

No known issues.

4.5.6 H5ESget_count

Name: H5ESget_count

Signature:

*hid_t H5ESget_count(hid_t es_id, size_t *count)*

Purpose:

Query the number of events in an event stack.

Description:

H5ESget_count retrieves the number of events in the event stack specified by `es_id`.

Parameters:

`hid_t es_id` IN: Identifier of the event stack.

`size_t *count` OUT: Number of events.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

History:

Quarter 5: Designed and documented but not yet implemented.

Quarter 6: Implemented.

Man Page Status:

No known issues.

4.5.7 H5ESget_event_info

Name: H5ESget_event_info

Signature:

```
herr_t H5ESget_event_info(hid_t es_id, size_t start_idx, size_t count,
const char *ev_trace_str_arr[], H5ES_status_t ev_status_arr[],
H5E_stack_id ev_err_stack_id_arr[])
```

Purpose:

Retrieve information about events in an event stack.

Description:

H5ESget_event_info retrieves a variety of information about events in the event stack specified by `es_id`.

`start_idx` gives the index of the first event to retrieve information about. The most recent event pushed on the event stack is at index 0.

`count` gives the number of events to retrieve information about.

The array parameters `ev_trace_str_arr`, `ev_status_arr`, and `ev_err_stack_id_arr` must each be large enough to hold at least `count` entries. Passing a NULL pointer for any of the array parameters indicates the information normally returned in that array will not be retrieved.

`ev_trace_str_arr` is an array that holds tracing strings that document the API call and parameters for each event.

`ev_status_arr` is an array that holds the completion status (`H5ES_STATUS_SUCCEED`, `H5ES_STATUS_FAIL`, `H5ES_STATUS_IN_PROGRESS`, or `H5ES_STATUS_CANCEL`) for each event.

`ev_err_stack_id_arr` is an array that holds HDF5 error stack IDs for the operations associated with each event. Events that are in progress, have been cancelled, or have completed successfully will have a -1 value for the error stack ID. Error stacks retrieved from this routine may be queried with the `H5E` API calls.

This call may be made before or after all events in the event stack have completed.

Parameters:

<code>hid_t es_id</code>	IN: Identifier of the event stack.
<code>size_t *start_idx</code>	IN: Index of the first event to retrieve information about.
<code>size_t count</code>	IN: Number of events to retrieve information about.
<code>const char *ev_trace_str_arr[]</code>	OUT: Array containing tracing strings documenting API call and parameters of events
<code>H5ES_status_t ev_status_arr[]</code>	OUT: Array containing status of events
<code>H5E_stack_id ev_err_stack_id_arr[]</code>	OUT: Array containing error stack IDs of events.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

History:

Use or disclosure of data contained on this sheet is subject to the restriction on the title page of this document.
Copyright © The HDF Group, 2014. All rights reserved

Quarter 5: Designed and documented but not yet implemented.

Man Page Status:

No known issues.

4.5.8 H5EStest

Name: H5EStest

Signature:

*herr_t H5EStest(hid_t es_id, size_t event_idx, H5ES_status_t *status)*

Purpose:

Test to determine the status of a particular event in an event stack.

Description:

`H5EStest` reports the completion status of the event specified by `event_idx` in the event stack specified by `es_id`. The `H5EStest` call returns immediately.

`event_idx` is the index of the event to be tested. The most recent event pushed onto the event stack is always at index 0.

`status` is set to indicate the status of the event. Possible values for `status` and their meanings are:

`H5ES_STATUS_IN_PROGRESS` Event is still active.

`H5ES_STATUS_SUCCEED` Event completed successfully.

`H5ES_STATUS_FAIL` Event completed unsuccessfully.

`H5ES_STATUS_CANCEL` Event was cancelled.

Parameters:

`hid_t es_id` IN: Identifier of the event stack.

`size_t event_idx` IN: Index of the event to be tested.

`H5ES_status_t *status` OUT: Status of the event.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

History:

Quarter 5: Designed and documented but not yet implemented.

Quarter 6: Implemented.

Man Page Status:

No known issues.

4.5.9 H5ESTest_all

Name: H5ESTest_all

Signature:

*herr_t H5ESTest_all(hid_t es_id, H5ES_status_t *status)*

Purpose:

Test to determine the status of all events in an event stack.

Description:

H5ESTest_all reports an overall completion status for all events in the event stack specified by `es_id`. The H5ESTest_all call returns immediately.

`status` is set to indicate the overall status of the events in the event stack. Possible values for `status` and their meanings are:

`H5ES_STATUS_IN_PROGRESS` At least one event is still active.

`H5ES_STATUS_SUCCEED` Some or all events completed, and all completed events succeeded.

`H5ES_STATUS_FAIL` Some or all events completed, and one or more completed event failed.

`H5ES_STATUS_CANCEL` All events were cancelled.

Parameters:

`hid_t es_id` IN: Identifier of the event stack.

`H5ES_status_t *status` OUT: Overall status of the events in the event stack.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

History:

Quarter 5: Designed and documented but not yet implemented.

Quarter 6: Implemented.

Man Page Status:

No known issues.

4.5.10 H5ESwait

Name: H5ESwait

Signature:

*herr_t H5ESwait(hid_t es_id, size_t event_idx, H5ES_status_t *status)*

Purpose:

Wait for a particular event in an event stack to complete or be cancelled.

Description:

`H5ESwait` waits for the completion or cancellation of the event specified by `event_idx` in the event stack specified by `es_id` and reports the event's completion status. The `H5ESwait` call does not return until the event being waited on completes or is cancelled.

`event_idx` is the index of the event to be waited on. The most recent event pushed on to the event stack is at index 0.

`status` is set to indicate the completion status of the event. Possible values for `status` and their meanings are:

`H5ES_STATUS_SUCCEED` Event completed successfully.

`H5ES_STATUS_FAIL` Event completed unsuccessfully.

`H5ES_STATUS_CANCEL` Event was cancelled.

Parameters:

`hid_t es_id` IN: Identifier of the event stack.

`size_t event_idx` IN: Index of the event to be tested.

`H5ES_status_t *status` OUT: Status of the event.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

History:

Quarter 5: Designed and documented but not yet implemented.

Quarter 6: Implemented.

Man Page Status:

No known issues.

4.5.11 H5ESwait_all

Name: H5ESwait_all

Signature:

*herr_t H5ESwait_all(hid_t es_id, H5ES_status_t *status)*

Purpose:

Wait for all events in an event stack to complete or be cancelled.

Description:

H5ESwait_all waits for the completion or cancellation of all events in the event stack specified by *es_id* and reports an overall completion status. The H5ESwait_all call does not return until all events in the event stack complete or are cancelled.

status is set to indicate the overall completion status of the events in the event stack. Possible values for *status* and their meanings are:

<i>H5ES_STATUS_SUCCEED</i>	Some or all events completed, and all completed events succeeded.
----------------------------	---

<i>H5ES_STATUS_FAIL</i>	Some or all events completed, and one or more completed event failed.
-------------------------	---

<i>H5ES_STATUS_CANCEL</i>	All events were cancelled.
---------------------------	----------------------------

Parameters:

hid_t es_id IN: Identifier of the event stack.

*H5ES_status_t *status* OUT: Overall completion status of the events in the event stack.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

History:

Quarter 5: Designed and documented but not yet implemented.

Quarter 6: Implemented

Man Page Status:

No known issues.

4.6 H5F: File (Container) APIs

These routines are used to operate on HDF5 File (Container) Objects.

The routines ending in _ff have different signatures than the standard HDF5 library routines.

Man pages for routines whose user interface is unchanged from the standard HDF5 implementation can be found at: http://www.hdfgroup.org/HDF5/doc/RM/RM_H5F.html.

Routine	Implemented	Notes
H5Fclose_ff	Quarter 4	
H5Fcreate_ff	Quarter 3	
H5Fopen_ff	Quarter 3	
H5Fget_access_plist	Quarter 3	Local operation, no need to do asynchronously. See standard HDF5 man page
H5Fget_create_plist	Quarter 3	Local operation, no need to do asynchronously. See standard HDF5 man page
H5Fget_intent	Quarter 3	Local operation, no need to do asynchronously. See standard HDF5 man page
H5Fget_name	Quarter 3	Local operation, no need to do asynchronously. See standard HDF5 man page
H5Fget_filesize		Probably will not implement in Prototype
H5Fget_info		Probably will not implement in Prototype
H5Flush		Doesn't make sense for FF storage; will not implement (<i>decided in Q5</i>)
H5Fget_atomicity		Doesn't make sense for FF storage; will not implement
H5Fget_file_image		Doesn't make sense for FF storage; will not implement
H5Fget_free_space		Doesn't make sense for FF storage; will not implement
H5Fget_mdc_config		Doesn't make sense for FF storage; will not implement
H5Fget_mdc_hit_rate		Doesn't make sense for FF storage; will not implement
H5Fget_mdc_size		Doesn't make sense for FF storage; will not implement

Use or disclosure of data contained on this sheet is subject to the restriction on the title page of this document.
Copyright © The HDF Group, 2014. All rights reserved

H5Fget_vfd_handle		Doesn't make sense for FF storage; will not implement
H5Fis_hdf5		Will not implement unless needed
H5Fmount		Doesn't make sense for FF storage; will not implement
H5Freopen		Doesn't make sense for FF storage; will not implement
H5Freset_mdc_hit_rate_stats		Doesn't make sense for FF storage; will not implement
H5Fset_mpi_atomicity		Doesn't make sense for FF storage; will not implement
H5Fset_mdc_config		Doesn't make sense for FF storage; will not implement
H5Funmount		Doesn't make sense for FF storage; will not implement
H5Fget_obj_count		Will not implement unless needed
H5Fget_obj_ids		Will not implement unless needed

Use or disclosure of data contained on this sheet is subject to the restriction on the title page of this document.
 Copyright © The HDF Group, 2014. All rights reserved

4.6.1 H5Fclose_ff

Name: H5Fclose_ff

Signature:

hid_t H5Fclose_ff(hid_t file_id, hbool_t persist_flag, hid_t es_id)

Purpose:

Terminate access to an HDF5 file (container), possibly asynchronously.

Description:

H5Fclose_ff terminates access to an HDF5 file on the Exascale FastForward storage system by terminating access to the container through `file_id`.

Applications are responsible for finishing transactions (H5TRfinish) before H5Fclose_ff is called. Applications should close all open objects in a container before calling H5Fclose_ff.

H5Fclose_ff will consume a transaction ID (the highest writable transaction) to write additional file metadata before closing the container. By default, H5Fclose_ff will call persist (H5RCpersist) on that last transaction. The user can opt not persist the final transaction created by H5Fclose_ff by setting `persist_flag` to FALSE (0). The container will still be readable and correct, but there will be some performance loss if the additional file metadata is not persisted.

The `es_id` parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in `H5_EVENT_STACK_NULL` for the `es_id` parameter.

Parameters:

hid_t file_id IN: Identifier of container to terminate access to.

hid_t es_id IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously. Use `H5_EVENT_STACK_NULL` for synchronous execution.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

When this routine is executed asynchronously, the return value indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

History:

Added in Quarter 4.

Quarter 5: Changed from event queue to event stack. Added text clarifying that call doesn't flush data to storage.

Quarter 7: Added persist flag to indicate whether the transaction consumed by H5Fclose should be persisted or not. Default is to persist.

Man Page Status:

No known issues.

4.6.2 H5Fcreate_ff

Name: H5Fcreate_ff

Signature:

```
hid_t H5Fcreate_ff( const char *name, unsigned flags, hid_t fcpl_id, hid_t fapl_id,  
hid_t es_id )
```

Purpose:

Create an HDF5 file (container), possibly asynchronously.

Description:

H5Fcreate_ff is the primary function for creating HDF5 files on the Exascale FastForward storage system. It creates a new HDF5 file (IOD container / DAOS container) with the specified name and property lists, and indicates whether an existing container of same name should be overwritten.

The `name` parameter specifies the name of the new container.

The `flags` parameter indicates whether an existing container is to be overwritten. It should be set to either `H5F_ACC_TRUNC` to overwrite an existing container or `H5F_ACC_EXCL`, instructing the function to fail if the container already exists.

New containers are always created in read-write mode, so the read-write and read-only flags, `H5F_ACC_RDWR` and `H5F_ACC_RDONLY`, respectively, are not relevant in this function. Further note that a specification of `H5F_ACC_RDONLY` will be ignored; the file will be created in read-write mode, regardless.

More complex behaviors of file creation and access are controlled through the file creation and file access property lists, `fcpl_id` and `fapl_id`, respectively. The value of `H5P_DEFAULT` for any property list value indicates that the library should use the default values for that appropriate property list.

Note: The `H5Pset_vol_iod()` routine must be invoked on the `fapl_id`, or the IOD VOL plugin will not be used.

The `es_id` parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in `H5_EVENT_STACK_NULL` for the `es_id` parameter.

The return value is a file identifier for the newly-created container; this file identifier should be closed by calling `H5Fclose_ff` when it is no longer needed.

Calling `H5Fcreate_ff` with an already opened file will fail with an error.

Parameters:

`const char *name` IN: Name of the file to create.

`uintn flags` IN: File access flags. Allowable values are:
`H5F_ACC_TRUNC`
Truncate file, if it already exists, erasing all data previously stored in the file.
`H5F_ACC_EXCL`
Fail if file already exists.

- `H5F_ACC_TRUNC` and `H5F_ACC_EXCL` are mutually exclusive; use exactly one.

hid_t fcpl_id IN: File creation property list identifier, used when modifying default file metadata. Use `H5P_DEFAULT` to specify default file creation properties.

hid_t es_id IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously. Use `H5_EVENT_STACK_NULL` for synchronous execution.

Returns:

Returns a file identifier if successful; otherwise returns a negative value. When executed asynchronously, a future ID for the new file is returned initially. Upon completion of the asynchronous operation, the future ID will be transparently modified to be a “normal” file identifier.

When this routine is executed asynchronously, the return value indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

History:

Added in Quarter 3.

Quarter 4: Updated to reflect switch from parameter `H5_request_t request_ptr` to `hid_t eq_id`.

Quarter 5: Changed from event queue to event stack.

Man Page Status:

No known issues.

4.6.3 H5Fopen_ff

Name: H5Fopen_ff

Signature:

```
hid_t H5Fopen_ff( const char *name, unsigned flags, hid_t fapl_id, hid_t rcntxt_id,  
hid_t es_id )
```

Purpose:

Open an existing HDF5 file (container), possibly asynchronously.

Description:

H5Fopen_ff is the primary function for accessing existing HDF5 files (IOD containers / DAOS containers). This function opens the named file in the specified access mode and with the specified access property list.

Note that H5Fopen does not create a container if it does not already exist; see H5Fcreate_ff.

The name parameter specifies the name of the container to be opened.

The flags parameter specifies whether the container will be opened in read-write or read-only mode, H5F_ACC_RDWR or H5F_ACC_RDONLY, respectively. More complex behaviors of file access are controlled through the file-access property list.

The fapl_id parameter specifies the file access property list. Use of H5P_DEFAULT specifies that default I/O access properties will be used.

Note: the H5Pset_vol_iod() routine must be invoked on the fapl_id, or the IOD VOL plugin will not be used.

The rcntxt_id parameter allows the user to pass a pointer to an hid_t datatype signaling that the function should acquire a read handle for the last (highest) container version for the successfully opened container. In addition to acquiring the read handle, a read context is created and the identifier for that context is returned in *rcntxt_id. The read context identifier will not be available until after the H5Fopen_ff call completes successfully. H5RCrelease must be called to release the read handle and close the context when it is no longer needed.

If the rcntxt_id parameter is set to NULL, no read handle is acquired and no read context is created.

The es_id parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in H5_EVENT_STACK_NULL for the es_id parameter

The return value is a file identifier for the newly-opened container; this file identifier should be closed by calling H5Fclose_ff when it is no longer needed.

Multiple opens:

A container can be opened with a new H5Fopen_ff call without closing an already-open identifier established in a previous H5Fopen_ff or H5Fcreate_ff call. Each H5Fopen_ff call will return a unique identifier and the container can be accessed through any of these identifiers as long as the identifier remains valid. A container can only have one open identifier with write privileges at any given time.

Parameters:

const char *name IN: Name of the file to be opened.

unsigned flags IN: File access flags. Allowable values are:
H5F_ACC_RDWR Allow read and write access to file.

Use or disclosure of data contained on this sheet is subject to the restriction on the title page of this document.
Copyright © The HDF Group, 2014. All rights reserved

`H5F_ACC_RDONLY` Allow read-only access to file.
`H5F_ACC_RDWR` and `H5F_ACC_RDONLY` are mutually exclusive; use exactly one.

`hid_t fapl_id` IN: Identifier for the file access properties list. If parallel file access is desired, this is a collective call according to the communicator stored in the `fapl_id`. Use `H5P_DEFAULT` for default file access properties.

`hid_t *rcntxt_id` IN/OUT: Pointer to read context for the last (highest) container version for the successfully opened container. Pass in `NULL` if no read context is desired.

`hid_t es_id` IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously. Use `H5_EVENT_STACK_NULL` for synchronous execution.

Returns:

Returns a file identifier if successful; otherwise returns a negative value. When executed asynchronously, a future ID for the file is returned initially; upon completion of the asynchronous operation, the future ID will be transparently modified to be a “normal” file identifier.

Note that when this routine is executed asynchronously, the return value from the routine only indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event queue.

History:

Added in Quarter 3.
Quarter 4: Updated to reflect switch from parameter `H5_request_t request_ptr` to `hid_t eq_id`.
Quarter 5: Changed from event queue to event stack. Added `rcntxt_id` parameter.

Man Page Status:

No known issues.

4.7 H5G: Group APIs

These routines are used to operate on HDF5 Group Objects.

The routines ending in _ff have different signatures than the standard HDF5 library routines.

Man pages for routines whose user interface is unchanged from the standard HDF5 implementation can be found at: http://www.hdfgroup.org/HDF5/doc/RM/RM_H5G.html.

Routine	Implemented	Notes
H5Gclose_ff	Quarter 4	
H5Gcreate_ff	Quarter 3	
H5Gevict_ff	Quarter 7	
H5Gopen_ff	Quarter 3	
H5Gprefetch_ff	Quarter 7	
H5Gget_create_plist	Quarter 3	Local operation, no need to do asynchronously. See standard HDF5 man page.
H5Gget_info_ff		Will not implement
H5Gget_info_by_name_ff		Will not implement
H5Gcreate_anon		Will not implement
H5Gget_info_by_idx		Will not implement

4.7.1 H5Gclose_ff

Name: H5Gclose_ff

Signature:

hid_t H5Gclose_ff(hid_t group_id, hid_t es_id)

Purpose:

Close the specified group, possibly asynchronously.

Description:

H5Gclose_ff releases resources used by a group that was opened by H5Gcreate_ff or H5Gopen_ff. After closing a group, the *group_id* cannot be used again.

Failure to release a group with this call will result in resource leaks.

The *es_id* parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in H5_EVENT_STACK_NULL for the *es_id* parameter.

Parameters:

hid_t group_id IN: Group identifier to release.

hid_t es_id IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously. Use H5_EVENT_STACK_NULL for synchronous execution.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

When this routine is executed asynchronously, the return value indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

History:

Added in Quarter 4.

Quarter 5: Changed from event queue to event stack

Man Page Status:

No known issues.

4.7.2 H5Gcreate_ff

Name: H5Gcreate_ff

Signature:

```
hid_t H5Gcreate_ff( hid_t loc_id, const char *name, hid_t lcpl_id, hid_t gcpl_id,
hid_t gapl_id, hid_t trans_id, hid_t es_id )
```

Purpose:

Create a new group and links it into the file, possibly asynchronously.

Description:

H5Gcreate_ff creates a new group in a file. After a group has been created, links to datasets and to other groups can be added.

The loc_id and name parameters specify where the group is located. loc_id may be a file identifier or a group identifier. name is the link to the group; name may be either an absolute path in the file (the links from the root group to the new group) or a relative path from loc_id (the link(s) from the group specified by loc_id to the new group). See the “[Accessing objects by location and name](#)” topic for more information. Both loc_id and name must be in scope for the transaction identified by trans_id.

lcpl_id, gcpl_id, and gapl_id are property list identifiers. These property lists govern how the link to the group is created, how the group is created, and how the group can be accessed in the future, respectively.

H5P_DEFAULT can be passed in if the default properties are appropriate for these property lists. Currently, there are no APIs for the group access property list; use H5P_DEFAULT. See “[H5P: Property List Interface](#)” for the functions that can be used with each property list. H5P_DEFAULT must be used for all three of these parameters in the current EFF implementation.

trans_id indicates the transaction this operation is part of.

The es_id parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in H5_EVENT_STACK_NULL for the es_id parameter.

The return value is the group identifier for the new group. Call H5Gclose with the group identifier to close the group and release resources when access is no longer required.

Parameters:

<i>hid_t loc_id</i>	IN: File or group identifier that is in scope for the transaction
<i>const char *name</i>	IN: Absolute or relative name of the link to the new group that is in scope for the transaction
<i>hid_t lcpl_id</i>	IN: Link creation property list identifier Must be H5P_DEFAULT in the current EFF implementation.
<i>hid_t gcpl_id</i>	IN: Group creation property list identifier Must be H5P_DEFAULT in the current EFF implementation.
<i>hid_t gapl_id</i>	IN: Group access property list identifier (No group access properties have been implemented at this time; use H5P_DEFAULT.)
<i>uint64_t trans</i>	IN: Value used to indicate transaction this operation is a part of.

hid_t es_id IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously. Use H5_EVENT_STACK_NULL for synchronous execution.

Returns:

Returns a group identifier if successful; otherwise returns a negative value. When executed asynchronously, a future ID for the new group is returned initially. Upon completion of the asynchronous operation, the future ID will be transparently modified to be a “normal” group identifier.

When this routine is executed asynchronously, the return value indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

History:

Added in Quarter 3.

Quarter 4: Updated to reflect switch from parameter *H5_request_t request_ptr* to *hid_t eq_id*.

Quarter 5: Changed from transaction to transaction id and from event queue to event stack. Added scope requirement.

Man Page Status:

No known issues.

4.7.3 H5Gevict_ff

Name: H5Gevict_ff

Signature:

```
herr_t H5Gevict_ff( hid_t grp_id, uint64_t *container_version, hid_t gapl_id,
hid_t es_id )
```

Purpose:

Evict group from the burst buffer, possibly asynchronously.

Description:

H5Gevict_ff evicts data associated with a group from the burst buffer.

The data to be evicted may be resident in the burst buffer under two different scenarios

In the first scenario, the data is resident in the burst buffer as the result of updates to the group made via the EFF transaction model. For example, through calls to H5Gcreate_ff or H5Dcreate_ff, where the new group is created or a new dataset is created in the group. These calls add updates to a transaction that are atomically applied to the group when the transaction is committed, and we refer to this data as *transaction update data*.

When evicting transaction update data, the container version being evicted should first be persisted to permanent storage (DAOS), with the H5RCpersist command. The group's transaction update data for the specified container version, as well as the group's transaction update data for all lower-numbered container versions that has not yet been evicted from the burst buffer, will be evicted as the result of this call. If evicting the data would result in container versions with open read contexts becoming inaccessible, the evict will fail.

In the second scenario, the data is resident in the burst buffer as the result of a call to H5Gprefetch_ff. This call replicates data from persistent storage (DAOS) to the burst buffer, and we refer to this data as *replica data*. When replica data is evicted, only data in the burst buffer as a result of the exact replica specified is evicted – transaction update data and other replicas for the group remain in the burst buffer.

The `grp_id` parameter specifies the dataset whose data is to be evicted.

The version of the dataset to be evicted is specified by the `container_version` property.

The evict replica property in the access property list, `gapl_id`, is used to specify that a group replica is to be evicted. H5Pset_evict_replica() sets the evict replica property. If this is set, then replica data will be evicted, otherwise transaction update data will be evicted.

The `es_id` parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in `H5_EVENT_STACK_NULL` for the `es_id` parameter.

Limitations / Future Considerations:

In Quarter 7, only the primary IOD KV object associated with the group is evicted. Auxiliary IOD objects will also be evicted in Quarter 8.

When evicting a replica, the `container_version` argument is redundant, as the replica identifier fully specifies the data to be evicted. Consider revisiting and possibly revising the API prior to production release. Possibly have separate evict commands for eviction of transaction update data and of replicas.

For other potential extensions that are beyond the scope of the EFF prototype project, refer to the document *Burst Buffer Space Management – Prototype to Production*.

Parameters:

Use or disclosure of data contained on this sheet is subject to the restriction on the title page of this document.
Copyright © The HDF Group, 2014. All rights reserved

<i>hid_t grp_id</i>	IN: Identifier of the group being evicted.
<i>uint64_t container_version</i>	IN: Container version specifying what version of the group to evict.
<i>hid_t gapl_id</i>	IN: Identifier of an access property list. If the access property list contains an evict replica property (set via <code>H5Pset_evict_replica()</code>), then the replica_id specified by that property will be evicted.
<i>hid_t es_id</i>	IN: The <code>es_id</code> parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in <code>H5_EVENT_STACK_NULL</code> for the <code>es_id</code> parameter.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

Note that when this routine is executed asynchronously, the return value from the routine only indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

History:

Added in Quarter 7.

Man Page Status:

Will need updates in Quarter 8 when additional features are implemented.

4.7.4 H5Gopen_ff

Name: H5Gopen_ff

Signature:

```
hid_t H5Gopen_ff( hid_t loc_id, const char * name, hid_t gapl_id,  
hid_t rctxt_id, hid_t es_id )
```

Purpose:

Open an existing group with a group access property list, possibly asynchronously.

Description:

H5Gopen_ff opens an existing group, `name`, at the location specified by `loc_id`.

`loc_id` may be a file identifier or a group identifier. `name` may be either an absolute path in the file or a relative path from `loc_id` to the group. Both `loc_id` and `name` must be in scope for the read context identified by `rcntxt_id`.

`gapl_id` is a group access property list identifier. No group access property lists have been implemented at this time; use `H5P_DEFAULT`.

`rcntxt_id` indicates the read context for this operation.

The `es_id` parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in `H5_EVENT_STACK_NULL` for the `es_id` parameter.

H5Gopen_ff returns a group identifier for the group that was opened. This group identifier should be released by calling `H5Gclose` when it is no longer needed.

Parameters:

<code>hid_t loc_id</code>	IN: Location identifier May be any HDF5 group identifier or file identifier that is in scope for the read context.
<code>const char *name</code>	IN: Name of the group to open The group name (path to the group) can be specified relative to <code>loc_id</code> or absolute from the file's root group, and must be in scope for the read context.
<code>hid_t gapl_id</code>	IN: Group access property list identifier (No group access properties have been implemented at this time; use <code>H5P_DEFAULT</code> .)
<code>hid_t rctxt_id</code>	IN: Read context identifier indicating the read context for this operation.
<code>hid_t es_id</code>	IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously. Use <code>H5_EVENT_STACK_NULL</code> for synchronous execution

Returns:

Returns a group identifier if successful; otherwise returns a negative value. When executed asynchronously, a future ID for the group is returned initially. Upon completion of the asynchronous operation, the future ID will be transparently modified to be a "normal" group identifier.

When this routine is executed asynchronously, the return value indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

History:

Added in Quarter 3.

Quarter 4: Updated to reflect switch from parameter *H5_request_t request_ptr* to *hid_t eq_id*.

Quarter 5: Changed from transaction to read context id and from event queue to event stack. Added scope requirement.

Man Page Status:

No known issues.

4.7.5 H5Gprefetch_ff

Name: H5Gprefetch_FF

Signature:

```
herr_t H5Gprefetch_ff( hid_t grp_id, hid_t rcntxt_id, hid_t *replica_id, hid_t gapl_id,  
hid_t es_id )
```

Purpose:

Prefetch a group from persistent storage to burst buffer storage, possibly asynchronously.

Description:

H5Gprefetch_ff prefetches a group, specified by its identifier grp_id, from persistent storage (DAOS) into burst buffer storage.

rcntxt_id indicates the read context for this operation.

replica_id, the replica identifier, is set to indicate where the pre-fetched data can be found in the burst buffer, and is passed to subsequent H5Gevict calls.

gapl_id, a group access property list identifier, is not currently used and should be set to H5P_DEFAULT.

The es_id parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in H5_EVENT_STACK_NULL for the es_id parameter.

The replica_id is not valid until the operation has completed, if it is executing asynchronously.

Limitations / Future Considerations:

For the EFF prototype project, only the primary IOD KV object associated with the HDF5 group will be prefetched; auxiliary IOD objects remain on persistent storage (DAOS). For more information, and other potential extensions, refer to the document *Burst Buffer Space Management – Prototype to Production*.

This function was implemented for completeness, but since no access routines accept the group's replica_id in this phase of the project, it has no practical value other than to demonstrate the ability to prefetch and evict groups.

Parameters:

hid_t grp_id	IN: Identifier of the group being prefetched. grp_id must be in scope for the read context.
hid_t rcntxt_id	IN: Read context identifier indicating the read context for this operation.
hid_t *replica_id	IN: Identifier of the replicated data in the burst buffer.
hid_t gapl_id	IN: Identifier of an access property list for this I/O operation. Should be H5P_DEFAULT.
hid_t es_id	IN: The es_id parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in H5_EVENT_STACK_NULL for the es_id parameter.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

Note that when this routine is executed asynchronously, the return value from the routine only indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

History:

Added in Quarter 7.

Man Page Status:

No known issues.

4.8 H5L: Link APIs

These routines are used to create and manipulate Links in an HDF5 group, and are designed to be used in conjunction with the Object APIs (H5O).

The routines ending in _ff have different signatures than the standard HDF5 library routines.

Man pages for routines whose user interface is unchanged from the standard HDF5 implementation can be found at: http://www.hdfgroup.org/HDF5/doc/RM/RM_H5L.html.

Routine	Implemented	Notes
H5Lcopy_ff	Quarter 4	
H5Lcreate_hard_ff	Quarter 4	
H5Lcreate_soft_ff	Quarter 4	
H5Ldelete_ff	Quarter 4	
H5Lexists_ff	Quarter 4	
H5Lmove_ff	Quarter 4	
H5Lget_info_ff	Quarter 6	Man page to be added in Q8
H5Lget_val_ff	Quarter 6	Man page to be added in Q8
H5Literate_ff		Will not implement
H5Literate_by_name_ff		Will not implement
H5Lvisit_ff		Will not implement
H5Lvisit_by_name_ff		Will not implement
H5Lcreate_external		Will not implement
H5Lcreate_ud		Will not implement
H5Ldelete_by_idx		Will not implement
H5Lget_info_by_idx		Will not implement
H5Lget_name_by_idx		Will not implement
H5Lget_val_by_idx		Will not implement
H5Lis_registered		Will not implement
H5Lregister		Will not implement
H5Lunregister		Will not implement

4.8.1 H5Lcopy_ff

Name: H5Lcopy_ff

Signature:

```
herr_t H5Lcopy_FF(hid_t src_loc_id, const char *src_name, hid_t dest_loc_id,  
const char *dest_name, hid_t lcp1_id, hid_t lap1_id, hid_t trans_id, hid_t es_id)
```

Purpose:

Copies a link from one location to another, possibly asynchronously.

Description:

H5Lcopy_ff copies the link specified by `src_name` from the file or group specified by `src_loc_id` to the file or group specified by `dest_loc_id`. The new copy of the link is created with the name `dest_name`.

`src_loc_id` and `dest_loc_id` must reference locations in the same file in the EFP stack.
`src_loc_id`, `src_name`, `dest_loc_id`, and `dest_name` must all be in scope for the transaction identified by `trans_id`.

If `dest_loc_id` is a file identifier, `dest_name` will be interpreted relative to that file's root group.

The new link is created with the creation and access property lists specified by `lcp1_id` and `lap1_id`. The interpretation of `lcp1_id` is limited in the manner described in the next paragraph.

H5Lcopy_ff retains the creation time and the target of the original link. However, since the link may be renamed, the character encoding is that specified in `lcp1_id` rather than that of the original link. Other link creation properties are ignored.

If the link is a soft link, also known as a symbolic link, its target is interpreted relative to the location of the copy.

Several properties are available to govern the behavior of H5Lcopy. These properties are set in the link creation and access property lists, `lcp1_id` and `lap1_id`, respectively. In the EFP stack, neither the link creation property list nor the link access property list have any effect. Both should be set to H5P_DEFAULT.

`trans_id` indicates the transaction this operation is part of.

The `es_id` parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in H5_EVENT_STACK_NULL for the `es_id` parameter.

H5Lcopy_ff does not affect the object that the link points to.

Parameters:

<code>hid_t src_loc_id</code>	IN: Location identifier of the source link May be any file or group identifier that is in scope for the transaction
<code>const char *src_name</code>	IN: Name of the link to be copied The name of the link can be specified relative to <code>src_loc_id</code> , absolute from the file's root group, or '.' (a dot), and must be in scope for the transaction.
<code>hid_t dest_loc_id</code>	IN: Location identifier specifying the destination of the copy May be any file or group identifier that is in scope for the transaction.
<code>const char *dest_name</code>	IN: Name to be assigned to the new copy The name of the new copy can be specified relative to <code>dest_loc_id</code> or

absolute from the file's root group, and must be in scope for the transaction.

hid_t lcpl_id

IN: Link creation property list identifier
Currently not used in EFF; specify H5P_DEFAULT.

hid_t lapl_id

IN: Link access property list identifier
Currently not used in EFF; specify H5P_DEFAULT.

hid_t trans_id

IN: Transaction identifier specifying the transaction this operation is a part of.

hid_t es_id

IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously. Use H5_EVENT_STACK_NULL for synchronous execution.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

When this routine is executed asynchronously, the return value indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

History:

Added in Quarter 4.

Quarter 5: Changed from transaction to transaction id and from event queue to event stack. Added scope requirement and noted source and destination must be in same file. Noted lcpl and lapl have no effect in EFF.

Man Page Status:

No known issues.

4.8.2 H5Lcreate_hard_ff

Name: H5Lcreate_hard_ff

Signature:

```
herr_t H5Lcreate_hard_ff( hid_t obj_loc_id, const char *obj_name, hid_t link_loc_id, const char *link_name,  
hid_t lcpl_id, hid_t lapl_id, hid_t trans_id, hid_t es_id )
```

Purpose:

Creates a hard link to an object, possibly asynchronously.

Description:

H5Lcreate_hard_ff creates a new hard link to a pre-existing object in an HDF5 file. The new link may be one of many that point to that object.

The target object must already exist in the scope for the transaction indicated by `trans_id`.

`obj_loc_id` and `obj_name` specify the location and name, respectively, of the target object, i.e., the object that the new hard link points to. Both `obj_loc_id` and `obj_name` must be in scope for the transaction identified by `trans_id`.

`link_loc_id` and `link_name` specify the location and name, respectively, of the new hard link. Both `link_loc_id` and `link_name` must be in scope for the transaction identified by `trans_id`.

`obj_name` and `link_name` are interpreted relative to `obj_loc_id` and `link_loc_id`, respectively.

If `obj_loc_id` and `link_loc_id` are the same location, the HDF5 macro H5L_SAME_LOC can be used for either parameter (but not both).

`lcpl_id` and `lapl_id` are the link creation and access property lists associated with the new link. In the EFF stack, neither the link creation property list nor the link access property list have any effect. Both should be set to H5P_DEFAULT.

`trans_id` indicates the transaction this operation is part of.

The `es_id` parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in H5_EVENT_STACK_NULL for the `es_id` parameter.

Hard and soft links are for use only if the target object is in the current file. If the desired target object is in a different file from the new link, an external link may be created with H5Lcreate_external. External links are not currently supported in the EFF stack.

The HDF5 library keeps a count of all hard links pointing to an object; if the hard link count reaches zero (0), the object will be deleted from the file. Creating new hard links to an object will prevent it from being deleted if other links are removed. The library maintains no similar count for soft links and they can dangle.

Parameters:

`hid_t obj_loc_id` IN: The file or group identifier for the target object.
The identifier must be in scope for the transaction.

`const char *obj_name` IN: Name of the target object, which must already exist.
The object name (path to the object) can be specified relative to `obj_loc_id` or absolute from the file's root group and must be in scope for the transaction.

<i>hid_t link_loc_id</i>	IN: The file or group identifier for the new link. The identifier must be in scope for the transaction.
<i>const char * link_name</i>	IN: The name of the new link. The link name can be specified relative to <code>link_loc_id</code> or absolute from the file's root group and must be in scope for the transaction.
<i>hid_t lcpl_id</i>	IN: Link creation property list identifier. <i>Currently not used in EFF; specify H5P_DEFAULT.</i>
<i>hid_t lapl_id</i>	IN: Link access property list identifier. <i>Currently not used in EFF; specify H5P_DEFAULT.</i>
<i>hid_t trans_id</i>	IN: Transaction identifier specifying the transaction this operation is a part of.
<i>hid_t es_id</i>	IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously. Use <code>H5_EVENT_STACK_NULL</code> for synchronous execution.
.	.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

When this routine is executed asynchronously, the return value indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

History:

Added in Quarter 4.

Quarter 5: Changed from transaction to transaction id and from event queue to event stack. Added scope requirement. Noted that lcpl and lapl are unused in EFF stack.

Man Page Status:

No known issues.

4.8.3 H5Lcreate_soft_ff

Name: H5Lcreate_soft_ff

Signature:

```
herr_t H5Lcreate_soft_ff(const char *target_path, hid_t link_loc_id,  
const char *link_name, hid_t lcpl_id, hid_t lapl_id, hid_t trans_id, hid_t es_id)
```

Purpose:

Creates a soft link to an object, possibly asynchronously.

Description:

H5Lcreate_soft_ff creates a new soft link to an object in an HDF5 file. The new link may be one of many that point to that object.

`target_path` specifies the path to the target object, i.e., the object that the new soft link points to.

`target_path` can be anything and is interpreted at lookup time. This path may be absolute in the file or relative to `link_loc_id`.

`link_loc_id` and `link_name` specify the location and name, respectively, of the new soft link.

`link_name` is interpreted relative to `link_loc_id`. Both `link_loc_id` and `link_name` must be in scope for the transaction identified by `trans_id`.

`lcpl_id` and `lapl_id` are the link creation and access property lists associated with the new link. In the EFF stack, neither the link creation property list nor the link access property list have any effect. Both should be set to `H5P_DEFAULT`.

`trans_id` indicates the transaction this operation is part of.

The `es_id` parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in `H5_EVENT_STACK_NULL` for the `es_id` parameter.

H5Lcreate_soft_ff is for use only if the target object is in the current file. If the desired target object is in a different file from the new link, use H5Lcreate_external to create an external link. External links are not currently supported in the EFF stack.

For instance, if `target_path` is `./foo`, `link_loc_id` specifies `./x/y/bar`, and the name of the new link is `new_link`, then a subsequent request for `./x/y/bar/new_link` will return same the object as would be found at `./foo`.

Soft links and external links are also known as symbolic links as they use a name to point to an object; hard links employ an object's address in the file.

Unlike hard links, a soft link in an HDF5 file is allowed to *dangle*, meaning that the target object need not exist at the time that the link is created.

The HDF5 library does not keep a count of soft links as it does of hard links.

Parameters:

`const char *target_path` IN: Path to the target object, which is not required to exist.

`hid_t link_loc_id` IN: The file or group identifier for the new link.
The identifier must be in scope for the transaction.

<code>const char *link_name</code>	IN: The name of the new link. The link name can be specified relative to <code>link_loc_id</code> or absolute from the file's root group and must be in scope for the transaction.
<code>hid_t lcpl_id</code>	IN: Link creation property list identifier. <i>Currently not used in EFF; specify H5P_DEFAULT.</i>
<code>hid_t lapl_id</code>	IN: Link access property list identifier. <i>Currently not used in EFF; specify H5P_DEFAULT.</i>
<code>hid_t trans_id</code>	IN: Transaction identifier specifying the transaction this operation is a part of.
<code>hid_t es_id</code>	IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously. Use <code>H5_EVENT_STACK_NULL</code> for synchronous execution.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

When this routine is executed asynchronously, the return value indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

History:

Added in Quarter 4.

Quarter 5: Changed from transaction to transaction id and from event queue to event stack. Added scope requirement. Noted that lcpl and lapl are unused in EFF stack and that external links are unsupported.

Man Page Status:

No known issues.

4.8.4 H5Ldelete_ff

Name: H5Ldelete_ff

Signature:

```
herr_t H5Ldelete_ff(hid_t loc_id, const char *name, hid_t lapl_id, hid_t trans_id,  
hid_t es_id )
```

Purpose:

Removes a link from a group.

Description:

H5Ldelete_ff removes the link specified by name from the location loc_id. Both name and loc_id must be in scope for the transaction indicated by trans_id.

lapl_id specifies the link access property list. The link access property list currently has no effect in the EFF stack and should be set to H5P_DEFAULT.

trans_id indicates the transaction this operation is part of.

The es_id parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in H5_EVENT_STACK_NULL for the es_id parameter.

If the link being removed is a hard link, H5Ldelete_ff also decrements the link count for the object to which name points. Unless there is a duplicate hard link in that group, this action removes the object to which name points from the group that previously contained it.

A reference count keeps track of how many hard links refer to an object; when the reference count reaches zero, the object can be removed from the file. Objects that are open are not removed until all identifiers to the object are closed.

Warning:

Exercise caution in the use of H5Ldelete_ff; if the link being removed is on the only path leading to an HDF5 object, that object may become permanently inaccessible in the file.

Parameters:

hid_t loc_id	IN: Identifier of the file or group containing the link object. Must be in scope for the transaction.
const char *name	IN: Name of the link to delete. Must be in scope for the transaction.
hid_t lapl_id	IN: Link access property list identifier. <i>Currently not used in EFF; specify H5P_DEFAULT.</i>
hid_t trans_id	IN: Transaction identifier specifying the transaction this operation is a part of.
hid_t es_id	IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously. Use H5_EVENT_STACK_NULL for synchronous execution.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

When this routine is executed asynchronously, the return value indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

History:

Added in Quarter 4.

Quarter 5: Changed from transaction to transaction id and from event queue to event stack. Added scope requirement. Noted that lapl is not used in EFF.

Man Page Status:

No known issues.

4.8.5 H5Lexists_ff

Name: H5Lexists_ff

Signature:

```
herr_t H5Lexists_ff(hid_t loc_id, const char *name, hid_t lapl_id, hbool_t *exists,  
hid_t rcntxt_id, hid_t es_id )
```

Purpose:

Determine whether a link with the specified name exists in a group.

Description:

H5Lexists_ff allows an application to determine whether the link name exists in the group or file specified with `loc_id`. The link may be of any type; only the presence of a link with that name is checked. Both `name` and `loc_id` must be in scope for the read context indicated by `rcntxt_id`.

The results of the existence test are put in `exists`. Note that this value will not be set until after the function completes, which may be later than when the call returns for asynchronous execution. Traditionally the existence results were indicated by the function call's return value, but with the support for asynchronous execution that is no longer possible and the FastForward version included the new `exists` parameter.

The link access property list, `lapl_id`, may provide information regarding the properties of links. The link access property list currently has no effect in the EFF stack and should be set to `H5P_DEFAULT`.

`rcntxt_id` indicates the read context for this operation.

The `es_id` parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in `H5_EVENT_STACK_NULL` for the `es_id` parameter.

Note that H5Lexists_ff verifies only that the target link exists. If `name` includes either a relative path or an absolute path to the target link, intermediate steps along the path must be verified before the existence of the target link can be safely checked. If the path is not verified and an intermediate element of the path does not exist, H5Lexists_ff will fail. The example in the next paragraph illustrates one step-by-step method for verifying the existence of a link with a relative or absolute path.

Example: Use the following steps to verify the existence of the link `datasetD` in the group `group1/group2/softlink_to_group3/`, where `group1` is a member of the group specified by `loc_id`:

- First use H5Lexists_ff to verify that `group1` exists.
- If `group1` exists, use H5Lexists_ff again, this time with `name` set to `group1/group2`, to verify that `group2` exists.
- If `group2` exists, use H5Lexists_ff with `name` set to `group1/group2/softlink_to_group3` to verify that `softlink_to_group3` exists.
- If `softlink_to_group3` exists, you can now safely use H5Lexists_ff with `name` set to `group1/group2/softlink_to_group3/datasetD` to verify that the target link, `datasetD`, exists.

If the link to be verified is specified with an absolute path, the same approach should be used, but starting with the first link in the file's root group. For instance, if `datasetD` were in `/group1/group2/softlink_to_group3`, the first call to H5Lexists_ff would have `name` set to `/group1`.

Note that this is an outline and does not include all necessary details. Depending on circumstances, for example, you may need to verify that an intermediate link points to a group and that a soft link points to an existing target.

Parameters:

<i>hid_t loc_id</i>	IN: Identifier of the file or group to query.
<i>const char *name</i>	IN: The name of the link to check.
<i>hid_t lapl_id</i>	IN: Link access property list identifier. <i>Currently not used in EFF; specify H5P_DEFAULT.</i>
<i>hbool_t *exists</i>	OUT: Pointer to returned results indicating existence of link. When successful, will be a positive value to indicate the link exists and 0 (zero) if the link does not exist. The value pointed to will not be modified if the existence test failed for some reason.
<i>hid_t rcntxt_id</i>	IN: Read context identifier indicating the read context for this operation.
<i>hid_t es_id</i>	IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously. Use H5_EVENT_STACK_NULL for synchronous execution.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

When this routine is executed asynchronously, the return value indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack. When the asynchronous execution completes successfully, `exists` will contain the existence test results.

History:

Added in Quarter 4.

Quarter 5: Changed from transaction to read context id and from event queue to event stack. Added scope requirement. Noted that link access property list is not used.

Man Page Status:

No known issues.

4.8.6 H5Lmove_ff

Name: H5Lmove_ff

Signature:

```
herr_t H5Lmove_ff(hid_t src_loc_id, const char *src_name, hid_t dest_loc_id,  
const char *dest_name, hid_t lcpl_id, hid_t lapl_id, hid_t trans_id, hid_t es_id )
```

Purpose:

Moves a link within an HDF5 file, possibly asynchronously.

Description:

H5Lmove_tt moves a link within an HDF5 file. The original link, `src_name`, is removed from `src_loc_id` and the new link, `dest_name`, is inserted at `dest_loc_id`. This change is accomplished as an atomic operation.

`src_loc_id` and `src_name` identify the original link. `src_loc_id` is either a file or group identifier; `src_name` is the path to the link and is interpreted relative to `src_loc_id`. `dest_loc_id` and `dest_name` identify the new link. `dest_loc_id` is either a file or group identifier; `dest_name` is the path to the link and is interpreted relative to `dest_loc_id`.

`src_loc_id` and `dest_loc_id` must reference locations in the same file in the EFF stack. `src_loc_id`, `src_name`, `dest_loc_id`, and `dest_name` must all be in scope for the transaction identified by `trans_id`.

Note that H5Lmove_tt does not modify the value of the link; the new link points to the same object as the original link pointed to. Furthermore, if the object pointed to by the original link was already open with a valid object identifier, that identifier will remain valid after the call to H5Lmove_tt.

`lcpl_id` and `lapl_id` are the link creation and link access property lists, respectively, associated with the new link, `dest_name`. Neither of these property lists currently have any effect in the EFF stack and

`trans_id` indicates the transaction this operation is part of.

The `es_id` parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in H5_EVENT_STACK_NULL for the `es_id` parameter.

Warning:

Exercise care in moving links, as it is possible to render data in a file inaccessible with H5Lmove. If the link being moved is on the only path leading to an HDF5 object, that object may become permanently inaccessible in the file.

Parameters:

<code>hid_t src_loc_id</code>	IN: Original file or group identifier.
<code>const char *src_name</code>	IN: Original link name.
<code>hid_t dest_loc_id</code>	IN: Destination file or group identifier.
<code>const char *dest_name</code>	IN: New link name.
<code>hid_t lcpl_id</code>	IN: Link creation property list identifier to be associated with the new link. <i>Currently not used in EFF; specify H5P_DEFAULT.</i>

<i>hid_t lapl_id</i>	IN: Link access property list identifier to be associated with the new link. <i>Currently not used in EFF; specify H5P_DEFAULT.</i>
<i>hid_t trans_id</i>	IN: Transaction identifier specifying the transaction this operation is a part of.
<i>hid_t es_id</i>	IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously. Use H5_EVENT_STACK_NULL for synchronous execution.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

When this routine is executed asynchronously, the return value indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

History:

Added in Quarter 4.

Quarter 5: Changed from transaction to transaction id and from event queue to event stack. Added scope requirement and noted source and destination must be in same file. Updated text in property list description to match what EFF stack supports.

Man Page Status:

No known issues.

4.9 H5M: Map APIs

These routines are used to read or write HDF5 Map objects, which were added to support ACG use cases. Map objects, which offer a key-value store data structure, should also have wide applicability to other application.

Map objects in HDF5 are similar to a typical “map” data structure in computer science, and are implemented in the EFF stack using IOD Key-Value objects.

HDF5 maps set/get a value in the object, according to the key provided, with a 1-1 mapping of keys to values. All keys for a given map object must be of the same HDF5 datatype. All values for a given map object must be of the same HDF5 datatype.

HDF5 maps are leaf objects in the group hierarchy within a container. Attributes can be attached to map objects.

Map APIs were designed for the EFF stack and also for the traditional HDF5 Binary HDF5 file format storage. EFF-specific versions have the “_ff” suffix and include one or more of (1) a read context id, (2) a transaction id, and/or (3) an event stack id. The Map APIs for the traditional HDF5 format do not have the _ff suffix or the parameters listed, and have not been implemented or explicitly included in this section.

These routines did not exist prior to the Exascale FastForward project.

Routine	Implemented	Notes
H5Mclose_ff	Quarter 5	
H5Mcreate_ff	Quarter 5	
H5Mdelete_ff	Quarter 5	
H5Mevict_ff	Quarter 7	
H5Mexists_ff	Quarter 5	
H5Mget_ff	Quarter 5	
H5Mget_count_ff	Quarter 5	
H5Mget_types_ff	Quarter 5	
H5Miterate_ff		Will not implement
H5Mopen_ff	Quarter 5	
H5Mprefetch_ff	Quarter 7	
H5Mset_ff	Quarter 5	

4.9.1 H5Mclose_ff

Name: H5Mclose_ff

Signature:

herr_t H5Mclose_ff(hid_t map_id, hid_t es_id)

Purpose:

Close the specified map object, possibly asynchronously.

Description:

H5Mclose_ff ends access to the map object specified by `map_id` and releases resources used by it. Further use of the map identifier is illegal in calls to the map API.

The `es_id` parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in `H5_EVENT_STACK_NULL` for the `es_id` parameter.

Parameters:

<code>hid_t map_id</code>	IN: Identifier of the map object to close access to.
<code>hid_t es_id</code>	IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously. Use <code>H5_EVENT_STACK_NULL</code> for synchronous execution.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

When this routine is executed asynchronously, the return value indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

History:

Added in Quarter 5.

Man Page Status:

No known issues.

4.9.2 H5Mcreate_ff

Name: H5Mcreate_ff

Signature:

```
hid_t H5Mcreate_ff ( hid_t loc_id, const char *name, hid_t keytype_id, hid_t valtype_id,
hid_t lcpl_id, hid_t mcpl_id, hid_t mapl_id, hid_t trans_id, hid_t es_id )
```

Purpose:

Create a new map object and link it into the file, possibly asynchronously.

Description:

H5Mcreate_ff creates a new map object named name at the location specified by loc_id.

loc_id may be a file identifier or a group identifier. name may be either an absolute path in the file or a relative path from loc_id naming the map object. Both loc_id and name must be in scope for the transaction identified by trans_id.

keytype_id specifies the datatype for all keys in the map object and valtype_id specifies the datatype for all values in the object.

If keytype_id or valtype_id are either fixed-length or variable-length strings, it is important to set the string length when defining the datatype. String datatypes are derived from H5T_C_S1, which defaults to 1 character in size. See [H5Tset_size](#) and “[Creating variable-length string datatypes](#).”

The EFF stack currently does not support variable-length datatypes for keys.

The link creation property list, lcpl_id, governs creation of the link(s) by which the new map object is accessed and the creation of any intermediate groups that may be missing. In the EFF stack, automatic creation of missing intermediate groups (controlled by H5Pset_create_intermediate_group) is not supported and none of the properties in the link creation property list currently have any effect.

The map creation property list, mcpl_id, and map access property list, mapl_id, modify the new map object’s behavior.

H5P_DEFAULT must be used for all three of the property list parameters in the current EFF implementation.

trans_id indicates the transaction this operation is part of.

The es_id parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in H5_EVENT_STACK_NULL for the es_id parameter.

To conserve and release resources, the map object id returned from this routine must be closed with H5Mcclose when access is no longer required.

Parameters:

hid_t loc_id IN: Location identifier
May be any HDF5 group identifier or file identifier that is in scope for the transaction

const char *name IN: name
The map object name (path to the map object) can be specified relative to loc_id or absolute from the file’s root group, and must be in scope for the transaction.

hid_t keytype_id IN: Datatype identifier for the map keys

Use or disclosure of data contained on this sheet is subject to the restriction on the title page of this document.
Copyright © The HDF Group, 2014. All rights reserved

<i>hid_t valtype_id</i>	IN: Datatype identifier for the map values
<i>hid_t lcpl_id</i>	IN: Link creation property list <i>Currently not used in EFF; specify H5P_DEFAULT.</i>
<i>hid_t mcpl_id</i>	IN: Map creation property list <i>Currently not used; specify H5P_DEFAULT.</i>
<i>hid_t mapl_id</i>	IN: Map access property list <i>Currently not used; specify H5P_DEFAULT.</i>
<i>hid_t trans_id</i>	IN: Transaction identifier specifying the transaction this operation is a part of.
<i>hid_t es_id</i>	IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously. Use <code>H5_EVENT_STACK_NULL</code> for synchronous execution.

Returns:

Returns a map object identifier if successful; otherwise returns a negative value. When executed asynchronously, a future ID for the new map object is returned initially. Upon completion of the asynchronous operation, the future ID will be transparently modified to be a “normal” map object identifier.

When this routine is executed asynchronously, the return value indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

History:

Added in Quarter 5.

Man Page Status:

No known issues.

4.9.3 H5Mdelete_ff

Name: H5Mdelete_ff

Signature:

```
herr_t H5Mdelete_ff(hid_t map_id, hid_t key_mem_type_id, const void *key, hid_t trans_id,  
hid_t es_id )
```

Purpose:

Delete a key/value pair from a map object.

Description:

H5Mdelete_ff removes a key/value pair from the map object given by `map_id`. `map_id` must be in scope for the transaction indicated by `trans_id`.

The `key` parameter points to the key of the map entry (the key/value pair) that will be deleted, and `key_mem_type_id` specifies the datatype of the `key` parameter.

If `key_mem_type_id` is not the same as the datatype of the keys in the map object, the `key` parameter will automatically undergo datatype conversion in order to locate the correct key in the map object.

If the key is not in the map file, a negative value will be returned if the routine is called synchronously (`es_id` is `H5_EVENT_STACK_NULL`). If the routine is called asynchronously, the completion status of the asynchronous operation will be `H5ES_STATUS_FAIL`.

`trans_id` indicates the transaction this operation is part of.

The `es_id` parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in `H5_EVENT_STACK_NULL` for the `es_id` parameter.

Parameters:

<code>hid_t map_id</code>	IN: Identifier of the map object. Must be in scope for the transaction.
<code>hid_t key_mem_type_id</code>	IN: Datatype of the <code>key</code> parameter.
<code>const void *key</code>	IN: Pointer to key of the map key/value pair that will be deleted.
<code>hid_t trans_id</code>	IN: Transaction identifier specifying the transaction this operation is a part of.
<code>hid_t es_id</code>	IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously. Use <code>H5_EVENT_STACK_NULL</code> for synchronous execution.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

When this routine is executed asynchronously, the return value indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

History:

Added in Quarter 5.

Use or disclosure of data contained on this sheet is subject to the restriction on the title page of this document.
Copyright © The HDF Group, 2014. All rights reserved

Man Page Status:

No known issues.

Use or disclosure of data contained on this sheet is subject to the restriction on the title page of this document.
Copyright © The HDF Group, 2014. All rights reserved

4.9.4 H5Mevict_ff

Name: H5Mevict_ff

Signature:

```
herr_t H5Mevict_ff( hid_t map_id, uint64_t *container_version, hid_t map1_id,  
hid_t es_id )
```

Purpose:

Evict map object from the burst buffer, possibly asynchronously.

Description:

H5Mevict_ff evicts data associated with a map object from the burst buffer.

The data to be evicted may be resident in the burst buffer under two different scenarios

In the first scenario, the data is resident in the burst buffer as the result of updates to the map object made via the EFF transaction model. For example, through calls to H5Mcreate_ff or H5Mset_ff. These calls add updates to a transaction that are atomically applied to the map object when the transaction is committed, and we refer to this data as *transaction update data*.

When evicting transaction update data, the container version being evicted should first be persisted to permanent storage (DAOS), with the H5RCpersist command. The map object's transaction update data for the specified container version, as well as the map object's transaction update data for all lower-numbered container versions that has not yet been evicted from the burst buffer, will be evicted as the result of this call. If evicting the data would result in container versions with open read contexts becoming inaccessible, the evict will fail.

In the second scenario, the data is resident in the burst buffer as the result of a call to H5Mprefresh_ff. This call replicates data from persistent storage (DAOS) to the burst buffer, and we refer to this data as *replica data*. When replica data is evicted, only data in the burst buffer as a result of the exact replica specified is evicted – transaction update data and other replicas for the dataset remain in the burst buffer.

The `map_id` parameter specifies the map object whose data is to be evicted.

The version of the map object to be evicted is specified by the `container_version` property.

The evict replica property in the access property list, `map1_id`, is used to specify that a map object replica is to be evicted. H5Pset_evict_replica() sets the evict replica property. If this is set, then replica data will be evicted, otherwise transaction update data will be evicted.

The `es_id` parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in `H5_EVENT_STACK_NULL` for the `es_id` parameter.

Limitations / Future Considerations:

In Quarter 7, only the primary IOD KV object associated with the HDF5 map object is evicted. Auxiliary IOD objects will also be evicted in Quarter 8.

When evicting a replica, the `container_version` argument is redundant, as the replica identifier fully specifies the data to be evicted. Consider revisiting and possibly revising the API prior to production release. Possibly have separate evict commands for eviction of transaction update data and of replicas.

For other potential extensions that are beyond the scope of the EFF prototype project, refer to the document *Burst Buffer Space Management – Prototype to Production*.

Parameters:

Use or disclosure of data contained on this sheet is subject to the restriction on the title page of this document.
Copyright © The HDF Group, 2014. All rights reserved

<i>hid_t map_id</i>	IN: Identifier of the map object being evicted.
<i>uint64_t container_version</i>	IN: Container version specifying what version of the map object to evict.
<i>hid_t map1_id</i>	IN: Identifier of an access property list. If the access property list contains an evict replica property (set via <code>H5Pset_evict_replica()</code>), then the replica_id specified by that property will be evicted.
<i>hid_t es_id</i>	IN: The <code>es_id</code> parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in <code>H5_EVENT_STACK_NULL</code> for the <code>es_id</code> parameter.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

Note that when this routine is executed asynchronously, the return value from the routine only indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

History:

Added in Quarter 7.

Man Page Status:

Will need updates in Quarter 8 when additional features are implemented.

4.9.5 H5Mexists_ff

Name: H5Mexists_ff

Signature:

```
herr_t H5Mexists_ff(hid_t map_id, hid_t key_mem_type_id, const void *key, hbool_t *exists,  
hid_t rctxt_id, hid_t es_id)
```

Purpose:

Determine whether a key exists in a map object, possibly asynchronously.

Description:

H5Mexists_ff determines whether the key, `key`, exists in the map object specified by `map_id`. `map_id` must be in scope for the read context identified by `rcntxt_id`.

The `key` parameter points to the key whose existence is being checked, and `key_mem_type_id` specifies the datatype of the `key` parameter.

If `key_mem_type_id` is not the same as the datatype of the keys in the map object, the `key` parameter will automatically undergo datatype conversion in order to locate the correct key in the map object.

The results of the existence test are put in `exists`. Note that this value will not be set until after the function completes, which may be later than when the call returns for asynchronous execution.

`rcntxt_id` indicates the read context for this operation.

The `es_id` parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in `H5_EVENT_STACK_NULL` for the `es_id` parameter.

Parameters:

<code>hid_t map_id</code>	IN: Identifier of the map object. Must be in scope for the read context.
<code>hid_t key_mem_type_id</code>	IN: Datatype of the <code>key</code> parameter.
<code>const void *key</code>	IN: Pointer to key whose existence in the map object is being checked.
<code>hbool_t *exists</code>	OUT: Pointer to returned results indicating existence of key. When successful, will be a positive value to indicate the key exists in the map object and 0 (zero) if the key does not exist. The value pointed to will not be modified if the existence test fails with an error.
<code>hid_t rctxt_id</code>	IN: Read context identifier indicating the read context for this operation.
<code>hid_t es_id</code>	IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously. Use <code>H5_EVENT_STACK_NULL</code> for synchronous execution.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

When this routine is executed asynchronously, the return value indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous

operation must be checked separately through the event stack. When the asynchronous execution completes successfully, `exists` will contain the existence test results.

History:

Added in Quarter 5.

Man Page Status:

No known issues.

4.9.6 H5Mget_ff

Name: H5Mget_ff

Signature:

```
herr_t H5Mget_ff(hid_t map_id, hid_t key_mem_type_id, const void *key,
hid_t val_mem_type_id, void *value, hid_t dxpl_id, hid_t rcntxt_id, hid_t es_id)
```

Purpose:

Retrieve the value for a given key from a map object, possibly asynchronously.

Description:

H5Mget_ff retrieves the value for the key, `key`, from the map object specified by `map_id`. `map_id` must be in scope for the read context identified by `rcntxt_id`.

The `key` parameter points to the memory buffer holding the key whose value is being retrieved; `key_mem_type_id` specifies the datatype of the `key` parameter.

The `value` parameter points to the memory buffer where the value should be stored; `val_mem_type_id` specifies the datatype of the `value` parameter.

If `key_mem_type_id` is not the same as the datatype of the keys in the map object, the `key` parameter will automatically undergo datatype conversion in order to locate the correct key in the map object. If `val_mem_type_id` is not the same as the datatype of the values in the map object, the value retrieved will automatically undergo datatype conversion when placed in the `value` buffer.

The data transfer property list, `dxpl_id`, may modify the operation's behaviour. When the user sets a property (via a call to H5Pset_dxpl_checksum_ptr) to pass a checksum pointer to the library, the checksum for `value` (not `key`) is placed at the pointer location when the operation completes. To get from a replica that was previously brought into the burst buffer by a call to H5Mprefetch_ff, use the read replica property in the data transfer property list. The call H5Pset_read_replica() sets the read replica property.

`rcntxt_id` indicates the read context for this operation.

The `es_id` parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in H5_EVENT_STACK_NULL for the `es_id` parameter.

Parameters:

`hid_t map_id` IN: Identifier of the map object.
Must be in scope for the read context.

`hid_t key_mem_type_id` IN: Datatype of the `key` parameter.

`const void *key` IN: Pointer to key whose value is being retrieved.

`hid_t val_mem_type_id` IN: Datatype of the `value` parameter.

`void *value` OUT: Pointer to buffer for the retrieved value.

`hid_t dxpl_id` IN: Data transfer property list. If specified, the read replica property directs the get to access pre-fetched data.

<i>hid_t rcntxt_id</i>	IN: Read context identifier indicating the read context for this operation.
<i>hid_t es_id</i>	IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously. Use H5_EVENT_STACK_NULL for synchronous execution.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

When this routine is executed asynchronously, the return value indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

History:

Added in Quarter 5.

Quarter 7: Added information about gets from prefetched replicas.

Man Page Status:

No known issues.

4.9.7 H5Mget_count_ff

Name: H5Mget_count_ff

Signature:

*herr_t H5Mget_count_ff(hid_t map_id, hsize_t *count, hid_t rctxt_id, hid_t es_id)*

Purpose:

Retrieve the number of key/value pairs in a map object, possibly asynchronously.

Description:

H5Mget_count_ff retrieves the number of key/value pairs in the map object specified by `map_id`. `map_id` must be in scope for the read context identified by `rcntxt_id`.

The number of pairs is returned in `count`.

`rcntxt_id` indicates the read context for this operation.

The `es_id` parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in `H5_EVENT_STACK_NULL` for the `es_id` parameter.

Parameters:

<code>hid_t map_id</code>	IN: Identifier of the map object. Must be in scope for the read context.
<code>hsize_t *count</code>	OUT: The number of key/value pairs in the map object.
<code>hid_t rctxt_id</code>	IN: Read context identifier indicating the read context for this operation.
<code>hid_t es_id</code>	IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously. Use <code>H5_EVENT_STACK_NULL</code> for synchronous execution.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

When this routine is executed asynchronously, the return value indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

History:

Added in Quarter 5.

Man Page Status:

No known issues.

4.9.8 H5Mget_types_ff

Name: H5Mget_types_ff

Signature:

```
herr_t H5Mget_types_ff(hid_t map_id, hid_t *key_type_id, hid_t *val_type_id,  
hid_t rctxt_id, hid_t es_id)
```

Purpose:

Retrieve the datatypes for the keys and values of a map object, possibly asynchronously.

Description:

H5Mget_types_ff retrieves the datatypes for the keys and values of the map object specified by `map_id`. `map_id` must be in scope for the read context identified by `rcntxt_id`.

The datatype of the keys is returned in `key_type_id` and the datatype of the values is returned in `val_type_id`.

`rcntxt_id` indicates the read context for this operation.

The `es_id` parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in `H5_EVENT_STACK_NULL` for the `es_id` parameter.

Parameters:

<code>hid_t map_id</code>	IN: Identifier of the map object. Must be in scope for the read context.
<code>hid_t *key_type_id</code>	OUT: Datatype of the keys in the map object.
<code>hid_t *val_type_id</code>	OUT: Datatype of the values in the map object.
<code>hid_t rctxt_id</code>	IN: Read context identifier indicating the read context for this operation.
<code>hid_t es_id</code>	IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously. Use <code>H5_EVENT_STACK_NULL</code> for synchronous execution.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

When this routine is executed asynchronously, the return value indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

History:

Added in Quarter 5.

Man Page Status:

No known issues.

4.9.9 H5Miterate_ff

Name: H5Miterate_ff

Signature:

```
herr_t H5Miterate_ff(hid_t map_id, hid_t key_mem_type_id, hid_t val_mem_type_id,  
H5M_iterate_func_t callback_func, void *func_data, hid_t rctxt_id )
```

Purpose:

Iterate over the key/value pairs in a map object, possibly asynchronously.

Description:

H5Miterate_ff iterates over the key/values pairs in the specified map object and invokes a user-defined callback routine for each pair in the map.

The map object to be iterated over is specified by `map_id`. `map_id` must be in scope for the read context identified by `rcntxt_id`.

`key_mem_type_id` and `val_mem_type_id` specify the datatypes for the keys and values that will be presented to the callback routine.

The user-defined callback function, specified by `callback_func`, is invoked for each pair in the map and passed the key, value, and user-provided `func_data`. Note that `func_data` can include a transaction id or read context, allowing the user's callback function to do updates to or reads from an HDF5 container as part of its operation.

The prototype for `H5M_iterate_func_t` is:

```
typedef herr_t (*H5M_iterate_func_t)(const void *key, const void *value,  
void *func_data);
```

The iteration callback routine should obey the same rules as other HDF5 iteration callbacks: return `H5_ITER_ERROR` for an error condition (which will stop iteration), `H5_ITER_CONT` for success (with continued iteration) and `H5_ITER_STOP` for success (and stop iteration).

`rcntxt_id` indicates the read context for this operation.

There is no asynchronous execution option for this function because the user callback function cannot be invoked asynchronously.

Parameters:

<code>hid_t map_id</code>	IN: Identifier of the map object. Must be in scope for the read context.
<code>hid_t key_mem_type_id</code>	IN: Datatype for the key that is presented to the callback routine.
<code>hid_t val_mem_type_id</code>	IN: Datatype for the value that is presented to the callback routine.
<code>H5M_iterate_func_t callback_func</code>	IN: User's function to pass each attribute to
<code>void *func_data</code>	IN/OUT: User's data to pass through to callback function
<code>hid_t rctxt_id</code>	IN: Read context identifier indicating the read context for this operation.

Returns:

Use or disclosure of data contained on this sheet is subject to the restriction on the title page of this document.
Copyright © The HDF Group, 2014. All rights reserved

Returns a non-negative value if successful; otherwise returns a negative value.

History:

Quarter 5: Man page added but not yet implemented.

Man Page Status:

No known issues.

Future consideration: Should this have a dxpl so that values can be checksummed?

Future consideration: Should you be able to specify a replica id so you can read from prefetched data?

4.9.10 H5Mopen_ff

Name: H5Mopen_ff

Signature:

```
hid_t H5Mopen_ff( hid_t loc_id, const char *name, hid_t mapl_id, hid_t rcntxt_id,  
hid_t es_id)
```

Purpose:

Open an existing map object possibly asynchronously.

Description:

H5Mopen_ff opens the existing map object specified by loc_id and name.

loc_id may be a file identifier or a group identifier. name may be either an absolute path in the file or a relative path from loc_id naming the map object. Both loc_id and name must be in scope for the read context identified by rcntxt_id.

The map access property list, mapl_id, provides information regarding access to the map object. The map access property list currently has no effect and should be set to H5P_DEFAULT.

rcntxt_id indicates the read context for this operation.

The es_id parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in H5_EVENT_STACK_NULL for the es_id parameter.

To conserve and release resources, the map object should be closed with H5Mclose_ff when access is no longer required.

Parameters:

hid_t loc_id IN: Location identifier
May be any HDF5 group identifier or file identifier that is in scope for the read context.

const char *name IN: Map object name
The map object name (path to the map object) can be specified relative to loc_id or absolute from the file's root group, and must be in scope for the read context.

hid_t mapl_id IN: Map object access property list
Currently not used; specify H5P_DEFAULT.

hid_t rcntxt_id IN: Read context identifier indicating the read context for this operation.

hid_t eq_id IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously.
Use H5_EVENT_STACK_NULL for synchronous execution.

Returns:

Returns a map object identifier if successful; otherwise returns a negative value. When executed asynchronously, a future ID for the map object is returned initially. Upon completion of the asynchronous operation, the future ID will be transparently modified to be a "normal" map object identifier.

When this routine is executed asynchronously, the return value indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

History:

Added in Quarter 5.

Man Page Status:

No known issues.

4.9.11 H5Mprefetch_ff

Name: H5Mprefetch_FF

Signature:

```
herr_t H5Mprefetch_ff( hid_t map_id, hid_t rcntxt_id, hid_t *replica_id, hid_t mapl_id,  
hid_t es_id )
```

Purpose:

Prefetch all or part of a map object from persistent storage to burst buffer storage, possibly asynchronously.

Description:

H5Mprefetch_ff prefetches all or part of a map object, specified by its identifier `map_id`, from persistent storage (DAOS) into burst buffer storage.

`rcntxt_id` indicates the read context for this operation.

`replica_id`, the replica identifier, is set to indicate where the pre-fetched key-value data can be found in the burst buffer, and is passed to subsequent H5Dread_ff and H5Devict_ff calls.

`mapl_id`, a map access property list identifier, is used to specify partial maps (key range selection) and control layout of the fetched data on the burst buffers.

The `es_id` parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in H5_EVENT_STACK_NULL for the `es_id` parameter.

The `replica_id` is not valid until the operation has completed, if it is executing asynchronously.

Limitations / Future Considerations:

For the EFF prototype project, only the primary IOD KV object associated with the HDF5 map will be prefetched; auxiliary IOD objects (including the Blob objects that hold variable-length data) remain on persistent storage (DAOS). For more information, and other potential extensions, refer to the document *Burst Buffer Space Management – Prototype to Production*.

In the EFF prototype project, only the H5Mget_ff routine accepts the `replica_id` and can read from prefetched map objects. Other routines, such as H5Mexists_ff, should also be able to access the prefetched data in a production release.

In Quarter 7, only complete maps and the default layout are supported.

Parameters:

<code>hid_t map_id</code>	IN: Identifier of the map being prefetched. <code>map_id</code> must be in scope for the read context.
<code>hid_t rcntxt_id</code>	IN: Read context identifier indicating the read context for this operation.
<code>hid_t *replica_id</code>	IN: Identifier of the replicated data in the burst buffer.
<code>hid_t mapl_id</code>	IN: Identifier of an access property list for this I/O operation.
<code>hid_t es_id</code>	IN: The <code>es_id</code> parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in H5_EVENT_STACK_NULL for the <code>es_id</code> parameter.

Returns:

Use or disclosure of data contained on this sheet is subject to the restriction on the title page of this document.

Copyright © The HDF Group, 2014. All rights reserved

Returns a non-negative value if successful; otherwise returns a negative value.

Note that when this routine is executed asynchronously, the return value from the routine only indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

History:

Added in Quarter 7.

Man Page Status:

Will need updates in Quarter 8 when additional features are implemented.

4.9.12 H5Mset_ff

Name: H5Mset_ff

Signature:

```
herr_t H5Mset_ff(hid_t map_id, hid_t key_mem_type_id, const void *key,  
hid_t val_mem_type_id, const void *value, hid_t dxpl_id, hid_t trans_id, hid_t es_id)
```

Purpose:

Set the value for a given key in a map object, possibly asynchronously.

Description:

H5Mset_ff sets the value for a key in a map object. If the key did not exist in the map object previously, this function creates a new key/value pair in the map. If the key already existed in the map object, this function overwrites the previous value in the existing key/value pair.

The map object is specified by `map_id`, which must be in scope for the transaction identified by `trans_id`.

The `key` parameter points to the memory buffer holding the key whose value is being set; `key_mem_type_id` specifies the datatype of the `key` parameter.

The `value` parameter points to the memory buffer of the value that will be stored; `val_mem_type_id` specifies the datatype of the `value` parameter.

If `key_mem_type_id` is not the same as the datatype of the keys in the map object, the `key` parameter will automatically undergo datatype conversion in order to locate the correct key in the map object. If `val_mem_type_id` is not the same as the datatype of the values in the map object, the `value` parameter will automatically undergo datatype conversion when stored in the map object.

The data transfer property list, `dxpl_id`, may modify the operation's behaviour. When the user sets a property (via a call to `H5Pset_dxpl_checksum`) to pass a checksum, the checksum given should be for `value` (not `key`).

`trans_id` indicates the transaction this operation is part of.

The `es_id` parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in `H5_EVENT_STACK_NULL` for the `es_id` parameter.

Parameters:

`hid_t map_id` IN: Identifier of the map object.
Must be in scope for the transaction.

`hid_t key_mem_type_id` IN: Datatype of the `key` parameter.

`const void *key` IN: Pointer to key whose value is being set.

`hid_t val_mem_type_id` IN: Datatype of the `value` parameter.

`const void *value` IN: Pointer to buffer holding the new value.

`hid_t dxpl_id` IN: Data transfer property list.

`hid_t trans_id` IN: Transaction identifier indicating the transaction this operation is part of.

Use or disclosure of data contained on this sheet is subject to the restriction on the title page of this document.
Copyright © The HDF Group, 2014. All rights reserved

hid_t es_id IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously. Use H5_EVENT_STACK_NULL for synchronous execution.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

When this routine is executed asynchronously, the return value indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

History:

Added in Quarter 5.

Man Page Status:

No known issues.

4.10 H5O: Object APIs

These routines are used to operate on HDF5 Objects, and are designed to be used in conjunction with the Link APIs (H5L).

The routines ending in _ff have different signatures than the standard HDF5 library routines.

Man pages for routines whose user interface is unchanged from the standard HDF5 implementation can be found at: http://www.hdfgroup.org/HDF5/doc/RM/RM_H5O.html.

Routine	Implemented	Notes
H5Oclose_ff	Quarter 4	
H5Oexists_by_name_ff	Quarter 4	
H5Oget_comment_ff	Quarter 4	
H5Oget_comment_by_name_ff	Quarter 4	
H5Olink_ff	Quarter 4	
H5Oopen_ff	Quarter 5	Q5: Changed from standard HDF5 H5Oopen, which was implemented in Q4.
H5Oset_comment_ff	Quarter 4	
H5Oset_comment_by_name_ff	Quarter 4	
H5Oget_token	Quarter 6	Man page will be added in Q8.
H5Oopen_by_token	Quarter 6	Man page will be added in Q8.
H5Oget_info_ff	Quarter 6	Man page will be added in Q8.
H5Oget_info_by_name_ff	Quarter 6	Man page will be added in Q8.
H5Ovisit_ff		Will not implement
H5Ovisit_by_name_ff		Will not implement
H5Oincr_ref_count		Will not implement unless needed
H5Odecr_ref_count		Will not implement unless needed
H5Ocopy_ff		Q5: Will not implement for prototype. Early version done in Q4, but very limited functionality or value without support for object copies from lower layers in the stack. Cross-container copies with transactions raise additional issues so removed.
H5Oopen_by_addr		Doesn't make sense for FF storage. Will not implement
H5Oget_info_by_idx		Will not implement
H5Oopen_by_idx		Will not implement

Use or disclosure of data contained on this sheet is subject to the restriction on the title page of this document.
Copyright © The HDF Group, 2014. All rights reserved

4.10.1 H5Oclose_ff

Name: H5Oclose_ff

Signature:

herr_t H5Oclose_ff(hid_t object_id, hid_t es_id)

Purpose:

Close an object in an HDF5 file, possibly asynchronously.

Description:

H5Oclose_ff closes the group, dataset, map, or named datatype specified by object_id.

The es_id parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in H5_EVENT_STACK_NULL for the es_id parameter.

This function is the companion to H5Oopen_ff, and has the same effect as calling H5Gclose_ff, H5Dclose_ff, H5Fclose_ff, or H5Tclose_ff.

H5Oclose_ff is *not* used to close a dataspace, attribute, property list, or file.

Parameters:

hid_t object_id IN: Object identifier

hid_t es_id IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously. Use H5_EVENT_STACK_NULL for synchronous execution.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

When this routine is executed asynchronously, the return value indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

History:

Added in Quarter 4.

Quarter 5: Changed from event queue to event stack. Added map object.

Man Page Status:

No known issues.

4.10.2 H5Oexists_by_name_ff

Name: H5Oexists_by_name_ff

Signature:

```
herr_t H5Oexists_by_name_ff(hid_t loc_id, const char * name, hid_t lapl_id, hbool_t  
*exists, hid_t rcntxt_id, hid_t es_id )
```

Purpose:

Determine whether a link resolves to an actual object, possibly asynchronously.

Description:

H5Oexists_by_name_ff allows an application to determine whether the link name in the group or file specified by `loc_id` resolves to an HDF5 object or if the link dangles. The link may be of any type, but hard links will always resolve to objects and do not need to be verified. The EFF stack currently does not support external or user-defined links. `loc_id` and `name` must both be in scope for the read context identified by `rcntxt_id`.

The link access property list, `lapl_id`, may provide information regarding the properties of links. The link access property list currently has no effect in the EFF stack and should be set to `H5P_DEFAULT`.

The results of the existence test are put in `exists`. Note that this value will not be set until after the function completes, which may be later than when the call returns for asynchronous execution. Traditionally the existence results were indicated by the function call's return value, but with the support for asynchronous execution that is no longer possible and the FastForward version included the new `exists` parameter.

`rcntxt_id` indicates the read context for this operation.

The `es_id` parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in `H5_EVENT_STACK_NULL` for the `es_id` parameter.

Note that H5Oexists_by_name_ff only verifies that the target object exists. If `name` includes either a relative path or an absolute path to the target link, intermediate steps along the path must be verified before the existence of the target link can be safely checked. If the path is not verified and an intermediate element of the path does not exist, H5Oexists_by_name_ff will fail. The example in the next paragraph illustrates one step-by-step method for verifying the existence of a link with a relative or absolute path.

Example: Use the following steps to verify the existence of the link `datasetD` in the group `group1/group2/softlink_to_group3/`, where `group1` is a member of the group specified by `loc_id`:

- First use H5Lexists_ff to verify that a link named `group1` exists.
- If `group1` exists, use H5Oexists_by_name_ff to verify that the link `group1` resolves to an object.
- If `group1` exists, use H5Lexists_ff again, this time with `name` set to `group1/group2`, to verify that the link `group2` exists in `group1`.
- If the `group2` link exists, use H5Oexists_by_name_ff to verify that `group1/group2` resolves to an object.
- If `group2` exists, use H5Lexists_ff again, this time with `name` set to `group1/group2/softlink_to_group3`, to verify that the link `softlink_to_group3` exists in `group2`.

- If the `softlink_to_group3` link exists, use `H5Oexists_by_name_ff` to verify that `group1/group2/softlink_to_group3` resolves to an object.
- If `softlink_to_group3` exists, you can now safely use `H5Lexists_ff` with `name` set to `group1/group2/softlink_to_group3/datasetD` to verify that the target link, `datasetD`, exists.
- And finally, if the link `datasetD` exists, use `H5Oexists_by_name_ff` to verify that `group1/group2/softlink_to_group3/datasetD` resolves to an object.

If the link to be verified is specified with an absolute path, the same approach should be used, but starting with the first link in the file's root group. For instance, if `datasetD` were in `/group1/group2/softlink_to_group3`, the first call to `H5Lexists_ff` would have `name` set to `/group1`.

Note that this is an outline and does not include all necessary details. Depending on circumstances, for example, an application may need to verify the type of an object also.

Parameters:

<code>hid_t loc_id</code>	IN: Identifier of the file or group to query. Must be in scope for the read context.
<code>const char *name</code>	IN: The name of the link to check. Must be in scope for the read context.
<code>hid_t lapl_id</code>	IN: Link access property list identifier. <i>Currently not used in EFF; specify H5P_DEFAULT.</i>
<code>hbool_t *exists</code>	IN: Pointer to returned results indicating existence of link. When successful, will be a positive value to indicate the link exists and 0 (zero) if the link does not exist. The value pointed to will not be modified if the existence test failed for some reason.
<code>hid_t rcntxt_id</code>	IN: Read context identifier indicating the read context for this operation.
<code>hid_t es_id</code>	IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously. Use <code>H5_EVENT_STACK_NULL</code> for synchronous execution.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

When this routine is executed asynchronously, the return value indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack. When the asynchronous execution completes successfully, `exists` will contain the existence test results.

History:

Added in Quarter 4.

Quarter 5: Changed from transaction to read context id and from event queue to event stack. Added scope requirement and note that property list currently not used in EFF.

Man Page Status:

No known issues.

4.10.3 H5Oget_comment_ff

Name: H5Oget_comment_ff

Signature:

```
hid_t H5Oget_comment_ff(hid_t object_id, char *comment, size_t bufsize,  
ssize_t *comment_size, hid_t rctxt_id, hid_t es_id )
```

Purpose:

Retrieve comment for specified object, possibly asynchronously.

Description:

H5Oget_comment_ff retrieves the comment for the specified object and puts it into the buffer `comment`.

The object with the comment is specified by the identifier `object_id`. `object_id` must be in scope for the read context specified by `rcntxt_id`.

The size in bytes of the buffer `comment`, including the NULL terminator, is specified in `bufsize`. If `bufsize` is unknown, a preliminary H5Oget_comment_ff call with the pointer `comment` set to NULL will report the size of the comment *without* the NULL terminator.

If `bufsize` is set to a smaller value than described above, only `bufsize` bytes of the comment, without a NULL terminator, are returned in `comment`.

If an object does not have a comment, the empty string is returned in `comment`.

Upon success, the number of characters in the comment, not including the NULL terminator, or zero (0) if the comment has no comment, is put in `comment_size`. The value returned in `comment_size` may be larger than `bufsize`. On failure, `comment_size` will contain a negative value. Note that this value will not be set until after the function completes, which may be later than when the call returns for asynchronous execution.

Traditionally the comment size was indicated by the function call's return value, but with the support for asynchronous execution that is no longer possible and the FastForward version included the new `comment_size` parameter.

`rcntxt_id` indicates the read context for this operation.

The `es_id` parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in H5_EVENT_STACK_NULL for the `es_id` parameter.

Parameters:

`hid_t object_id` IN: Identifier for the object whose comment is to be retrieved.
 Must be in scope for the read context.

`char *comment` OUT: The comment.

`size_t bufsize` IN: Anticipated required size of the `comment` buffer.

`ssize_t *comment_size` OUT: Pointer to returned results indicating comment size. When successful, will contain the number of characters in the comment, not including the NULL terminator, or zero (0) if the comment has no comment. On failure, will contain a

negative value.

hid_t rcntxt_id IN: Read context identifier indicating the read context for this operation.

hid_t es_id IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously. Use `H5_EVENT_STACK_NULL` for synchronous execution

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

Note that when this routine is executed asynchronously, the return value from the routine only indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack. When the asynchronous execution completes successfully, `comment_size` will contain the number of characters in the comment, not including the NULL terminator, or zero (0) if the comment has no comment. On failure, `comment_size` will be a negative number.

History:

Added in Quarter 4.

Quarter 5: Changed from transaction to read context id and from event queue to event stack. Added scope requirement.

Man Page Status:

No known issues.

4.10.4 H5Oget_comment_by_name_ff

Name: H5Oget_comment_by_name_ff

Signature:

```
hid_t H5Oget_comment_by_name_ff(hid_t loc_id, const char *name, char *comment, size_t  
bufsize, hid_t lapl_id, ssize_t *comment_size, hid_t rcntxt_id, hid_t es_id )
```

Purpose:

Retrieve comment for specified object, possibly asynchronously.

Description:

H5Oget_comment_by_name_ff H5Oget_comment_ff retrieves the comment for the specified object and puts it into the buffer `comment`.

The object with the comment is specified by `loc_id` and `name`. `loc_id` can specify any object in the file. `name` can be one of the following:

- The name of the object relative to `loc_id`
- An absolute name of the object, starting from /, the file's root group
- A dot (.), if `loc_id` fully specifies the object

Both `loc_id` and `name` must be in scope for the read context specified by `rcntxt_id`.

The size in bytes of the buffer `comment`, including the NULL terminator, is specified in `bufsize`. If `bufsize` is unknown, a preliminary H5Oget_comment_by_name_ff call with the pointer `comment` set to NULL will report the size of the comment *without* the NULL terminator.

If `bufsize` is set to a smaller value than described above, only `bufsize` bytes of the comment, without a NULL terminator, are returned in `comment`.

If an object does not have a comment, the empty string is returned in `comment`.

Upon success, the number of characters in the comment, not including the NULL terminator, or zero (0) if the comment has no comment, is put in `comment_size`. The value returned in `comment_size` may be larger than `bufsize`. On failure, `comment_size` will contain a negative value. Note that this value will not be set until after the function completes, which may be later than when the call returns for asynchronous execution. Traditionally the comment size was indicated by the function call's return value, but with the support for asynchronous execution that is no longer possible and the FastForward version included the new `comment_size` parameter.

The link access property list, `lapl_id`, may provide information regarding the properties of links. The link access property list currently has no effect in the EFF stack and should be set to H5P_DEFAULT.

`rcntxt_id` indicates the read context for this operation.

The `es_id` parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in H5_EVENT_STACK_NULL for the `es_id` parameter.

Parameters:

`hid_t loc_id`

IN: Location identifier

May be any HDF5 object identifier (group, dataset, map, or named datatype) or file identifier that is in scope for the read context

<code>const char *name</code>	IN: Object name The object name (path to the object) can be specified relative to <code>loc_id</code> , absolute from the file's root group, or '.' (a dot), and must be in scope for the read context.
<code>char *comment</code>	OUT: The comment.
<code>size_t bufsize</code>	IN: Anticipated required size of the <code>comment</code> buffer.
<code>hid_t lapl_id</code>	IN: Link access property list identifier. <i>Currently not used in EFF; specify H5P_DEFAULT.</i>
<code>ssize_t *comment_size</code>	OUT: Pointer to returned results indicating comment size. When successful, will contain the number of characters in the comment, not including the NULL terminator, or zero (0) if the comment has no comment. On failure, will contain a negative value.
<code>hid_t rcntxt_id</code>	IN: Read context identifier indicating the read context for this operation.
<code>hid_t es_id</code>	IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously. Use <code>H5_EVENT_STACK_NULL</code> for synchronous execution.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

Note that when this routine is executed asynchronously, the return value from the routine only indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack. When the asynchronous execution completes successfully, `comment_size` will contain the number of characters in the comment, not including the NULL terminator, or zero (0) if the comment has no comment. On failure, `comment_size` will be a negative number.

History:

Added in Quarter 4.

Quarter 5: Changed from transaction to read context id and from event queue to event stack. Added scope requirement and note about link access property list under EFF.

Man Page Status:

No known issues.

4.10.5 H5Olink_ff

Name: H5Olink_ff

Signature:

```
herr_t H5Olink_ff(hid_t object_id, hid_t new_loc_id, const char *new_link_name, hid_t  
lcp1, hid_t lap1, hid_t trans_id, hid_t es_id )
```

Purpose:

Create a hard link to an object in an HDF5 file, possibly asynchronously.

Description:

H5Olink_ff creates a new hard link to an object in an HDF5 file.

`new_loc_id` and `new_link_name` specify the location and name of the new link, while `object_id` identifies the object that the link points to. `new_loc_id` and `object_id` must both be in scope for the transaction identified by `trans_id`.

H5Olink_ff is designed for two purposes:

- To create the first hard link to an object that has just been created with one of the H5*create_anon functions or with H5Tcommit_anon. (The prototype EFF stack does not support the _anon functions.)
- To add additional structure to an existing file so that, for example, an object can be shared among multiple groups.

`lcp1` and `lap1` are the link creation and access property lists associated with the new link. Neither of these property lists currently have any effect in the EFF stack and both should be set to H5P_DEFAULT.

`trans_id` indicates the transaction this operation is part of.

The `es_id` parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in H5_EVENT_STACK_NULL for the `es_id` parameter.

Parameters:

<code>hid_t object_id</code>	IN: Object to be linked. May be any HDF5 object identifier (group, dataset, map, or named datatype) that is in scope for the transaction.
<code>hid_t new_loc_id</code>	IN: File or group identifier specifying location at which object is to be linked. Must be in scope for the transaction.
<code>const char *new_link_name</code>	IN: Name of link to be created, relative to <code>new_loc_id</code> .
<code>hid_t lcp1_id</code>	IN: Link creation property list identifier. <i>Currently not used in EFF; specify H5P_DEFAULT.</i>
<code>hid_t lap1_id</code>	IN: Link access property list identifier. <i>Currently not used in EFF; specify H5P_DEFAULT.</i>
<code>hid_t trans_id</code>	IN: Transaction identifier specifying the transaction this operation is a part of.
<code>hid_t es_id</code>	IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function

Use or disclosure of data contained on this sheet is subject to the restriction on the title page of this document.
Copyright © The HDF Group, 2014. All rights reserved

call when executed asynchronously. Use H5_EVENT_STACK_NULL for synchronous execution.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

When this routine is executed asynchronously, the return value indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack

History:

Added in Quarter 4.

Quarter 5: Changed from transaction to transaction id and from event queue to event stack. Added scope requirement. Noted what is not supported in EFF stack and removed example that demonstrated functionality that isn't available in EFF.

Man Page Status:

No known issues.

4.10.6 H5Oopen_ff

Name: H5Oopen_ff

Signature:

```
hid_t H5Oopen_ff( hid_t loc_id, const char *name, hid_t lapl_id, hid_t rcntxt_id )
```

Purpose:

Open an existing object.

Description:

H5Oopen_ff opens the existing HDF5 object specified by loc_id and name.

loc_id is a location identifier and may be any HDF5 object identifier (group, dataset, map, or named datatype) or file identifier.

name is the path to the object relative to loc_id. If loc_id fully specifies the object that is to be opened, name should be '.' (a dot).

Both loc_id and name must be in scope for the read context identified by rcntxt_id.

The link access property list, lapl_id, currently has no effect in the EFF stack and should be set to H5P_DEFAULT.

rcntxt_id indicates the read context for this operation.

There is no asynchronous execution option for this function because the object type must be known before the object id (even a future object id) can be returned. Since the object type is not known until the operation completes, executing the operation asynchronously is not possible.

To conserve and release resources, the object should be closed with H5Oclose_ff when access is no longer required.

Parameters:

hid_t loc_id

IN: Location identifier

May be any HDF5 object identifier (group, dataset, map, or named datatype) or file identifier that is in scope for the read context.

const char *name

IN: Object name

The object name (path to the map object) can be specified relative to loc_id or absolute from the file's root group, and must be in scope for the read context.

hid_t lapl_id

IN: Link object access property list

Currently not used in EFF; specify H5P_DEFAULT.

hid_t rcntxt_id

IN: Read context identifier indicating the read context for this operation.

Returns:

Returns an object identifier if successful; otherwise returns a negative value.

History:

Quarter 4: Extended H5Open to support EFF stack.

Quarter 5: New signature, H5Oopen_ff, to include the read context. Man page created.

Man Page Status:

No known issues.

Use or disclosure of data contained on this sheet is subject to the restriction on the title page of this document.

Copyright © The HDF Group, 2014. All rights reserved

4.10.7 H5Oset_comment_ff

Name: H5Oset_comment_ff

Signature:

```
herr_t H5Oset_comment_ff(hid_t object_id, const char *comment, hid_t trans_id, hid_t es_id)
```

Purpose:

Set comment for specified object, possibly asynchronously.

Deprecated Function:

This function is deprecated in favor of object attributes.

Description:

H5Oset_comment_ff sets the comment for the specified object to the contents of `comment`. Any previously existing comment is overwritten. Comments should be relatively short, null-terminated, ASCII strings.

The target object is specified by an identifier, `object_id`, which must be in scope for the transaction indicated by `trans_id`.

If `comment` is the empty string or a null pointer, any existing comment message is removed from the object.

`trans_id` indicates the transaction this operation is part of.

The `es_id` parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in `H5_EVENT_STACK_NULL` for the `es_id` parameter.

Comments can be attached to any object that has an object header. Datasets, groups, maps, and named datatypes have object headers. Symbolic links do not have object headers. In the EFF stack, comments cannot be attached to attributes.

Parameters:

`hid_t object_id` IN: Identifier of the target object.

`const char *comment` IN: The new comment.

`hid_t trans_id` IN: Transaction identifier specifying the transaction this operation is a part of.

`hid_t eq_id` IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously. Use `H5_EVENT_STACK_NULL` for synchronous execution.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

When this routine is executed asynchronously, the return value indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

History:

Added in Quarter 4.

Quarter 5: Changed from transaction to transaction id and from event queue to event stack. Noted that you can't add comments to attributes in EFF.

Man Page Status:

Use or disclosure of data contained on this sheet is subject to the restriction on the title page of this document.

Copyright © The HDF Group, 2014. All rights reserved

No known issues.

Use or disclosure of data contained on this sheet is subject to the restriction on the title page of this document.
Copyright © The HDF Group, 2014. All rights reserved

4.10.8 H5Oset_comment_by_name_ff

Name: H5Oset_comment_by_name_ff

Signature:

```
herr_t H5Oset_comment_by_name_ff(hid_t loc_id, const char *name, const char *comment,
hid_t lapl_id, uint64_t trans, hid_t eq_id)
```

Purpose:

Set comment for specified object, possibly asynchronously.

Deprecated Function:

This function is deprecated in favor of object attributes.

Description:

H5Oset_comment_by_name_ff sets the comment for the specified object to the contents of `comment`. Any previously existing comment is overwritten. Comments should be relatively short, null-terminated, ASCII strings.

The target object is specified by `loc_id` and `name`. `loc_id` can specify any object in the file. `name` can be one of the following:

- The name of the object relative to `loc_id`
- An absolute name of the object, starting from /, the file's root group
- A dot (.), if `loc_id` fully specifies the object

Both `loc_id` and `name` must be in scope for the read context specified by `rcntxt_id`.

If `comment` is the empty string or a null pointer, any existing comment message is removed from the object.

`trans_id` indicates the transaction this operation is part of.

The `es_id` parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in `H5_EVENT_STACK_NULL` for the `es_id` parameter.

Comments can be attached to any object that has an object header. Datasets, groups, maps, and named datatypes have object headers. Symbolic links do not have object headers.

Comments can be attached to any object that has an object header. Datasets, groups, maps, and named datatypes have object headers. Symbolic links do not have object headers. In the EFF stack, comments cannot be attached to attributes. In the EFF stack, comments cannot be attached to attributes.

`lapl_id` contains a link access property list identifier. A link access property list can come into play when traversing links to access an object. The link access property list currently has no effect in the EFF stack and should be set to `H5P_DEFAULT`.

Parameters:

`hid_t loc_id` IN: Identifier of a file, group, dataset, map, or named datatype. Must be in scope for the transaction.

`const char *name` IN: Name of the object whose comment is to be set or reset, specified as a path relative to `loc_id`. `name` can be '.' (a dot) if `loc_id` fully specifies the object for which the comment is to be set. Must be in scope for the transaction.

<i>const char *comment</i>	IN: The new comment.
<i>hid_t lapl_id</i>	IN: Link access property list identifier. <i>Currently not used in EFF; specify H5P_DEFAULT.</i>
<i>hid_t trans_id</i>	IN: Transaction identifier specifying the transaction this operation is a part of.
<i>hid_t es_id</i>	IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously. Use H5_EVENT_STACK_NULL for synchronous execution.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

When this routine is executed asynchronously, the return value indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

History:

Added in Quarter 4.

Quarter 5: Changed from transaction to transaction id and from event queue to event stack. Added scope requirement and noted that link access property lists not supported in EFF. Noted that you can't add comments to attributes in EFF.

Man Page Status:

No known issues.

4.11 H5P: Property APIs

These routines allow applications to set values in HDF5 Property Lists. The routines appear in the table below grouped by the property list they apply to; the reference manual pages are organized alphabetically.

These routines did not exist prior to the Exascale FastForward project.

Routine	Implemented	Notes
<i>Data Transfer</i>		
H5Pset_dxpl_checksum	Quarter 4	
H5Pset_dxpl_checksum_ptr	Quarter 4	
H5Pset_dxpl_inject_corruption	Quarter 4	
H5Pset_rawdata_integrity_scope	Quarter 5	H5Pget_rawdata_integrity_scope also exists but has not yet been documented.
H5Pset_read_replica	Quarter 7	H5Pget_read_replica also exists but has not yet been documented.
<i>Object Evict from Burst Buffer</i>		
H5Pset_evict_replica	Quarter 7	H5Pget_evict_replica also exists but has not yet been documented.
<i>File Access</i>		
H5Pset_fapl_iod	Quarter 3	
H5Pset_metadata_integrity_scope	Quarter 5	H5Pget_metadata_integrity_scope also exists but has not yet been documented.
<i>Read Context Acquire</i>		
H5Pget_rcapl_version_request	Quarter 5	
H5Pset_rcapl_version_request	Quarter 5	
<i>Transaction Start</i>		
H5Pget_trspl_num_peers	Quarter 5	
H5Pset_trspl_num_peers	Quarter 5	

4.11.1 H5Pget_rcapl_version_request

Name: H5Pget_rcapl_version_request

Signature:

```
herr_t H5Pget_rcapl_version_request(hid_t rcapl_id, H5RC_request_t *acquire_req)
```

Purpose:

Retrieves a version request modifier from a read context acquire property list.

Description:

H5Pget_rcapl_version_request gets a property from the rcapl_id read context acquire property list and puts it in *acquire_req. The property is a container version request modifier for a read handle acquired with a call to H5RCacquire using rcapl_id.

Parameters:

<i>hid_t</i> rcapl_id	IN: Read context acquire property list identifier
<i>H5RC_request_t</i> *acquire_req	OUT: Modifier for container version specified in H5RCacquire. Valid values include: H5RC_EXACT Acquire a read handle for the exact container version specified. (Default) H5RC_PREV Acquire a read handle for the container version specified, or the highest previous version if the specified version is not available. H5RC_NEXT Acquire a read handle for the container version specified, or the lowest next version if the specified version is not available. H5RC_LAST Acquire a read handle for the last (highest) container version possible. With this setting, the input value pointed to by container_version is ignored.

Returns:

Returns a non-negative value if successful. Otherwise returns a negative value.

History:

Added in Quarter 5.

Man Page Status:

No known issues.

4.11.2 H5Pget_trspl_num_peers

Name: H5Pget_trspl_num_peers

Signature:

*hid_t H5Pget_trspl_num_peers(hid_t trspl_id, unsigned *num_peers)*

Purpose:

Retrieves the leader count property from a transaction start property list.

Description:

H5Pget_trspl_version_request gets a property from the `trspl_id` transaction start property list and puts it in `*num_peers`. The property is the number of leaders who will call `H5TRstart` for a given transaction using `trspl_id`.

Parameters:

<i>hid_t trspl_id</i>	IN: Transaction start property list identifier
<i>unsigned *num_peers</i>	OUT: Number of leader processes that will call <code>H5TRstart</code> for a given transaction.

Returns:

Returns a non-negative value if successful. Otherwise returns a negative value.

History:

Added in Quarter 5.

Man Page Status:

No known issues.

Later: Consider renaming function from ...num_peers to ...leader_count since terminology has changed re: transaction leaders. Also, confirm unsigned type –vs- unsigned w/ explicit length of bytes.

4.11.3 H5Pset_dxpl_checksum

Name: H5Pset_dxpl_checksum

Signature:

herr_t H5Pset_dxpl_checksum(hid_t dxpl_id, uint32_t value)

Purpose:

Specify a user-supplied checksum for a write data transfer.

Description:

H5Pset_dxpl_checksum sets a property in the dxpl_id data transfer property list specifying value as a user-supplied checksum for data written with a call to H5Dwrite_ff or H5Mset using dxpl_id.

When this is set, the HDF5 IOD VOL client will create a checksum for the data and verify that checksum matches the value supplied by the user before sending the data on to the HDF5 VOL IOD server.

Regardless of whether the user supplies a checksum value, the HDF5 IOD VOL client will create a checksum that is compared with the value generated on the HDF5 IOD VOL server.

The user must call H5checksum to obtain the value used in the call to H5Pset_dxpl_checksum. This ensures that all levels of the stack are using compatible checksum algorithms.

See *Design and Implementation of FastForward Features in HDF5* for information on the interaction between transformation operations, such as datatype conversion, and end-to-end integrity support with checksums.

Parameters:

hid_t dxpl_id IN: Data transfer property list identifier

uint32_t value IN: User-supplied checksum obtained via call to H5checksum

Returns:

Returns a non-negative value if successful. Otherwise returns a negative value.

History:

Added in Quarter 4.

Quarter 5: Added H5Mset to routines that are affected by the dxpl setting. Added reference for interaction between transformations and checksums.

Man Page Status:

No known issues.

4.11.4 H5Pset_dxpl_checksum_ptr

Name: H5Pset_dxpl_checksum_ptr

Signature:

*herr_t H5Pset_dxpl_checksum_ptr(hid_t dxpl_id, uint32_t *value)*

Purpose:

Specify a memory location to receive the checksum from a read data transfer.

Description:

H5Pset_dxpl_checksum_ptr sets a property in the dxpl_id data transfer property list specifying value as a memory location to receive a checksum for data read with a call to H5Dread_ff or H5Mget using dxpl_id.

When this is set, the HDF5 IOD VOL client will put the checksum for the data into the memory location (*value) supplied by the user, allowing the user to compare the client's checksum to a value it generates using the H5checksum routine.

Regardless of whether the user supplies a memory location for the checksum, the HDF5 IOD VOL client will create a checksum and compare it with the value generated on the HDF5 IOD VOL server before the operation completes.

The user must call H5checksum to obtain the value for its comparison with the received value. This ensures that all levels of the stack are using compatible checksum algorithms.

See *Design and Implementation of FastForward Features in HDF5* for information on the interaction between transformation operations, such as datatype conversion, and end-to-end integrity support with checksums.

Parameters:

hid_t dxpl_id IN: Data transfer property list identifier

*uint32_t *value* OUT: User-supplied memory location to receive checksum value

Returns:

Returns a non-negative value if successful. Otherwise returns a negative value.

History:

Added in Quarter 4.

Quarter 5: Added H5Mget to routines that are affected by the dxpl setting. Added reference for interaction between transformations and checksums.

Man Page Status:

No known issues.

4.11.5 H5Pset_dxpl_inject_corruption

Name: H5Pset_dxpl_inject_corruption

Signature:

herr_t H5Pset_dxpl_inject_corruption(hid_t dxpl_id, hbool_t flag)

Purpose:

Specify that data should be corrupted prior to transfer so that data integrity pipeline can be tested. This routine is for testing purposes only and should not be used in a real application.

Description:

H5Pset_dxpl_inject_corruption sets a property in the `dxpl_id` data transfer property list that can tell the HDF5 IOD VOL client (for writes) and HDF5 IOD VOL server (for reads) to corrupt the data being transferred with `H5Dwrite_ff` or `H5Dread_ff` using `dxpl_id`. Corruption will occur when `flag` is TRUE.

Parameters:

hid_t dxpl_id IN: Data transfer property list identifier

hbool_t flag IN: Boolean (TRUE/FALSE), indicating if corruption should be injected. TRUE means inject corruption.

Returns:

Returns a non-negative value if successful. Otherwise returns a negative value.

History:

Added in Quarter 4.

Man Page Status:

No known issues.

4.11.6 H5Pset_read_replica

Name: H5Pset_read_replica

Signature:

herr_t H5Pset_read_replica(hid_t dxpl_id, hrpl_t replica_id)

Purpose:

Specify a property that identifies the object replica to use when reading data.

Description:

H5Pset_read_replica sets the read replica property in the `dxpl_id` data transfer property list, specifying an object replica identifier.

When this property is set, read operations on the object will first look for requested data in the object replica, which resides in the Burst Buffer, before looking for the data in transaction updates or on persistent store (DAOS). If the requested data is not found in the specified replica, it will be retrieved from DAOS or transaction update data.

Replica IDs are returned by the `H5Dprefetch_ff`, `H5Gprefetch_ff`, `H5Mprefetch_ff`, and `H5Tprefetch_ff` routines.

The read replica property is currently recognized and processed by `H5Dread_ff` and `H5Mget_ff`.

Limitations / Future Considerations:

In a production version, additional routines that access data should also recognize and process the read replica property.

Parameters:

<i>hid_t dxpl_id</i>	IN: Data transfer property list identifier
<i>hrpl_t replica_id</i>	IN: Identifier for an object replica that was previously prefetched into the Burst Buffer via a call to <code>H5?prefetch_ff</code> , where ? may be “D”, “M”, “G”, or “T”, depending on the object type that was prefetched.

Returns:

Returns a non-negative value if successful. Otherwise returns a negative value.

History:

Added in Quarter 7.

Man Page Status:

No known issues.

4.11.7 H5Pset_evict_replica

Name: H5Pset_evict_replica

Signature:

herr_t H5Pset_evict_replica(hid_t lapl_id, hrpl_t replica_id)

Purpose:

Specify a property that identifies the object replica to evict.

Description:

H5Pset_evict_replica sets the evict replica property in the `lapl_id` access property list, specifying an object replica identifier.

When this property is set, evict operations on the object will evict the specified replica rather than the transaction update data for the object.

Replica IDs are returned by the `H5Dprefetch_ff`, `H5Gprefetch_ff`, `H5Mprefetch_ff`, and `H5Tprefetch_ff` routines.

Evict replica properties are recognized by `H5Devict_ff`, `H5Mevict_ff`, `H5Gevict_ff`, and `H5Tevict_ff`.

Limitations / Future Considerations:

For other potential extensions that are beyond the scope of the EFF prototype project, refer to the document *Burst Buffer Space Management – Prototype to Production*.

Parameters:

hid_t lapl_id IN: Access property list identifier

hrpl_t replica_id IN: Identifier for an object replica that was previously prefetched into the Burst Buffer via a call to `H5?prefetch_ff`, where ? may be “D”, “M”, “G”, or “T”, depending on the object type that was prefetched.

Returns:

Returns a non-negative value if successful. Otherwise returns a negative value.

History:

Added in Quarter 7.

Man Page Status:

No known issues.

4.11.8 H5Pset_fapl_iod

Name: H5Pset_fapl_iod

Signature:

herr_t H5Pset_fapl_iod(hid_t fapl_id, MPI_Comm comm, MPI_Info info)

Purpose:

Specify that the IOD VOL plugin should be used to access the HDF5 container.

Description:

H5Pset_fapl_iod sets a property in the `fapl_id` file access property list specifying that the IOD VOL plugin should be used to perform I/O. The IOD VOL plugin is used to interact with the Exascale Fast Forward I/O stack, where HDF5 files are stored as IOD and DAOS containers and objects rather than in the native binary HDF5 File Format or another storage/access mechanism.

Calling this routine is *mandatory* to use the HDF5 FastForward (H5*_ff) APIs described in this document.

H5Pset_fapl_iod also stores the user-supplied MPI parameters `comm`, for communicator, and `info`, for information, in the file access property list `fapl_id`. That property list can then be used to create and/or open a file. The communicator and info parameters are used to set up communication channels for collective operations on the HDF5 container.

`comm` is the MPI communicator to be used for accessing the file. This function makes a duplicate of the communicator, so modifications to `comm` after this function call returns have no effect on the file access property list.

`info` is the MPI Info object used to pass MPI info parameters to the MPI library. This function makes a duplicate copy of the Info object, so modifications to the Info object after this function call returns will have no effect on the file access property list.

If the file access property list already contains previously-set communicator and Info values, those values will be replaced and the old communicator and Info object will be freed.

Parameters:

`hid_t fapl_id` IN: File access property list identifier

`MPI_Comm comm` IN: MPI communicator

`MPI_Info info` IN: MPI info object

Returns:

Returns a non-negative value if successful. Otherwise returns a negative value.

History:

Added in Quarter 3.

Man Page Status:

No known issues.

4.11.9 H5Pset_metadata_integrity_scope

Name: H5Pset_metadata_integrity_scope

Signature:

```
herr_t H5Pset_metadata_integrity_scope(hid_t fapl_id, uint32_t scope )
```

Purpose:

Specify the scope of checksum generation and verification for metadata transfer.

Description:

H5Pset_metadata_integrity_scope sets a property in the `fapl_id` file access property list specifying the scope of checksum generation and verification for metadata transfer in the Exascale Fast Forward stack. The specified scope will be used in all metadata transfer operations for containers opened with `fapl_id`. Note that different scopes may be specified for different containers, or for different opens of the same container.

The `scope` is a bitflag parameter that specifies which stages of the metadata transfer process will have checksums generated and verified. The possible values are:

<code>H5_CHECKSUM_NONE</code>	No checksums are computed or verified at any part of the stack.
<code>H5_CHECKSUM_TRANSFER</code>	Metadata is verified after transfer through mercury (the function shipper).
<code>H5_CHECKSUM_IOD</code>	Checksums are computed for the metadata, given to IOD when written, and verified when read.
<code>H5_CHECKSUM_MEMORY</code>	Metadata is verified when moved into memory.
<code>H5_CHECKSUM_ALL</code>	Metadata is checksummed and verified at all levels.

Multiple individual selections can be combined using the OR operator. For example:

```
status = H5_CHECKSUM_IOD || H5_CHECKSUM_MEMORY
```

Parameters:

`hid_t fapl_id` IN: File access property list identifier

`uint32_t scope` IN: Bitfield specifying the scope of checksums for raw data transfers

Returns:

Returns a non-negative value if successful. Otherwise returns a negative value.

History:

Added in Quarter 5.

Man Page Status:

To do: Update with description of what is considered metadata. Improve descriptions of scope options, maybe adding an ASCII picture of data flow.

4.11.10 H5Pset_rawdata_integrity_scope

Name: H5Pset_rawdata_integrity_scope

Signature:

herr_t H5Pset_rawdata_integrity_scope(hid_t dxpl_id, uint32_t scope)

Purpose:

Specify the scope of checksum generation and verification for raw data transfer.

Description:

H5Pset_rawdata_integrity_scope sets a property in the `dxpl_id` data transfer property list specifying the scope of checksum generation and verification for raw data transfer in the Exascale Fast Forward stack. The specified scope will be used in all data transfer operations that are called with `dxpl_id`. In particular, `H5Dwrite_ff`, `H5Mset`, `H5Dread_ff`, and `H5Mget`. Note that different scopes may be specified for each data transfer operations.

The `scope` is a bitflag parameter that specifies which stages of the raw data transfer process will have checksums generated and verified. The possible values are:

<code>H5_CHECKSUM_NONE</code>	No checksums are computed or verified at any part of the stack.
<code>H5_CHECKSUM_TRANSFER</code>	Raw data is verified after transfer through mercury (the function shipper).
<code>H5_CHECKSUM_IOD</code>	Checksums are computed for raw data, given to IOD when written, and verified when read.
<code>H5_CHECKSUM_MEMORY</code>	Raw data is verified when moved into memory.
<code>H5_CHECKSUM_ALL</code>	Raw data is checksummed and verified at all levels.

Multiple individual selections can be combined using the OR operator. For example:

```
status = H5_CHECKSUM_IOD || H5_CHECKSUM_MEMORY
```

Parameters:

`hid_t dxpl_id` IN: Data transfer property list identifier

`uint32_t scope` IN: Bitfield specifying the scope of checksums for raw data transfers

Returns:

Returns a non-negative value if successful. Otherwise returns a negative value.

History:

Added in Quarter 5.

Man Page Status:

To do: Update with list of functions that are affected by the property's setting. Improve descriptions of scope options, maybe adding an ASCII picture of data flow.

4.11.11 H5Pset_rcapl_version_request

Name: H5Pset_rcapl_version_request

Signature:

herr_t H5Pset_rcapl_version_request(hid_t rcapl_id, H5RC_request_t acquire_req)

Purpose:

Specify a version request modifier in a read context acquire property list.

Description:

H5Pset_rcapl_version_request sets a property in the `rcapl_id` read context acquire property list specifying `acquire_req` as a container version request modifier for a read handle acquired with a call to H5RCacquire using `rcapl_id`.

When this is set, the container version specified in the call to H5RCacquire will be interpreted based on the value of the property. The possible values and interpretations are:

H5RC_EXACT	Acquire a read handle for the exact container version specified. (Default)
H5RC_PREV	Acquire a read handle for the container version specified, or the highest previous version if the specified version is not available.
H5RC_NEXT	Acquire a read handle for the container version specified, or the lowest next version if the specified version is not available. With this setting, and <code>*container_version</code> set to 0, the handle will be acquired for the first (lowest) readable version of the container.
H5RC_LAST	Acquire a read handle for the last (highest) container version possible. With this setting, the input value pointed to by <code>container_version</code> is ignored.

H5Pcreate(H5P_RC_ACQUIRE) will create a read context acquire property list whose identifier can be passed to this function.

Parameters:

<i>hid_t rcapl_id</i>	IN: Read context acquire property list identifier
<i>H5RC_request_t acquire_req</i>	IN: Modifier for container version specified in H5RCacquire

Returns:

Returns a non-negative value if successful. Otherwise returns a negative value.

History:

Added in Quarter 5.

Man Page Status:

No known issues.

4.11.12 H5Pset_trspl_num_peers

Name: H5Pset_trspl_num_peers

Signature:

herr_t H5Pset_trspl_num_peers(hid_t trspl_id, unsigned num_peers)

Purpose:

Set the leader count property in a transaction start property list.

Description:

H5Pset_trspl_version_request sets a property in the `trspl_id` transaction start property list to `num_peers`. The property is the number of leaders who will call `H5TRstart` for a given transaction using `trspl_id`.

`H5Pcreate(H5P_TR_START)` will create a transaction start property list whose identifier can be passed to this function.

Parameters:

`hid_t trspl_id` IN: Transaction start property list identifier

`unsigned num_peers` IN: Number of leader processes that will call `H5TRstart` for a given transaction.

Returns:

Returns a non-negative value if successful. Otherwise returns a negative value.

History:

Added in Quarter 5.

Man Page Status:

No known issues.

Later: Consider renaming function from ...num_peers to ...leader_count since terminology has changed re: transaction leaders. Also, confirm unsigned type –vs- unsigned w/ explicit length of bytes

4.12 H5Q: Query APIs

These routines are used to specify a query for analysis shipping operations on HDF5 Files (containers) in the Exascale Fast Forward stack.

Query objects specify constraints on HDF5 objects that restrict an analysis shipping operation to certain objects or portions of objects in an HDF5 File. As of the deliverables for Q6, query operations can only be performed on HDF5 dataset objects, but future deliverables will expand query operation capabilities to allow queries to be specified on attribute and group objects as well.

These routines did not exist prior to the Exascale FastForward project.

Routine	Implemented	Notes
H5Qcreate	Quarter 6	
H5Qcombine	Quarter 6	
H5Qclose	Quarter 6	

4.12.1 H5Qcreate

Name: H5Qcreate

Signature:

```
hid_t H5Qcreate ( H5Q_type_t query_type, H5Q_match_op_t match_op, ... )
```

Purpose:

Create a new query object, specifying a type of object to apply the query to and the condition for matching.

Description:

`H5Qcreate` creates a new query object, and specifies the type of object in the HDF5 file that the query applies to, along with the condition of matching and the match value. Query objects created with this routine can be used for analysis shipping operations (with `H5ASexecute`) or combined into more complicated queries (with `H5Qcombine`).

The `query_type` parameter indicates the type of HDF5 object that the query will apply to. This parameter will determine the interpretation of the varargs parameters. The possible values and meanings are:

<code>H5Q_TYPE_DATA_ELEM</code>	Query will be applied to dataset elements.
<code>H5Q_TYPE_ATTR_NAME</code>	Query will be applied to attribute names. (Not yet supported)
<code>H5Q_TYPE_LINK_NAME</code>	Query will be applied to link names. (Not yet supported)

The `match_op` parameter indicates the conditions for query matching. The possible values and meanings are:

<code>H5Q_MATCH_EQUAL</code>	Match values for the HDF5 object that are equal to the value specified for the query.
<code>H5Q_MATCH_NOT_EQUAL</code>	Match values for the HDF5 object that are not equal to the value specified for the query.
<code>H5Q_MATCH_LESS_THAN</code>	Match values for the HDF5 object that are less than the value specified for the query.
<code>H5Q_MATCH_GREATER_THAN</code>	Match values for the HDF5 object that are greater than the value specified for the query.

The values for the varargs parameter (...) are interpreted according to the `query_type` parameter. The possible types and meanings for each value of `query_type` are:

<code>H5Q_TYPE_DATA_ELEM</code>	<code>hid_t value_datatype_id, value_datatype_id</code> specifies the datatype of the query value and <code>value</code> is a pointer to the query value itself.
<code>H5Q_TYPE_ATTR_NAME</code>	<code>const char * attr_name</code> <code>attr_name</code> is a pointer to the query value.
<code>H5Q_TYPE_LINK_NAME</code>	<code>const char * link_name</code> <code>link_name</code> is a pointer to the query

value.

For example, to create a query that specifies matching on dataset elements that are greater than 53.4, the following example code could be used:

```
float f = 53.4;  
hid_t query_id = H5Qcreate(H5Q_TYPE_DATA_ELEM, H5Q_MATCH_GREATER_THAN,  
    H5T_NATIVE_FLOAT, &f);
```

This routine executes immediately and locally, without need for asynchronous operations. Query objects created by this routine must be closed by H5Qclose.

Parameters:

<i>H5Q_type_t</i> <code>query_type</code>	IN: Specifies the type of HDF5 object that the query will apply to. This parameter will determine interpretation of the varargs parameters.
<i>H5Q_match_op_t</i> <code>match_op</code>	IN: Specifies the conditions for matching a query value.
...	IN: The varargs parameters are interpreted according to the query type.

Returns:

Returns a query identifier if successful; otherwise returns a negative value. Query identifiers returned from this routine must be released with the H5Qclose routine.

History:

Added in Quarter 6.

Man Page Status:

No known issues.

4.12.2 H5Qcombine

Name: H5Qcombine

Signature:

```
hid_t H5Qcombine ( hid_t query_id_1, H5Q_combine_op_t combine_op, hid_t query_id_2)
```

Purpose:

Create a new query object by combining two existing query objects.

Description:

H5Qcombine creates a new query object by combining the conditions specified for two existing query objects. Query objects created with this routine can be used for analysis shipping operations (with H5ASexecute) or combined again into more complicated queries (with H5Qcombine).

The `query_id_1` and `query_id_2` parameters indicate the two query objects to be combined with the `combine_op` parameter.

The `combine_op` parameter indicates how the two query objects are combined. The possible values and meanings are:

`H5Q_COMBINE_AND` Both of the conditions specified in the `query_id_1` and `query_id_2` parameters must be fulfilled for the resulting query.

`H5Q_COMBINE_OR` Either, or both, of the conditions specified in the `query_id_1` and `query_id_2` parameters must be fulfilled for the resulting query.

For example, to create a query that specifies matching on dataset elements that are greater than 53.4 and less than 93.6, the following example code could be used:

```
float f1 = 53.4;
hid_t query_id_1 = H5Qcreate(H5Q_TYPE_DATA_ELEM,
    H5Q_MATCH_GREATER_THAN, H5T_NATIVE_FLOAT, &f1);
float f2 = 93.6;
hid_t query_id_2 = H5Qcreate(H5Q_TYPE_DATA_ELEM,
    H5Q_MATCH_LESS_THAN, H5T_NATIVE_FLOAT, &f2);
hid_t query_id_3 = H5Qcombine(query_id_1, H5Q_COMBINE_AND, query_id_2);
```

This routine executes immediately and locally, without need for asynchronous operations. Query objects created by this routine must be closed by H5Qclose.

Parameters:

`hid_t query_id_1` IN: Specifies one of the queries to be combined.

`H5Q_combine_op_t combine_op` IN: Specifies the how the two queries are to be combined.

`hid_t query_id_2` IN: Specifies one of the queries to be combined.

Returns:

Returns a query identifier if successful; otherwise returns a negative value. Query identifiers returned from this routine must be released with the H5Qclose routine.

History:

Added in Quarter 6.

Man Page Status:

No known issues.

4.12.3 H5Qclose

Name: H5Qclose

Signature:

herr_t H5Qclose (hid_t query_id)

Purpose:

Release the resources for a query object.

Description:

`H5Qclose` closes the query specified by `query_id` and releases resources used by it. Further use of the query identifier is illegal in HDF5 API calls.

The actions performed by this function are all local to the compute node so an asynchronous execution option is not provided.

Parameters:

`hid_t query_id` IN: Identifier of the query to close.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

History:

Added in Quarter 6.

Man Page Status:

No known issues.

4.13 H5RC: Read Context APIs

These routines are used to establish the read context for operations on HDF5 Files (containers) in the Exascale Fast Forward stack.

In the Exascale Fast Forward stack it is possible that multiple versions of a container may be readable at any given time, where a container version results when a transaction is committed, and earlier versions may be 'flattened' by later versions as the data is reorganized for optimized access.

The application must indicate what container version it wants to read from, not simply what container. A read context is used to open a handle on a particular version, guaranteeing that the version will remain available until the context is closed. A read context ID is passed to HDF5 operations to indicate what version of the container should be read from. All reads made using the same read context will see a consistent version of the container data. Multiple read contexts can be open on a single container at any given time.

These routines did not exist prior to the Exascale FastForward project.

Routine	Implemented	Notes
H5RCacquire	Quarter 5	
H5RCacquire_wait		Added in response to Q5 feedback; May be implemented in future quarter if time allows.
H5RCclose	Quarter 5	
H5CCcreate	Quarter 5	
H5RCget_version	Quarter 5	
H5RCpersist	Quarter 5	
H5RCrelease	Quarter 5	
H5RCsnapshot	Quarter 5	

4.13.1 H5RCacquire

Name: H5RCacquire

Signature:

```
hid_t H5RCacquire ( hid_t file_id, uint64_t *container_version, hid_t rcapl_id,  
hid_t es_id )
```

Purpose:

Acquire a read handle for a container at a given version and create a read context associated with the container and version.

Description:

`H5RCacquire` acquires a read handle on a container version and creates a read context associated with that container version. The container version is guaranteed to remain readable and consistent until the context is closed and the handle released (with `H5RCrelease`).

The `file_id` parameter indicates the container the read context will be associated with.

The `container_version` parameter indicates the container version that the read handle will be acquired for and the read context associated with. Parameters in the read context acquire property list control how the value pointed to by `container_version` will be interpreted by the function, since it is not always possible to acquire a particular container version. Upon successful completion of the function, the value pointed to by the `container_version` parameter contains the actual version that the read handle was acquired for and the read context associated with.

The version request property in the read context acquire property list, `rcapl_id`, controls how the function interprets the value pointed to by `container_version`. `H5Pset_rcapl_version_request()` sets the version request property, whose possible values and meanings are:

<code>H5RC_EXACT</code>	Acquire a read handle for the exact container version specified. (Default)
<code>H5RC_PREV</code>	Acquire a read handle for the container version specified, or the highest previous version if the specified version is not available.
<code>H5RC_NEXT</code>	Acquire a read handle for the container version specified, or the lowest next version if the specified version is not available. With this setting, and <code>*container_version</code> set to 0, the handle will be acquired for the first (lowest) readable version of the container.
<code>H5RC_LAST</code>	Acquire a read handle for the last (highest) container version possible. With this setting, the input value pointed to by <code>container_version</code> is ignored.

The `es_id` parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in `H5_EVENT_STACK_NULL` for the `es_id` parameter, but the returned read context should not be used until the asynchronous function completes.

This function must be called by one process in a parallel application to ensure that the container version remains readable. That process can notify other processes in the application when the read handle is acquired, and the other processes can each call `H5RCcreate` to create local read context identifiers for the container version. It is not necessary for every process that will be doing reads on a given container version to call `H5RCacquire`.

Parameters:

Use or disclosure of data contained on this sheet is subject to the restriction on the title page of this document.
Copyright © The HDF Group, 2014. All rights reserved

<i>hid_t file_id</i>	IN: File identifier indicating the container that the read handle is being acquired for and read context associated with.
<i>uint64_t *container_version</i>	IN/OUT: On input, pointer to value that, in conjunction with the read context acquire property list, specifies the container version to open the handle for. Upon successful completion of the function, the value pointed to by this parameter will contain the actual container_version that the read handle was acquired. The value is unchanged if the function does not complete successfully.
<i>hid_t rcapl_id</i>	IN: Read context acquire property list. Can contain the version request property, set with <code>H5Pset_rcapl_version_request()</code> , that controls how <code>container_version</code> is interpreted.
<i>hid_t es_id</i>	IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously. Use <code>H5_EVENT_STACK_NULL</code> for synchronous execution.

Returns:

Returns a read context identifier if successful; otherwise returns a negative value. When executed asynchronously, a future ID for the new read context is returned initially. Upon completion of the asynchronous operation, the future ID will be transparently modified to be a “normal” read context identifier. Since the success of the read handle acquisition, and the actual container version acquired, is not known until completion, the application should not use the future ID for the read context in any other functions as a read context ID parameter.

When this routine is executed asynchronously, the return value indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

History:

Added in Quarter 5.

Man Page Status:

No known issues.

4.13.2 H5RCclose

Name: H5RCclose

Signature:

herr_t H5RCclose (hid_t rcntxt_id)

Purpose:

Close a read context.

Description:

`H5RCclose` closes the read context specified by `rcntxt_id` and releases resources used by it. Further use of the read context identifier is illegal in HDF5 API calls. `H5RCclose` does not release the read handle for the container version associated with the read context that is being closed. Typically the read context specified by `rcntxt_id` was obtained by a call to `H5RCcreate`, but in some circumstances `H5RCacquire` may have returned the read context.

`H5RCclose` should be called before the handle for the container version is released via a call to `H5RCrelease`. Typically the process that called `H5RCacquire` will call `H5RCrelease`, but in some circumstances, such as premature process termination, another process with a read context for the container version obtained via a call to `H5RCcreate` may make the call to `H5RCrelease`. Regardless of which process calls `H5RCrelease`, it closes the read context and releases the read handle on the associated container version.

For a given `rcntxt_id`, either `H5RCclose` or `H5RCrelease` should be called, but not both.

The actions performed by this function are all local to the compute node so an asynchronous execution option is not provided.

Parameters:

hid_t rcntxt_id IN: Identifier of the read context to close.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

History:

Added in Quarter 5.

Man Page Status:

No known issues.

4.13.3 H5RCcreate

Name: H5RCcreate

Signature:

```
hid_t H5RCcreate ( hid_t file_id, uint64_t container_version )
```

Purpose:

Create a read context associated with a container and version.

Description:

`H5RCcreate` creates a read context associated with a specified container and version after a read handle has already been acquired on the container version by another process in the application. The created read context is passed to other HDF5 functions that perform reads on the container.

The `file_id` parameter indicates the container the read context is associated with.

The `container_version` parameter indicates the container version that the read context is associated with.

Prior to this call, a read handle must be acquired on the container version by another process in the parallel application via a call to `H5RCacquire`. The container version is guaranteed to remain readable and consistent until the handle released (with `H5RCrelease`).

To conserve and release resources, the read context created by this function should be closed, when access is no longer required, usually with `H5RCclose`. The read context acquired by this function cannot be used after the handle for the container version has been released via a call to `H5RCrelease`. Typically the process that called `H5RCacquire` will call `H5RCrelease`, but in some circumstances (such as premature process termination), another process with a read context for the container version may make the call to `H5RCrelease`, closing the read context and releasing the read handle.

The actions performed by this function are all local to the compute node so an asynchronous execution option is not provided.

Parameters:

<code>hid_t file_id</code>	IN: File identifier indicating the container that the read context is associated with.
----------------------------	--

<code>uint64_t container_version</code>	IN: The container version that the read context is associated with..
---	--

Returns:

Returns a read context identifier if successful; otherwise returns a negative value.

History:

Added in Quarter 5.

Man Page Status:

No known issues.

4.13.4 H5RCget_version

Name: H5RCget_version

Signature:

```
herr_t H5RCget_version( hid_t rcntxt_id, uint64_t *container_version )
```

Purpose:

Retrieve the container version associated with a read context.

Description:

The `rcntxt_id` specifies the read context whose associated container version is to be retrieved.

The value pointed to by the `container_version` parameter will be set to the container version that the read context is associated with.

The actions performed by this function are all local to the compute node so an asynchronous execution option is not provided.

Parameters:

<code>hid_t rcntxt_id</code>	IN: Read context identifier whose associated container version is being retrieved.
<code>uint64_t *container_version</code>	OUT: Container version associated with the read context. Value is unchanged if operation is unsuccessful.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

History:

Added in Quarter 5.

Man Page Status:

No known issues.

4.13.5 H5RCpersist

Name: H5RCpersist

Signature:

herr_t H5RCpersist (hid_t rcntxt_id, hid_t es_id)

Purpose:

Copy data for a container from IOD to DAOS, bringing DAOS up to specified container version.

Description:

`H5RCpersist` requests that IOD update a container on DAOS to a specific version.

The `rcntxt_id` indicates both the container that is to be updated and the version that is to be persisted. The indicated container must be open for write.

Updates to the container between the last persisted version and the version currently being persisted will be copied to DAOS as part of this persist request.

The `es_id` parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in `H5_EVENT_STACK_NULL` for the `es_id` parameter.

Parameters:

<code>hid_t rcntxt_id</code>	IN: Read context identifier indicating the container (that is open for write) and the version that is to be persisted.
<code>hid_t es_id</code>	IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously. Use <code>H5_EVENT_STACK_NULL</code> for synchronous execution.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

When this routine is executed asynchronously, the return value indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

History:

Quarter 4: Documented as H5Fpersist but not yet implemented.

Quarter 5: Renamed from H5Fpersist to H5RCpersist, using `rcntxt_id` instead of `file_id` and `container_version`. Implemented.

Man Page Status:

No known issues.

4.13.6 H5RCrelease

Name: H5RCrelease

Signature:

herr_t H5RCrelease (hid_t rcntxt_id, hid_t es_id)

Purpose:

Close a specified read context and release a read handle for the associated container version

Description:

`H5RCrelease` closes the read context specified by `rcntxt_id` and releases resources used by it. It also releases the read handle for the container version associated with the read context. Further use of the read context identifier is illegal in HDF5 API calls. Typically the read context specified by `rcntxt_id` was obtained by a call to `H5RCacquire`, `H5Fopen_ff`, or `H5TRfinish`, but in some circumstances `H5RCcreate` may have returned the read context.

`H5RCrelease` should be called after all other read contexts for the container version have been released via calls to `H5RCclose`. Typically the process that called `H5RCacquire`, `H5Fopen_FF`, or `H5TRfinish` will call `H5RCrelease`, but in some circumstances, such as premature process termination, another process with a read context for the container version obtained via a call to `H5RCcreate` may make the call to `H5RCrelease`. Regardless of which process calls `H5RCrelease`, the read context is closed and the read handle on the associated container version is released.

For a given `rcntxt_id`, either `H5RCrelease` or `H5RCclose` should be called, but not both.

The `es_id` parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in `H5_EVENT_STACK_NULL` for the `es_id` parameter.

Parameters:

hid_t rcntxt_id IN: Identifier of the read context that will be closed; the read handle for the associated container version will also be released.

hid_t es_id IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously. Use `H5_EVENT_STACK_NULL` for synchronous execution.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

When this routine is executed asynchronously, the return value indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

History:

Added in Quarter 5.

Man Page Status:

No known issues.

4.13.7 H5RCsnapshot

Name: H5RCsnapshot

Signature:

```
herr_t H5RCsnapshot( hid_t rcntxt_id, const char* snapshot_name, hid_t es_id )
```

Purpose:

Make a snapshot of a container on DAOS.

Description:

`H5RCsnapshot` requests that DAOS make a copy of the container and version specified by the read context, and give it the indicated container name. The snapshot is a permanently visible copy of the container's state at the given container version.

The `rcntxt_id` indicates both the container and the container version that the new copy will duplicate. If the requested version is not readable on DAOS for the container, the snapshot request will fail.

`snapshot_name` is the name given to the newly created snapshot.

The `es_id` parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in `H5_EVENT_STACK_NULL` for the `es_id` parameter.

Parameters:

`hid_t rcntxt_id` IN: Read context identifier indicating the container and the version that is to be persisted for with the snapshot should be taken.

`const char* snapshot_name` IN: Name of the snapshot.

`hid_t es_id` IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously. Use `H5_EVENT_STACK_NULL` for synchronous execution.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

When this routine is executed asynchronously, the return value indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

History:

Quarter 4: Documented but not yet implemented.

Quarter 5: Renamed from `H5Fsnapshot` to `H5RCsnapshot`, using `rcntxt_id` instead of `file_id` and `container_version`. Implemented.

Man Page Status:

No known issues.

4.14 H5T: Datatype APIs

These routines are used to operate on HDF5 Datatype Objects.

The routines ending in _ff have different signatures than the standard HDF5 library routines.

Man pages for routines whose user interface is unchanged from the standard HDF5 implementation can be found at: http://www.hdfgroup.org/HDF5/doc/RM/RM_H5T.html.

Routine	Implemented	Notes
H5Tclose_ff	Quarter 4	
H5Tcommit_ff	Quarter 4	
H5Tevict_ff	Quarter 7	
H5Topen_ff	Quarter 4	This is implemented synchronously for now.
H5Tprefetch_ff	Quarter 7	
H5Tget_create_plist	Quarter 4	Local operation, no need to do asynchronously. See standard HDF5 man page.
H5Tcommit_anon		Will not implement

4.14.1 H5Tclose_ff

Name: H5Tclose_ff

Signature:

```
hid_t H5Tclose_ff( hid_t dtype_id, hid_t es_id )
```

Purpose:

Releases a datatype, possibly asynchronously.

Description:

H5Tclose_ff releases the datatype specified by `dtype_id`. Further access through the datatype identifier is illegal. Failure to release a datatype with this call will result in resource leaks.

The `es_id` parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in `H5_EVENT_STACK_NULL` for the `es_id` parameter.

Parameters:

`hid_t dtype_id` IN: Identifier of the datatype to release.

`hid_t es_id` IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously. Use `H5_EVENT_STACK_NULL` for synchronous execution.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

When this routine is executed asynchronously, the return value indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

History:

Added in Quarter 4.

Quarter 5: Changed from event queue to event stack.

Man Page Status:

No known issues.

4.14.2 H5Tcommit_ff

Name: H5Tcommit_ff

Signature:

```
herr_t H5Tcommit_ff(hid_t loc_id, const char *name, hid_t dtype_id, hid_t lcpl_id,  
hid_t tcpl_id, hid_t tapl_id, hid_t trans_id, hid_t es_id)
```

Purpose:

Commit a transient datatype, linking it into the file and creating a new named datatype, possibly asynchronously.

Description:

H5Tcommit_ff saves a transient datatype as an immutable named datatype in a file. The datatype specified by `dtype_id` is committed to the file with the name `name` at the location specified by `loc_id` and with the datatype creation and access property lists `tcpl_id` and `tapl_id`, respectively.

`loc_id` may be a file identifier, or a group identifier within that file. `loc_id` must be in scope for the transaction identified by `trans_id`.

`name` may be either an absolute path in the file or a relative path from `loc_id` naming the newly-committed datatype. `name` must be in scope for the transaction identified by `trans_id`.

The link creation property list, `lcpl_id`, governs creation of the link(s) by which the new named datatype is accessed and the creation of any intermediate groups that may be missing. The link creation property list currently has no effect in the EFF stack and should be set to `H5P_DEFAULT`.

The datatype creation property list, `tcpl_id`, and dataset access property list, `tapl_id`, are currently not used and should be set to `H5P_DEFAULT`.

`trans_id` indicates the transaction this operation is part of.

The `es_id` parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in `H5_EVENT_STACK_NULL` for the `es_id` parameter.

Once committed, this datatype may be used to define the datatype of any other dataset or attribute in the file.

This function will not accept a datatype that cannot actually hold data. This currently includes compound datatypes with no fields and enumerated datatypes with no members.

Parameters:

<code>hid_t loc_id</code>	IN: Location identifier May be any HDF5 group identifier or file identifier file identifier that is in scope for the transaction.
<code>const char *name</code>	IN: Name given to named datatype The name can be specified relative to <code>loc_id</code> , or absolute from the file's root group, and must be in scope for the transaction.
<code>hid_t dtype_id</code>	IN: Identifier of datatype to be committed and, upon function's return, identifier for the named datatype
<code>hid_t lcpl_id</code>	IN: Link creation property list

Currently not used in EFF; specify H5P_DEFAULT.

<i>hid_t tcpl_id</i>	IN: Datatype creation property list <i>Currently not used; specify H5P_DEFAULT.</i>
<i>hid_t tapl_id</i>	IN: Datatype access property list <i>Currently not used; specify H5P_DEFAULT.</i>
<i>hid_t trans_id</i>	IN: Transaction identifier specifying the transaction this operation is a part of.
<i>hid_t es_id</i>	IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously. Use H5_EVENT_STACK_NULL for synchronous execution.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

When this routine is executed asynchronously, the return value indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

History:

Added in Quarter 4.

Quarter 5: Changed from transaction to transaction id and from event queue to event stack. Added scope requirement and noted property lists that are not used.

Man Page Status:

No known issues.

4.14.3 H5Tevict_ff

Name: H5Tevict_ff

Signature:

```
herr_t H5Tevict_ff( hid_t dtype_id, uint64_t *container_version, hid_t tapl_id,  
hid_t es_id )
```

Purpose:

Evict named datatype from the burst buffer, possibly asynchronously.

Description:

H5Tevict_ff evicts data associated with a named datatype from the burst buffer.

The data to be evicted may be resident in the burst buffer under two different scenarios

In the first scenario, the data is resident in the burst buffer as the result of updates to the datatype made via the EFF transaction model. For example, through a call to H5Tcommit_ff. This calls add updates to a transaction that are atomically applied when the transaction is committed, and we refer to this data as *transaction update data*.

When evicting transaction update data, the container version being evicted should first be persisted to permanent storage (DAOS), with the H5RCpersist command. The named datatype's transaction update data for the specified container version, as well as the named datatype's transaction update data for all lower-numbered container versions that has not yet been evicted from the burst buffer, will be evicted as the result of this call. If evicting the data would result in container versions with open read contexts becoming inaccessible, the evict will fail.

In the second scenario, the data is resident in the burst buffer as the result of a call to H5Tprefetch_ff. This call replicates data from persistent storage (DAOS) to the burst buffer, and we refer to this data as *replica data*. When replica data is evicted, only data in the burst buffer as a result of the exact replica specified is evicted – transaction update data and other replicas for the named datatype remain in the burst buffer.

The `dtype_id` parameter specifies the named datatype whose data is to be evicted.

The version of the named datatype to be evicted is specified by the `container_version` property.

The evict replica property in the access property list, `tapl_id`, is used to specify that a named datatype replica is to be evicted. H5Pset_evict_replica() sets the evict replica property. If this is set, then replica data will be evicted, otherwise transaction update data will be evicted.

The `es_id` parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in H5_EVENT_STACK_NULL for the `es_id` parameter.

Limitations / Future Considerations:

In Quarter 7, only the primary IOD Blob object associated with the named datatype is evicted. Auxiliary IOD objects will also be evicted in Quarter 8.

When evicting a replica, the `container_version` argument is redundant, as the replica identifier fully specifies the data to be evicted. Consider revisiting and possibly revising the API prior to production release. Possibly have separate evict commands for eviction of transaction update data and of replicas.

For other potential extensions that are beyond the scope of the EFF prototype project, refer to the document *Burst Buffer Space Management – Prototype to Production*.

Parameters:

<i>hid_t</i> <code>datatype_id</code>	IN: Identifier of the named datatype being evicted.
<i>uint64_t</i> <code>container_version</code>	IN: Container version specifying what version of the named datatype to evict.
<i>hid_t</i> <code>tapl_id</code>	IN: Identifier of an access property list. If the access property list contains an evict replica property (set via <code>H5Pset_evict_replica()</code>), then the <code>replica_id</code> specified by that property will be evicted.
<i>hid_t</i> <code>es_id</code>	IN: The <code>es_id</code> parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in <code>H5_EVENT_STACK_NULL</code> for the <code>es_id</code> parameter.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

Note that when this routine is executed asynchronously, the return value from the routine only indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

History:

Added in Quarter 7.

Man Page Status:

Will need updates in Quarter 8 when additional features are implemented.

4.14.4 H5Topen_ff

Name: H5Topen_ff

Signature:

```
hid_t H5Topen_ff(hid_t loc_id, const char * name, hid_t tapl_id, hid_t rcntxt_id,  
hid_t es_id)
```

Purpose:

Opens a named datatype, possibly asynchronously.

Description:

H5Topen_ff opens a named datatype at the location specified by `loc_id` and returns an identifier for the datatype. `loc_id` is either a file or group identifier. `name` is the path to the named datatype object relative to `loc_id`. Both `loc_id` and `name` must be in scope for the read context identified by `rcntxt_id`.

The identifier should eventually be closed by calling H5Tclose_ff to release resources.

The named datatype is opened with the datatype access property list `tapl_id`. The datatype access property list currently has no effect and should be set to H5P_DEFAULT.

`rcntxt_id` indicates the read context for this operation.

The `es_id` parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in H5_EVENT_STACK_NULL for the `es_id` parameter. *NOTE: Currently, only synchronous execution is supported. We will continue to evaluate whether asynchronous execution makes sense for this function, as the user needs the information returned in order to proceed.*

Parameters:

<code>hid_t loc_id</code>	IN: A file or group identifier that is in scope for the read context.
<code>const char * name</code>	IN: A datatype name, defined within the file or group identified by <code>loc_id</code> . Must be in scope for the read context.
<code>hid_t tapl_id</code>	IN: Datatype access property list identifier. <i>Currently not used; specify H5P_DEFAULT.</i>
<code>hid_t rcntxt_id</code>	IN: Read context identifier indicating the read context for this operation.
<code>hid_t es_id</code>	IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously. Use H5_EVENT_STACK_NULL for synchronous execution. <i>Currently will be treated as H5_EVENT_STACK_NULL regardless of what value is passed in.</i>

Returns:

Returns a named datatype identifier if successful; otherwise returns a negative value.

History:

Added in Quarter 4.

4.14.5 H5Tprefetch_ff

Name: H5Tprefetch_FF

Signature:

```
herr_t H5Tprefetch_ff( hid_t dtype_id, hid_t rcntxt_id, hid_t *replica_id, hid_t tapl_id,  
hid_t es_id )
```

Purpose:

Prefetch a named datatype from persistent storage to burst buffer storage, possibly asynchronously.

Description:

H5Tprefetch_ff prefetches a named datatype, specified by its identifier `dtype_id`, from persistent storage (DAOS) into burst buffer storage.

`rcntxt_id` indicates the read context for this operation.

`replica_id`, the replica identifier, is set to indicate where the pre-fetched data can be found in the burst buffer, and is passed to subsequent H5Tevict calls.

`tapl_id`, a datatype access property list identifier, is not currently used and should be set to H5P_DEFAULT.

The `es_id` parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in H5_EVENT_STACK_NULL for the `es_id` parameter.

The `replica_id` is not valid until the operation has completed, if it is executing asynchronously.

Limitations / Future Considerations:

For the EFF prototype project, only the primary IOD Blob object associated with the HDF5 named datatype will be prefetched; auxiliary IOD objects remain on persistent storage (DAOS). For more information, and other potential extensions, refer to the document *Burst Buffer Space Management – Prototype to Production*.

This function was implemented for completeness, and to demonstrate that IOD Blob objects can be prefetched, but since no access routines accept the named datatype's `replica_id` in this phase of the project, it has no practical value other than to demonstrate the ability to prefetch and evict HDF5 named datatypes and IOD Blobs.

Parameters:

<code>hid_t dtype_id</code>	IN: Identifier of the named datatype being prefetched. <code>grp_id</code> must be in scope for the read context.
<code>hid_t rcntxt_id</code>	IN: Read context identifier indicating the read context for this operation.
<code>hid_t *replica_id</code>	IN: Identifier of the replicated data in the burst buffer.
<code>hid_t tapl_id</code>	IN: Identifier of an access property list for this I/O operation. Should be H5P_DEFAULT.
<code>hid_t es_id</code>	IN: The <code>es_id</code> parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in H5_EVENT_STACK_NULL for the <code>es_id</code> parameter.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

Note that when this routine is executed asynchronously, the return value from the routine only indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

History:

Added in Quarter 7.

Man Page Status:

No known issues.

4.15 H5TR: Transaction Operation APIs

These routines are used to create and manage HDF5 transactions in the Exascale Fast Forward stack.

An HDF5 transaction consists of a set of updates that modify an HDF5 file (container).

Transactions are applied atomically - either all of the transaction's updates are applied and become readable (the transaction is committed) or none of them are applied (the transaction is discarded). Transactions are numbered and are committed in strict numerical sequence.

Multiple transactions may be in progress concurrently, and multiple processes may participate in a single transaction.

A transaction is associated with a particular container and transaction number. The transaction ID returned when a transaction is created is used to start the transaction, and is passed to HDF5 functions that update the container. This allows every update to be added to the correct transaction, and all of the updates for a transaction to be applied atomically when the transaction is committed.

These routines did not exist prior to the Exascale FastForward project.

Routine	Implemented	Notes
H5TRabort	Quarter 5	
H5TRclose	Quarter 5	
H5TRcreate	Quarter 5	
H5TRfinish	Quarter 5	
H5TRget_trans_num	Quarter 7	
H5TRget_version	Quarter 7	
H5TRset_dependency	Quarter 5	
H5TRskip	Quarter 5	
H5TRstart	Quarter 5	

4.15.1 H5TRabort

Name: H5TRabort

Signature:

```
herr_t H5TRabort( hid_t trans_id, hid_t es_id )
```

Purpose:

Abort a transaction.

Description:

H5TRabort signals that all updates in the transaction identified by `trans_id` should be discarded.

H5TRabort can be called by any process that has a transaction identifier for the transaction to be aborted. All past and future updates in the transaction identified by `trans_id` will be discarded. The aborted transaction will never be committed to the container and the transaction number associated with the aborted transaction cannot be reused. An aborted transaction is classified as *discarded*. Discarded transactions do not block the commitment of higher-numbered transactions.

If a higher-numbered transaction registered a dependency on the aborted transaction, the dependent transaction will also be aborted. The abort of the dependent transaction may occur at any time before it commits.

The `es_id` parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in `H5_EVENT_STACK_NULL` for the `es_id` parameter.

Parameters:

`hid_t trans_id` IN: Transaction identifier for the transaction that is being aborted.

`hid_t es_id` IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously. Use `H5_EVENT_STACK_NULL` for synchronous execution.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

When this routine is executed asynchronously, the return value indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

History:

Added in Quarter 5.

Man Page Status:

No known issues.

4.15.2 H5TRclose

Name: H5TRclose

Signature:

herr_t H5TRclose (hid_t trans_id)

Purpose:

Close the specified transaction.

Description:

`H5TRclose` closes the transaction specified by `trans_id` and releases resources used by it. Further use of the transaction identifier is illegal in HDF5 API calls. This function does not finish the transaction and allow it to be committed – that is done by `H5TRfinish`.

All processes that called `H5TRcreate` to create a transaction and get a transaction id must call `H5TRclose` to release the resources used by the transaction. For a transaction that is to be committed, the function `H5TRfinish` must be called L times, where L is the number of leader processes that called `H5TRstart` to start the transaction. Typically the leader processes each call `H5TRfinish`, but in some circumstances – such as if a process dies – other processes may call `H5TRfinish`. Any process with a transaction id for a given transaction may abort the transaction with a call to `H5TRabort`. On a given process for a given transaction, calls to `H5TRfinish` or `H5TRabort` should be made – and completed – before `H5TRclose` is called.

The actions performed by this function are all local to the compute node so an asynchronous execution option is not provided.

Parameters:

`hid_t trans_id` IN: Identifier of the transaction to close.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

History:

Added in Quarter 5.

Man Page Status:

No known issues.

4.15.3 H5TRcreate

Name: H5TRcreate

Signature:

```
hid_t H5TRcreate ( hid_t file_id, hid_t rcntxt_id, uint64_t trans_num )
```

Purpose:

Create a transaction associated with a specified container, read context, and number.

Description:

`H5TRcreate` creates a transaction associated with a specified container, read context, and transaction number. Every process that will add updates to the transaction must call `H5TRcreate` and use the transaction identifier returned by `H5TRcreate` in subsequent calls related to the transaction.

After a transaction has been created, one or more transaction leaders must start the transaction by calling `H5TRstart`, effectively notifying lower layers of the I/O stack that the transaction is in-progress. After the leader processes start the transaction, it is their responsibility to notify the processes that called `H5TRcreate` but did not call `H5TRstart` (the delegate processes) that updates to the transaction may begin.

The `file_id` parameter indicates the container the transaction is associated with, and must be open for write. Updates added to the transaction will be applied atomically to the container identified by `file_id` when the transaction is committed.

The `rcntxt_id` parameter indicates read context for the transaction. Any read operations that occur in conjunction with updates to the transaction will be made on the container version associated with this read context.

The `trans_num` parameter indicates the number that is associated with the transaction. A single container can have multiple transactions in progress at any given time. Transactions are committed in strict numerical order, and it is the application's responsibility to manage transaction numbers. Transaction numbers uniquely identify a transaction and once a transaction is started or skipped the number can never be used again.

Note that although reads will be made from the container version associated with the read context, the transaction's updates will be atomically applied to the latest container version at the time the transaction is committed.

The actions performed by this function are all local to the compute node so an asynchronous execution option is not provided.

Parameters:

`hid_t file_id` IN: File identifier indicating the container that the transaction is associated with.

`hid_t rcntxt_id` IN: Read context identifier indicating the read context for the transaction.

`uint64_t trans_num` IN: The transaction number for the transaction.

Returns:

Returns a transaction identifier if successful; otherwise returns a negative value.

History:

Added in Quarter 5.

Man Page Status:

No known issues.

4.15.4 H5TRfinish

Name: H5TRfinish

Signature:

```
herr_t H5TRfinish( hid_t trans_id, hid_t trfp1_id, hid_t *rcntxt_id, hid_t es_id )
```

Purpose:

Finish a transaction that was started with H5TRstart.

Description:

H5TRfinish signals that no more updates will be added to the transaction identified by `trans_id`. If updates to the transaction have been made with asynchronous function calls, then all of the operations must complete before the last call to H5TRfinish is made for the transaction.

For a given transaction, H5TRfinish must be called the same number of times that H5TRstart was called to start the transaction (assuming that the transaction was not aborted with H5TRabort). Typically, the leader processes that called H5TRstart to start the transaction each call H5TRfinish; in some circumstances – such as when a process dies – other processes may call H5TRfinish.

The transaction is considered finished when H5TRfinish has been called the required number of times.

After the transaction is finished and all lower-numbered transactions are finished or discarded, the transaction is *committed* and all updates that are part of the transaction will be applied atomically to the container and become readable in a new container version with the same number as the committed transaction.

The transaction finish property list, `trfp1_id`, is currently not used and should be set to H5P_DEFAULT.

The `rcntxt_id` parameter allows the user to pass a pointer to an `hid_t` datatype signaling that the function should acquire a read handle and create a read context immediately after the transaction is committed. This ensures that a read handle on the container version resulting from the committed transaction can be acquired before a subsequent transaction commit occurs.

In addition to acquiring the read handle, a read context is created and the identifier for that context is returned in `*rcntxt_id`. The read context identifier will not be available until after the H5TRfinish call completes successfully, but a future id is returned immediately and can be used in subsequent HDF5 calls. H5RCrelease must be called to release the read handle and close the context when it is no longer needed.

If the `rcntxt_id` parameter is set to NULL, no read handle is acquired and no read context is created.

The `es_id` parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in H5_EVENT_STACK_NULL for the `es_id` parameter.

When executed asynchronously, the function does not complete until after the finished transaction has been committed. If `rcntxt_id` is not NULL, acquisition of the read handle, creation of the read context, and updating of `*rcntxt_id` will also occur before the function completes.

After H5TRfinish completes, the transaction should be closed by all processes that called H5TRcreate.

Parameters:

`hid_t trans_id` IN: Transaction identifier for the transaction being finished.

`hid_t trfp1_id` IN: Transaction finish property list.

Currently not used; specify H5P_DEFAULT.

*hid_t *rcntxt_id* IN/OUT: Pointer to read context for container version resulting from successfully committed transaction. Pass in NULL if no read context is desired.

hid_t es_id IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously. Use H5_EVENT_STACK_NULL for synchronous execution.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value. When executed asynchronously, a future ID for the read context is returned initially in **rcntxt_id*; upon completion of the asynchronous operation, the future ID will be transparently modified to be a “normal” read context identifier.

When this routine is executed asynchronously, the return value indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

History:

Added in Quarter 5.

Man Page Status:

No known issues.

4.15.5 H5TRget_trans_num

Name: H5TRget_trans_num

Signature:

*herr_t H5TRget_trans_num(hid_t trans_id, uint64_t *trans_num)*

Purpose:

Retrieve the transaction number associated with a transaction.

Description:

The `trans_id` specifies the transaction whose associated transaction number is to be retrieved.

The value pointed to by the `trans_num` parameter will be set to the transaction number that the transaction is associated with.

The actions performed by this function are all local to the compute node so an asynchronous execution option is not provided.

Parameters:

`hid_t trans_id`

IN: Identifier of the transaction whose associated transaction number is being retrieved.

`uint64_t *trans_num`

OUT: Transaction number associated with the transaction.
Value is unchanged if operation is unsuccessful.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

History:

Added in Quarter 7.

Man Page Status:

No known issues.

4.15.6 H5TRget_version

Name: H5TRget_version

Signature:

```
herr_t H5TRget_version( hid_t trans_id, uint64_t *container_version )
```

Purpose:

Retrieve the container version associated with a transaction.

Description:

The `trans_id` specifies the transaction whose associated container version is to be retrieved.

When a transaction is started via the `H5TRcreate` call, a read context is associated with the transaction. The container version returned by this `H5TRget_version` call is the one associated with the transaction is the container version associated with that read context.

The value pointed to by the `container_version` parameter will be set to the container version that transaction is associated with.

The actions performed by this function are all local to the compute node so an asynchronous execution option is not provided.

Parameters:

<code>hid_t trans_id</code>	IN: Identifier of the transaction whose associated container version is being retrieved.
<code>uint64_t *container_version</code>	OUT: Container version associated with the transaction. Value is unchanged if operation is unsuccessful.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

History:

Added in Quarter 7.

Man Page Status:

No known issues.

4.15.7 H5TRset_dependency

Name: H5TRset_dependency

Signature:

```
herr_t H5TRset_dependency( hid_t trans_id, uint64_t trans_num, hid_t es_id )
```

Purpose:

Register the dependency of a transaction on a lower-numbered transaction.

Description:

H5TRset_dependency registers a dependency of the transaction identified by `trans_id` on a lower-numbered transaction whose transaction number is given by `trans_num`. The transaction with the dependency is referred to as *dependent transaction* (identified by `trans_id`) and the lower-numbered transaction it depends on is the *prerequisite transaction* (identified by `trans_num`).

A dependency must be registered when a commit of the *dependent transaction* will leave the container in an inconsistent state unless the *prerequisite transaction* commits.

If a prerequisite transaction is discarded (either aborted or skipped) the dependent transaction will be aborted. The abort of the dependent transaction may occur at any time before it commits.

A given transaction may depend on multiple prerequisite transactions. It is an error to register a dependency on a future (higher-numbered) transaction.

The `es_id` parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in `H5_EVENT_STACK_NULL` for the `es_id` parameter.

Parameters:

<code>hid_t trans_id</code>	IN: Transaction identifier for the transaction with the dependency – the dependent transaction.
<code>uint64_t trans_num</code>	IN: Transaction number of the lower-numbered transaction that must commit successfully -- the prerequisite transaction.
<code>hid_t es_id</code>	IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously. Use <code>H5_EVENT_STACK_NULL</code> for synchronous execution.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

When this routine is executed asynchronously, the return value indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

History:

Added in Quarter 5.

Man Page Status:

No known issues.

4.15.8 H5TRskip

Name: H5TRskip

Signature:

```
herr_t H5TRskip( hid_t file_id, uint64_t start_trans_num, uint64_t count, hid_t es_id )
```

Purpose:

Explicitly skip one or more transaction numbers for a container.

Description:

H5TRskip signals that the application will not be using the identified transaction numbers for the specified container. Since transactions are numbered and are committed in strict numeric order, having unused transaction numbers that are not explicitly skipped will prevent higher-numbered transactions from committing.

The `file_id` is the file identifier for a container that is open for write.

`start_trans_num` is the first transaction number that should be skipped.

`count` specifies how many transaction numbers should be skipped.

The `es_id` parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in `H5_EVENT_STACK_NULL` for the `es_id` parameter.

H5TRskip should always be called by a single process for any given container and transaction number.

H5TRskip should never be called with the same `file_id` and `transaction_num` parameters used in a call to H5TRcreate. The (uncreated) transactions with the skipped transaction numbers are classified as *discarded*. Discarded transactions do not block the commitment of higher-numbered transactions.

Parameters:

<code>hid_t file_id</code>	IN: File identifier for container open for write.
<code>uint64_t start_trans_num</code>	IN: The first transaction number to skip.
<code>uint64_t count</code>	IN: How many transaction numbers to skip.
<code>hid_t es_id</code>	IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously. Use <code>H5_EVENT_STACK_NULL</code> for synchronous execution.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

When this routine is executed asynchronously, the return value indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

History:

Added in Quarter 5.

Man Page Status:

No known issues.

4.15.9 H5TRstart

Name: H5TRstart

Signature:

```
herr_t H5TRstart( hid_t trans_id, hid_t trspl, hid_t es_id )
```

Purpose:

Start a created transaction.

Description:

H5TRstart starts the previously created transaction identified by trans_id.

After a transaction has been created with H5TRcreate, one or more transaction leaders must start the transaction by calling H5TRstart, effectively notifying lower layers of the I/O stack that the transaction is in-progress. The transaction is considered in-progress when the first H5TRstart for the transaction completes successfully.

Updates of the container associated with the transaction may be added to the transaction after the transaction has been started. The transaction leaders may issue updates at any time after H5TRstart had been called. The transaction leader(s) must notify the delegates (processes that called H5TRcreate but did not call H5TRstart) when they can begin adding updates to the transaction. Unlike the leaders, the delegates cannot add updates until after at least one of the asynchronous H5TRstart operations for the transaction completes successfully.

The trspl parameter identifies the transaction start property list for the call. Currently the list can contain only the leader count property that specifies the total number of leader processes that will call H5TRstart for the transaction identified by trans_id. If there is a single transaction leader, trspl should be set to H5P_DEFAULT. If there are multiple transaction leaders, then each must pass in a transaction start property list specifying the same number of leaders. H5Pset_trspl_num_peers() is used to set the leader count property.

The es_id parameter indicates the event stack the event object for this call should be pushed onto when the function is executed asynchronously. The function may be executed synchronously by passing in H5_EVENT_STACK_NULL for the es_id parameter.

For a given transaction, H5TRfinish must be called the same number of times that H5TRstart was called to start the transaction (assuming that the transaction was not aborted with H5TRabort). Typically, the leader processes that called H5TRstart to start the transaction each call H5TRfinish; in some circumstances – such as when a process dies – other processes may call H5TRfinish.

The transaction is considered finished when H5TRfinish has been called the required number of times.

After the transaction is finished and all lower-numbered transactions are finished or discarded (by H5TRabort or H5TRskip), the transaction is *committed* and all updates that are part of the transaction will be applied atomically to the container and become readable in a new container version with the same number as the committed transaction.

Parameters:

hid_t trans_id IN: Transaction identifier indicating the transaction to start.

hid_t trspl IN: Transaction start property list. Should be set to H5P_DEFAULT if only one process will call H5TRstart for the transaction. If multiple processes will make the call, the transaction start property list must contain the leader count property, set with

`H5Pset_trspl_num_peers()`.

hid_t es_id IN: Event stack identifier specifying the event stack that will be used to monitor the status of the event associated with this function call when executed asynchronously. Use `H5_EVENT_STACK_NULL` for synchronous execution.

Returns:

Returns a non-negative value if successful; otherwise returns a negative value.

When this routine is executed asynchronously, the return value indicates whether the operation has been successfully scheduled for asynchronous execution. The actual success or failure of the asynchronous operation must be checked separately through the event stack.

History:

Added in Quarter 5.

Man Page Status:

No known issues.

5 Description of example programs

All the example programs that are included in the source code are heavily commented to describe the functionality being demonstrated. A higher-level description of the example programs is included in this section.

Figure 1 shows the relationship of the client and server example programs to the layers of the EFF stack being demonstrated in this milestone.

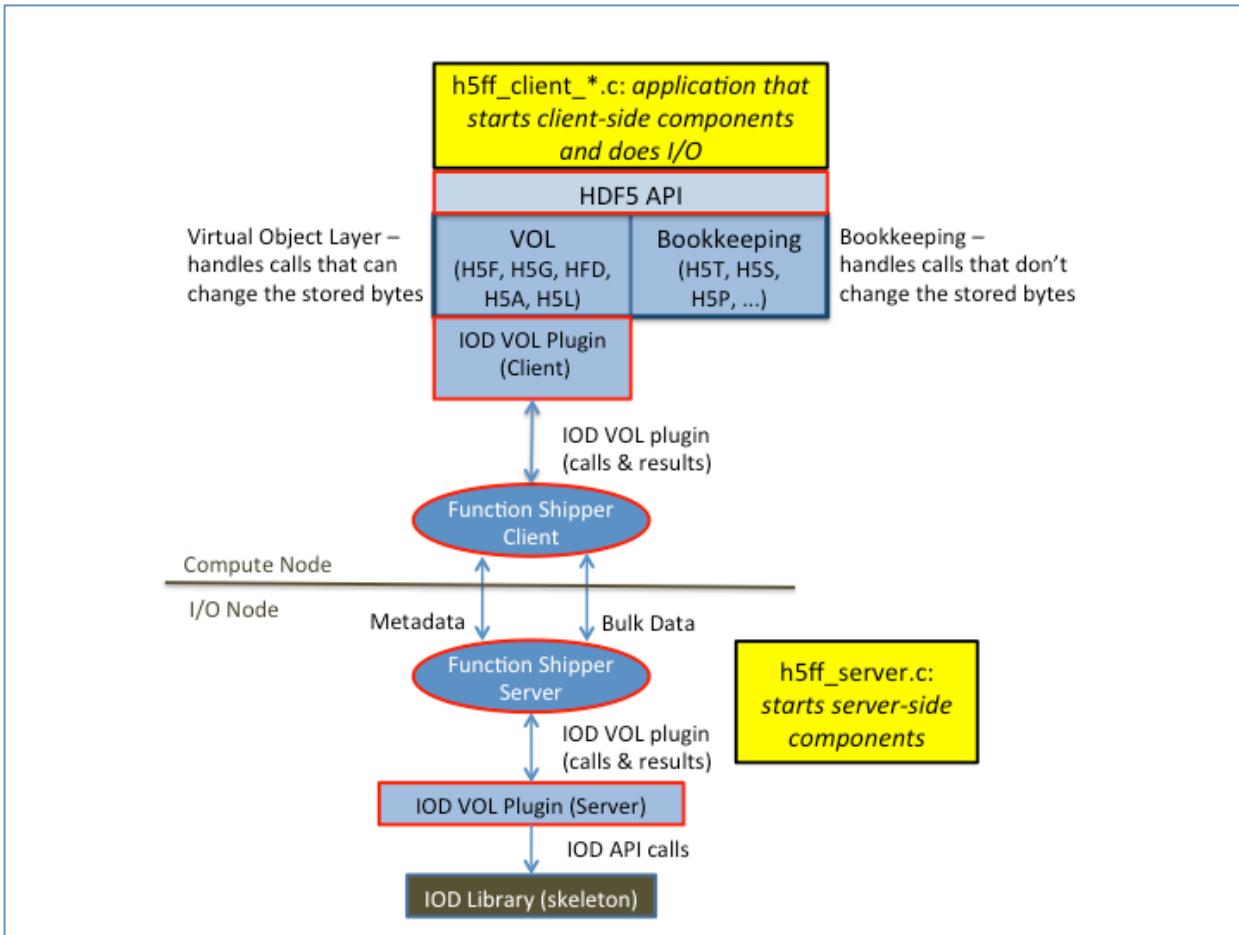


Figure 1: Client and server programs demonstration components

5.1 h5ff_server.c

The h5ff_server.c program launches an MPI application that acts as a server. This program will run on the IONs in the EFF architecture, but is being run on the same multi-processor system as clients for Q5 Milestones.

h5ff_server.c initializes the function shipper interfaces for metadata and bulk data, registers all the required callbacks to handle HDF5 VOL operations, and listens to incoming requests from clients. The first request received from clients is an initialization call that starts up the IOD library for the client and creates an AXE engine to insert future operations into. The last call

Use or disclosure of data contained on this sheet is subject to the restriction on the title page of this document.
Copyright © The HDF Group, 2014. All rights reserved

received from clients is a terminate call that will eventually shut down all the initialized libraries, once all clients have issued this call.

Note that the server program is independent from the client program, meaning that any client that uses HDF5 with the HDF5 IOD VOL plugin can connect to the server.

5.2 h5ff_client_*.c

Each of the client programs in the examples directory of the hdf5_ff source tree launches an MPI application that acts as a client. These programs will run on the CNs in the EFF architecture, but are being run on the same multi-processor system as h5ff_server.c for Q5 Milestones. The client programs demonstrate how an application would use the EFF I/O stack. All client programs have been updated in Q5 to use transactions and read contexts. This is a list of all the client programs included with a note to indicate what each one is testing (note that the comments inside each program explain in much more detail the flow of execution and usage of the HDF5 routines and semantics):

- h5ff_client_analysis.c: This tests the analysis shipping functionality (H5AS).
- h5ff_client_attr.c : This tests attribute routines (H5A).
- h5ff_client_dset.c : This tests dataset routines (H5D).
- h5ff_client_evict_deltas.c: This tests eviction of transaction updates (deltas).
- h5ff_client_links.c : This tests Links routines (H5L).
- h5ff_client_map.c : This tests the new Map routines (H5M) added to support Dynamic Data Structures.
- h5ff_client_multiple_cont.c : This tests access to multiple containers.
- h5ff_client_obj.c : This tests generic object routines (H5O).
- h5ff_client_prefetch.c: This tests prefetch and eviction of replicas.
- h5ff_client_view.c: This tests query and view routines (H5Q and H5V)
- h5ff_client_vl_data.c: This tests variable length data.
- h5ff_client_M6.2_demo.c: HDF5 and I/O Dispatcher Container Versioning Demonstration.
- h5ff_client_M7.2-pep_demo.c: Persist, Evict, and Prefetch data movement demonstration.

6 Instructions for building and running demo code

These instructions can also be found in the file "EFF_INSTALL.txt" in the top directory of the distributed tarball.

System pre-requisites:

```
An MPI3 implementation that supports MPI_THREAD_MULTIPLE  
[i.e. MPICH3 or later - we built with MPICH 3.0.4 in Q6]  
Pthread  
BOOST
```

build opa:

```
It is located in the tarball in the 'openpa' subdirectory, OR get it  
from here: git clone git://git.mcs.anl.gov/radix/openpa.git
```

```
./configure --prefix=/path/to/opa/install/directory  
make  
make check  
make install
```

build AXE:

```
It is located in the tarball in the 'axe' subdirectory, OR get it from  
here: svn checkout https://svn.hdfgroup.uiuc.edu/axe/trunk
```

```
./configure --prefix=/path/to/axe/install/directory --with-  
opa=/path/to/opa/install/directory  
make  
make check  
make install
```

build DAOS, PLFS, and IOD:

```
Please refer to the IOD tarball for instruction on how to build and setup the  
three libraries.
```

build Mercury (Function Shipper)

```
The code is located in tarball in the 'mercury' directory.
```

```
refer to the mercury build recipe in:  
mercury/README
```

build HDF5 IOD VOL plugin:

```
The code is located in the tarball in the 'hdf5_ff' subdirectory, OR  
get it from here:  
svn checkout http://svn.hdfgroup.uiuc.edu/hdf5/features/hdf5_ff
```

```
./configure --with-daos=/path/to/daos posix --with-plfs=/path/to/plfs --with-  
iod=/path/to/iod/ --with-axe=/path/to/axe/install/directory  
PKG_CONFIG_PATH=/path/to/mercury/install/directory/lib/pkgconfig:/path/to/mchecksu  
m/install/dir --enable-parallel --enable-debug --enable-trace --enable-threadsafe -  
-enable-unsupported --with-pthread=/usr --enable-eff --enable-shared --enable-  
python
```

```
If you want indexing to be built in add --enable-indexing.
```

```
Note in that case all 3rd party libraries have to be build shared or with -fPIC.  
You should also have devel python devel libraries and numpy installed on your  
system.
```

You should see in the configure summary at the end if the EFF plugin in HDF5 was successfully configured.

```
make  
make check  
make install
```

build the example programs:

The examples are located in hdf5_ff/examples/.

The server is h5ff_server.

The client programs are

- h5ff_client_attr.c : This tests attribute routines (H5A).
- h5ff_client_dset.c : This tests dataset routines (H5D).
- h5ff_client_links.c : This tests Links routines (H5L).
- h5ff_client_map.c : This tests the new Map routines (H5M) added this quarter to support Dynamic Data Structures.
- h5ff_client_multiple_cont.c : This tests access to multiple containers.
- h5ff_client_obj.c : This tests generic object routines (H5O).
- h5ff_client_analysis.c : This tests the analysis shipping functionality (H5AS).
- h5ff_client_M6.2_demo.c: HDF5 and I/O Dispatcher Container Versioning Demonstration
- h5ff_client_M7.2-pep_demo.c: Prefetch, evict, persist data movement demo.

```
cd path/where/hdf5_ff/is/built/examples/  
make  
chmod 775 run_ff_tests.sh  
./run_ff_tests.sh num_server_procs num_client_procs
```

Or you can run each test manually:

The client and server need to be launched from the same directory for now.

Launch the server first:

```
mpiexec -n <x> ./h5ff_server  
then launch one of the clients  
mpiexec -n <x> ./h5ff_client_xx
```

Note, for now, the number of clients must be greater than or equal to the number of servers.

END